

The Mercury User's Guide

Version rotd-2010-06-17

Fergus Henderson
Thomas Conway
Zoltan Somogyi
Peter Ross
Tyson Dowd
Mark Brown
Ian MacLarty
Paul Bone

Copyright © 1995–2010 The University of Melbourne.

Permission is granted to make and distribute verbatim copies of this manual provided the copyright notice and this permission notice are preserved on all copies.

Permission is granted to copy and distribute modified versions of this manual under the conditions for verbatim copying, provided also that the entire resulting derived work is distributed under the terms of a permission notice identical to this one.

Permission is granted to copy and distribute translations of this manual into another language, under the above conditions for modified versions.

Table of Contents

1 Introduction

This document describes the compilation environment of Mercury. It describes how to use ‘`mmc`’, the Mercury compiler; a build tool integrated into the compiler called ‘`mmc --make`’; an older tool, ‘`mmake`’, built on top of ordinary or GNU make to simplify the handling of Mercury programs; how to use ‘`mdb`’, the Mercury debugger; and how to use ‘`mprof`’, the Mercury profiler.

We strongly recommend that programmers use ‘`mmc --make`’ rather than invoking ‘`mmc`’ directly, because ‘`mmc --make`’ is generally easier to use and avoids unnecessary recompilation.

2 File naming conventions

Mercury source files must be named ‘`*.m`’. Each Mercury source file should contain a single Mercury module whose module name should be the same as the filename without the ‘`.m`’ extension.

The Mercury implementation uses a variety of intermediate files, which are described below. But all you really need to know is how to name source files. For historical reasons, the default behaviour is for intermediate files to be created in the current directory, but if you use the ‘`--use-subdirs`’ option to ‘`mmc`’ or ‘`mmake`’, all these intermediate files will be created in a ‘`Mercury`’ subdirectory, where you can happily ignore them. Thus you may wish to skip the rest of this chapter.

In cases where the source file name and module name don’t match, the names for intermediate files are based on the name of the module from which they are derived, not on the source file name.

Files ending in ‘`.int`’, ‘`.int0`’, ‘`.int2`’ and ‘`.int3`’ are interface files; these are generated automatically by the compiler, using the ‘`--make-interface`’ (or ‘`--make-int`’), ‘`--make-private-interface`’ (or ‘`--make-priv-int`’), ‘`--make-short-interface`’ (or ‘`--make-short-int`’) options. Files ending in ‘`.opt`’ are interface files used in inter-module optimization, and are created using the ‘`--make-optimization-interface`’ (or ‘`--make-opt-int`’) option. Similarly, files ending in ‘`.trans_opt`’ are interface files used in transitive inter-module optimization, and are created using the ‘`--make-transitive-optimization-interface`’ (or ‘`--make-trans-opt-int`’) option.

Since the interface of a module changes less often than its implementation, the ‘`.int`’, ‘`.int0`’, ‘`.int2`’, ‘`.int3`’, ‘`.opt`’, and ‘`.trans_opt`’ files will remain unchanged on many compilations. To avoid unnecessary recompilations of the clients of the module, the timestamps on these files are updated only if their contents change. ‘`.date`’, ‘`.date0`’, ‘`.date3`’, ‘`.optdate`’, and ‘`.trans_opt_date`’ files associated with the module are used as timestamp files; they are used when deciding whether the interface files need to be regenerated.

‘`.c_date`’, ‘`.il_date`’, ‘`.java_date`’, ‘`.erl_date`’, ‘`.s_date`’ and ‘`.pic_s_date`’ files perform a similar function for ‘`.c`’, ‘`.il`’, ‘`.java`’, ‘`.erl`’, ‘`.s`’ and ‘`.pic_s`’ files respectively. When smart recompilation (see [\[Auxiliary output options\]](#), page [\(undefined\)](#)) works out that a module does not need to be recompiled, the timestamp file for the target file is updated, and the timestamp of the target file is left unchanged.

‘.used’ files contain dependency information for smart recompilation (see [\[Auxiliary output options\]](#), page [\[undefined\]](#)).

Files ending in ‘.d’ are automatically-generated Makefile fragments which contain the dependencies for a module. Files ending in ‘.dep’ are automatically-generated Makefile fragments which contain the rules for an entire program. Files ending in ‘.dv’ are automatically-generated Makefile fragments which contain variable definitions for an entire program.

As usual, ‘.c’ files are C source code, and ‘.o’ files are object code. In addition, ‘.pic_o’ files are object code files that contain position-independent code (PIC). ‘.lpic_o’ files are object code files that can be linked with shared libraries, but don’t necessarily contain position-independent code themselves. ‘.mh’ and ‘.mih’ files are C header files generated by the Mercury compiler. The non-standard extensions are necessary to avoid conflicts with system header files. ‘.s’ files and ‘.pic_s’ files are assembly language. ‘.java’, ‘.class’ and ‘.jar’ files are Java source code, Java bytecode and Java archives respectively. ‘.il’ files are Intermediate Language (IL) files for the .NET Common Language Runtime. ‘.erl’ and ‘.beam’ files are Erlang source code and bytecode (object) files respectively. ‘.beams’ directories are collections of ‘.beam’ files which act like a library archive.

3 Using the Mercury compiler

Following a long Unix tradition, the Mercury compiler is called ‘mmc’ (for “Melbourne Mercury Compiler”). Some of its options (e.g. ‘-c’, ‘-o’, and ‘-I’) have a similar meaning to that in other Unix compilers.

Arguments to ‘mmc’ may be either file names (ending in ‘.m’), or module names, with ‘.’ (rather than ‘__’ or ‘:’) as the module qualifier. For a module name such as ‘foo.bar.baz’, the compiler will look for the source in files ‘foo.bar.baz.m’, ‘bar.baz.m’, and ‘baz.m’, in that order. Note that if the file name does not include all the module qualifiers (e.g. if it is ‘bar.baz.m’ or ‘baz.m’ rather than ‘foo.bar.baz.m’), then the module name in the ‘:- module’ declaration for that module must be fully qualified. To make the compiler look in another file for a module, use ‘mmc -f *sources-files*’ to generate a mapping from module name to file name, where *sources-files* is the list of source files in the directory (see [\[Output options\]](#), page [\[undefined\]](#)).

Arguments to ‘mmc’ may also be in ‘@file’. The ‘@file’ argument is replaced with arguments representing the contents of the file. This argument processing is done recursively. The contents of the ‘@file’ is split into arguments one per line in the file.

To compile a program which consists of just a single source file, use the command

```
mmc filename.m
```

Unlike traditional Unix compilers, however, ‘mmc’ will put the executable into a file called ‘*filename*’, not ‘a.out’.

For programs that consist of more than one source file, you can use Mmake (see [\[Using Mmake\]](#), page [\[undefined\]](#)) or the ‘--make’ option to ‘mmc’. Currently, the use of ‘mmc --make’ is recommended:

```
mmc --make filename
```

If you use Mmake or ‘mmc --make’, then you don’t need to understand the details of how the Mercury implementation goes about building programs. Thus you may wish to skip the rest of this chapter.

To compile a source file to object code without creating an executable, use the command

```
mmc -c filename.m
```

‘mmc’ will put the object code into a file called ‘*module.o*’, where *module* is the name of the Mercury module defined in ‘*filename.m*’. It also will leave the intermediate C code in a file called ‘*module.c*’. If the source file contains nested modules, then each sub-module will get compiled to separate C and object files.

Before you can compile a module, you must make the interface files for the modules that it imports (directly or indirectly). You can create the interface files for one or more source files using the following commands:

```
mmc --make-short-int filename1.m filename2.m ...
mmc --make-priv-int filename1.m filename2.m ...
mmc --make-int filename1.m filename2.m ...
```

If you are going to compile with ‘--intermodule-optimization’ enabled, then you also need to create the optimization interface files.

```
mmc --make-opt-int filename1.m filename2.m ...
```

If you are going to compile with ‘--transitive-intermodule-optimization’ enabled, then you also need to create the transitive optimization files.

```
mmc --make-trans-opt filename1.m filename2.m ...
```

Given that you have made all the interface files, one way to create an executable for a multi-module program is to compile all the modules at the same time using the command

```
mmc filename1.m filename2.m ...
```

This will by default put the resulting executable in ‘*filename1*’, but you can use the ‘-o *filename*’ option to specify a different name for the output file, if you so desire.

The other way to create an executable for a multi-module program is to compile each module separately using ‘mmc -c’, and then link the resulting object files together. The linking is a two stage process.

First, you must create and compile an *initialization file*, which is a C source file containing calls to automatically generated initialization functions contained in the C code of the modules of the program:

```
c2init module1.c module2.c ... > main-module_init.c,
mgnuc -c main-module_init.c
```

The ‘c2init’ command line must contain the name of the C file of every module in the program. The order of the arguments is not important. The ‘mgnuc’ command is the Mercury GNU C compiler; it is a shell script that invokes the GNU C compiler ‘gcc’ with the options appropriate for compiling the C programs generated by Mercury.

You then link the object code of each module with the object code of the initialization file to yield the executable:

```
m1 -o main-module module1.o module2.o ... main-module_init.o
```

‘m1’, the Mercury linker, is another shell script that invokes a C compiler with options appropriate for Mercury, this time for linking. ‘m1’ also pipes any error messages from the linker through ‘mdemangle’, the Mercury symbol demangler, so that error messages refer to predicate and function names from the Mercury source code rather than to the names used in the intermediate C code.

The above command puts the executable in the file ‘*main-module*’. The same command line without the ‘-o’ option would put the executable into the file ‘*a.out*’.

‘*mmc*’ and ‘*ml*’ both accept a ‘-v’ (verbose) option. You can use that option to see what is actually going on. For the full set of options of ‘*mmc*’, see [\[Invocation\]](#), page [\[undefined\]](#).

4 Running programs

Once you have created an executable for a Mercury program, you can go ahead and execute it. You may however wish to specify certain options to the Mercury runtime system. The Mercury runtime accepts options via the ‘*MERCURY_OPTIONS*’ environment variable. The most useful of these are the options that set the size of the stacks. (For the full list of available options, see [\[Environment\]](#), page [\[undefined\]](#).)

The det stack and the nondet stack are allocated fixed sizes at program start-up. The default size is 1024k times the word size (in bytes) for the det stack and 64k times the word size (in bytes) for the nondet stack, but these can be overridden with the ‘*--detstack-size*’ and ‘*--nondetstack-size*’ options, whose arguments are the desired sizes of the det and nondet stacks respectively, in units of kilobytes. On operating systems that provide the appropriate support, the Mercury runtime will ensure that stack overflow is trapped by the virtual memory system.

With conservative garbage collection (the default), the heap will start out with a zero size, and will be dynamically expanded as needed. When not using conservative garbage collection, the heap has a fixed size like the stacks. The default size is 8Mb times the word size (in bytes), but this can be overridden with the ‘*--heap-size*’ option.

5 Using Mmake

Mmake, short for “Mercury Make”, is a tool for building Mercury programs. The same functionality is now provided in ‘*mmc*’ directly by using the ‘*--make*’ option:

```
mmc --make main-module
```

The usage of Mmake is discouraged.

Mmake is built on top of ordinary or GNU Make¹. With Mmake, building even a complicated Mercury program consisting of a number of modules is as simple as

```
mmc -f source-files
mmake main-module.depend
mmake main-module
```

Mmake only recompiles those files that need to be recompiled, based on automatically generated dependency information. Most of the dependencies are stored in ‘.d’ files that are automatically recomputed every time you recompile, so they are never out-of-date. A little bit of the dependency information is stored in ‘.dep’ and ‘.dv’ files which are more expensive to recompute. The ‘*mmake main-module.depend*’ command which recreates the

¹ We might eventually add support for ordinary “Make” programs, but currently only GNU Make is supported.

`'main-module.dep'` and `'main-module.dv'` files needs to be repeated only when you add or remove a module from your program, and there is no danger of getting an inconsistent executable if you forget this step — instead you will get a compile or link error.

The `'mmc -f'` step above is only required if there are any source files for which the file name does not match the module name. `'mmc -f'` generates a file `'Mercury.modules'` containing a mapping from module name to source file. The `'Mercury.modules'` file must be updated when a source file for which the file name does not match the module name is added to or removed from the directory.

`'mmake'` allows you to build more than one program in the same directory. Each program must have its own `'dep'` and `'dv'` files, and therefore you must run `'mmake program.depend'` for each program. The `'Mercury.modules'` file is used for all programs in the directory.

If there is a file called `'Mmake'` or `'Mmakefile'` in the current directory, Mmake will include that file in its automatically-generated Makefile. The `'Mmake'` file can override the default values of various variables used by Mmake's builtin rules, or it can add additional rules, dependencies, and actions.

Mmake's builtin rules are defined by the file `'prefix/lib/mercury/mmake/Mmake.rules'` (where *prefix* is `'/usr/local/mercury-version'` by default, and *version* is the version number, e.g. `'0.6'`), as well as the rules and variables in the automatically-generated `'dep'` and `'dv'` files. These rules define the following targets:

`'main-module.depend'`

Creates the files `'main-module.dep'` and `'main-module.dv'` from `'main-module.m'` and the modules it imports. This step must be performed first. It is also required whenever you wish to change the level of inter-module optimization performed (see [\[Overall optimization options\]](#), [page \[Overall optimization options\]](#)).

`'main-module.ints'`

Ensure that the interface files for *main-module* and its imported modules are up-to-date. (If the underlying `'make'` program does not handle transitive dependencies, this step may be necessary before attempting to make `'main-module'` or `'main-module.check'`; if the underlying `'make'` is GNU Make, this step should not be necessary.)

`'main-module.check'`

Perform semantic checking on *main-module* and its imported modules. Error messages are placed in `'err'` files.

`'main-module'`

Compiles and links *main-module* using the Mercury compiler. Error messages are placed in `'err'` files.

`'main-module.javas'`

Compiles *main-module* to Java source files (`'*.java'`).

`'main-module.classes'`

Compiles *main-module* to Java bytecode files (`'*.class'`).

- `'main-module.ils'`
Compiles *main-module* to IL files (`*.il`) for the .NET Common Language Runtime.
- `'main-module.erls'`
Compiles *main-module* to Erlang source files (`*.erl`).
- `'main-module.beams'`
Compiles *main-module* to Erlang bytecode (object) files (`*.beam`).
- `'libmain-module'`
Builds a library whose top-level module is *main-module*. This will build a static object library, a shared object library (for platforms that support it), and the necessary interface files. For more information, see [\[Libraries\]](#), page [\[Libraries\]](#).
- `'libmain-module.install'`
Builds and installs a library whose top-level module is *main-module*. This target will build and install a static object library and (for platforms that support it) a shared object library, for the default grade and also for the additional grades specified in the `LIBGRADES` variable. It will also build and install the necessary interface files. The variable `INSTALL` specifies the name of the command to use to install each file, by default `'cp'`. The variable `INSTALL_MKDIR` specifies the command to use to create directories, by default `'mkdir -p'`.
- `'main-module.clean'`
Removes the automatically generated files that contain the compiled code of the program and the error messages produced by the compiler. Specifically, this will remove all the `.c`, `.s`, `.o`, `.pic_o`, `.prof`, `.no`, `.ql`, `.used`, `.mih`, and `.err` files belonging to the named *main-module* or its imported modules. Use this target whenever you wish to change compilation model (see [\[Compilation model options\]](#), page [\[Compilation model options\]](#)). This target is also recommended whenever you wish to change the level of inter-module optimization performed (see [\[Overall optimization options\]](#), page [\[Overall optimization options\]](#)) in addition to the mandatory *main-module.depend*.
- `'main-module.realclean'`
Removes all the automatically generated files. In addition to the files removed by *main-module.clean*, this removes the `.int`, `.int0`, `.int2`, `.int3`, `.opt`, `.trans_opt`, `.date`, `.date0`, `.date3`, `.optdate`, `.trans_opt_date`, `.mh` and `.d` files belonging to one of the modules of the program, and also the various possible executables, libraries and dependency files for the program as a whole — *'main-module'*, *'libmain-module.a'*, *'libmain-module.so'*, *'main-module.init'*, *'main-module.dep'* and *'main-module.dv'*.
- `'clean'`
This makes *'main-module.clean'* for every *main-module* for which there is a *'main-module.dep'* file in the current directory, as well as deleting the profiling files *'Prof.CallPair'*, *'Prof.Counts'*, *'Prof.Decl'*, *'Prof.MemWords'* and *'Prof.MemCells'*.

‘realclean’

This makes ‘*main-module.realclean*’ for every *main-module* for which there is a ‘*main-module.dep*’ file in the current directory, as well as deleting the profiling files as per the ‘*clean*’ target.

The variables used by the builtin rules (and their default values) are defined in the file ‘*prefix/lib/mercury/mmake/Mmake.vars*’, however these may be overridden by user ‘Mmake’ files. Some of the more useful variables are:

MAIN_TARGET

The name of the default target to create if ‘*mmake*’ is invoked with any target explicitly named on the command line.

MC

The executable that invokes the Mercury compiler.

GRADEFLAGS and EXTRA_GRADEFLAGS

Compilation model options (see [\[Compilation model options\]](#), page [\[undefined\]](#)) to pass to the Mercury compiler, linker, and other tools (in particular *mmc*, *mgnuc*, *ml*, and *c2init*).

MCFLAGS and EXTRA_MCFLAGS

Options to pass to the Mercury compiler. (Note that compilation model options should be specified in *GRADEFLAGS*, not in *MCFLAGS*.)

MGNUC

The executable that invokes the C compiler.

MGNUCFLAGS and EXTRA_MGNUCFLAGS

Options to pass to the *mgnuc* script.

CFLAGS and EXTRA_CFLAGS

Options to pass to the C compiler.

JAVACFLAGS and EXTRA_JAVACFLAGS

Options to pass to the Java compiler (if you are using it).

ERLANG_FLAGS and EXTRA_ERLANG_FLAGS

Options to pass to the Erlang compiler (if you are using it).

ML

The executable that invokes the linker.

LINKAGE

Can be set to either ‘*shared*’ to link with shared libraries, or ‘*static*’ to always link statically. The default is ‘*shared*’. This variable only has an effect with ‘*mmc --make*’.

MERCURY_LINKAGE

Can be set to either ‘*shared*’ to link with shared Mercury libraries, or ‘*static*’ to always link with the static versions of Mercury libraries. The default is system dependent. This variable only has an effect with ‘*mmc --make*’. See [\[Using installed libraries with mmc -make\]](#), page [\[undefined\]](#).

MLFLAGS and EXTRA_MLFLAGS

Options to pass to the *ml* and *c2init* scripts. (Note that compilation model options should be specified in *GRADEFLAGS*, not in *MLFLAGS*.) These variables have no effect with ‘*mmc --make*’.

LDFLAGS and EXTRA_LDFLAGS

Options to pass to the command used by the `ml` script to link executables (use `ml --print-link-command` to find out what command is used, usually the C compiler).

LD_LIBFLAGS and EXTRA_LD_LIBFLAGS

Options to pass to the command used to by the `ml` script to link shared libraries (use `ml --print-shared-lib-link-command` to find out what command is used, usually the C compiler or the system linker, depending on the platform).

MLLIBS and EXTRA_MLLIBS

A list of ‘-l’ options specifying libraries used by the program (or library) that you are building. See [\[Using libraries with Mmake\]](#), page [\[undefined\]](#). See [\[Using installed libraries with mmc -make\]](#), page [\[undefined\]](#).

MLOBS and EXTRA_MLOBS

A list of extra object files to link into the program (or library) that you are building.

C2INITFLAGS and EXTRA_C2INITFLAGS

Options to pass to the linker and the `c2init` program. `C2INITFLAGS` and `EXTRA_C2INITFLAGS` are obsolete synonyms for `MLFLAGS` and `EXTRA_MLFLAGS` (`ml` and `c2init` take the same set of options). (Note that compilation model options and extra files to be processed by `c2init` should not be specified in `C2INITFLAGS` — they should be specified in `GRADEFLAGS` and `C2INITARGS`, respectively.)

C2INITARGS and EXTRA_C2INITARGS

Extra files to be processed by `c2init`. These variables should not be used for specifying flags to `c2init` (those should be specified in `MLFLAGS`) since they are also used to derive extra dependency information.

EXTRA_LIBRARIES

A list of extra Mercury libraries to link into any programs or libraries that you are building. Libraries should be specified using their base name; that is, without any ‘lib’ prefix or extension. For example the library including the files ‘libfoo.a’ and ‘foo.init’ would be referred to as just ‘foo’. See [\[Using libraries with Mmake\]](#), page [\[undefined\]](#). See [\[Using installed libraries with mmc -make\]](#), page [\[undefined\]](#).

EXTRA_LIB_DIRS

A list of extra Mercury library directory hierarchies to search when looking for extra libraries. See [\[Using libraries with Mmake\]](#), page [\[undefined\]](#). See [\[Using installed libraries with mmc -make\]](#), page [\[undefined\]](#).

INSTALL_PREFIX

The path to the root of the directory hierarchy where the libraries, etc. you are building should be installed. The default is to install in the same location as the Mercury compiler being used to do the install.

INSTALL The command used to install each file in a library. The command should take a list of files to install and the location to install them. The default command is ‘`cp`’.

INSTALL_MKDIR

The command used to create each directory in the directory hierarchy where the libraries are to be installed. The default command is ‘`mkdir -p`’.

LIBGRADES

A list of additional grades which should be built when installing libraries. The default is to install the Mercury compiler’s default set of grades. Note that this may not be the set of grades in which the standard libraries were actually installed. Note also that any **GRADEFLAGS** settings will also be applied when the library is built in each of the listed grades, so you may not get what you expect if those options are not subsumed by each of the grades listed.

LIB_LINKAGES

A list of linkage styles (‘**shared**’ or ‘**static**’) for which libraries should be built and installed. The default is to install libraries for both static and shared linking. This variable only has an effect with ‘`mmc --make`’.

Other variables also exist — see ‘`prefix/lib/mercury/mmake/Mmake.vars`’ for a complete list.

If you wish to temporarily change the flags passed to an executable, rather than setting the various ‘**FLAGS**’ variables directly, you can set an ‘**EXTRA_**’ variable. This is particularly intended for use where a shell script needs to call mmake and add an extra parameter, without interfering with the flag settings in the ‘`Mmakefile`’.

For each of the variables for which there is version with an ‘**EXTRA_**’ prefix, there is also a version with an ‘**ALL_**’ prefix that is defined to include both the ordinary and the ‘**EXTRA_**’ version. If you wish to *use* the values any of these variables in your Mmakefile (as opposed to *setting* the values), then you should use the ‘**ALL_**’ version.

It is also possible to override these variables on a per-file basis. For example, if you have a module called say ‘`bad_style.m`’ which triggers lots of compiler warnings, and you want to disable the warnings just for that file, but keep them for all the other modules, then you can override **MCFLAGS** just for that file. This is done by setting the variable ‘`MCFLAGS-bad_style`’, as shown here:

```
MCFLAGS-bad_style = --inhibit-warnings
```

Mmake has a few options, including ‘`--use-subdirs`’, ‘`--use-mmc-make`’, ‘`--save-makefile`’, ‘`--verbose`’, and ‘`--no-warn-undefined-vars`’. For details about these options, see the man page or type ‘`mmake --help`’.

Finally, since Mmake is built on top of Make or GNU Make, you can also make use of the features and options supported by the underlying Make. In particular, GNU Make has support for running jobs in parallel, which is very useful if you have a machine with more than one CPU.

As an alternative to Mmake, the Mercury compiler now contains a significant part of the functionality of Mmake, using ‘`mmc`’s ‘`--make`’ option.

The advantages of the ‘`mmc --make`’ over Mmake are that there is no ‘`mmake depend`’ step and the dependencies are more accurate.

Note that ‘`--use-subdirs`’ is automatically enabled if you specify ‘`mmc --make`’.

The Mmake variables above can be used by ‘`mmc --make`’ if they are set in a file called ‘`Mercury.options`’. The ‘`Mercury.options`’ file has the same syntax as an Mmakefile, but only variable assignments and ‘`include`’ directives are allowed. All variables in ‘`Mercury.options`’ are treated as if they are assigned using ‘`:=`’. Variables may also be set in the environment, overriding settings in options files.

‘`mmc --make`’ can be used in conjunction with Mmake. This is useful for projects which include source code written in languages other than Mercury. The ‘`--use-mmc-make`’ Mmake option disables Mmake’s Mercury-specific rules. Mmake will then process source files written in other languages, but all Mercury compilation will be done by ‘`mmc --make`’. The following variables can be set in the Mmakefile to control the use of ‘`mmc --make`’.

MERCURY_MAIN_MODULES

The top-level modules of the programs or libraries being built in the directory. This must be set to tell Mmake to use ‘`mmc --make`’ to rebuild the targets for the main modules even if those files already exist.

MC_BUILD_FILES

Other files which should be built with ‘`mmc --make`’. This should only be necessary for header files generated by the Mercury compiler which are included by the user’s C source files.

MC_MAKE_FLAGS and EXTRA_MC_MAKE_FLAGS

Options to pass to the Mercury compiler only when using ‘`mmc --make`’.

6 Libraries

Often you will want to use a particular set of Mercury modules in more than one program. The Mercury implementation includes support for developing libraries, i.e. sets of Mercury modules intended for reuse. It allows separate compilation of libraries and, on many platforms, it supports shared object libraries.

6.1 Writing libraries

A Mercury library is identified by a top-level module, which should contain all of the modules in that library as sub-modules. It may be as simple as this ‘`mypackage.m`’ file:

```
:- module mypackage.
:- interface.
:- include_module foo, bar, baz.
```

This defines a module ‘`mypackage`’ containing sub-modules ‘`mypackage:foo`’, ‘`mypackage:bar`’, and ‘`mypackage:baz`’.

It is also possible to build libraries of unrelated modules, so long as the top-level module imports all the necessary modules. For example:

```
:- module blah.
:- import_module fee, fie, foe, fum.
```

This example defines a module ‘`blah`’, which has no functionality of its own, and which is just used for grouping the unrelated modules ‘`fee`’, ‘`fie`’, ‘`foe`’, and ‘`fum`’.

Generally it is better style for each library to consist of a single module which encapsulates its sub-modules, as in the first example, rather than just a group of unrelated modules, as in the second example.

6.2 Building with `mmc -make`

6.2.1 Building and installing libraries with `mmc -make`

To build a library from the source `'mypackage.m'` (and other included modules), run `'mmc'` with the following arguments:

```
mmc --make libmypackage
```

`'mmc'` will create static (non-shared) object libraries and, on most platforms, shared object libraries; however, we do not yet support the creation of dynamic link libraries (DLLs) on Windows. Use the `'mmc'` option `'--lib-linkage'` to specify which versions of the library should be created: `'shared'` or `'static'`. The `'--lib-linkage'` option can be specified multiple times. In our example, the files `'libmypackage.a'` and `'libmypackage.so'` should appear in the current directory.

Other programs can more easily use a library that is installed. To install the library, issue the following command:

```
mmc --make --install-prefix <dir> libmypackage.install
```

`'mmc'` will create the directory `'<dir>/lib/mercury'` and install the library there. The library will be compiled in all valid grades and with all interface files. Because several grades are usually compiled, installing the library can be a lengthy process. You can specify the set of installed grades using the option `'--no-libgrade'` followed by `'--libgrade <grade>'` for all grades you wish to install.

If no `'--install-prefix <dir>'` is specified, the library will be installed in the standard location, next to the Mercury standard library.

6.2.2 Using installed libraries with `mmc -make`

Once a library is installed, it can be used by running `'mmc'` with the following options:

```
mmc ... --ml mypackage ... --ml myotherlib ... --ml my_yet_another_lib ...
```

If a library was installed in a different place (using `'--install-prefix <dir>'`), you will also need to add this option:

```
mmc ... --mld <dir>/lib/mercury ...
```

Note that `'/lib/mercury'` has to be added to the searched path. The `'--mld'` option can be used several times to add more directories to the library search path.

You can also specify whether to link executables with the shared or static versions of Mercury libraries using `'--mercury-linkage shared'` or `'--mercury-linkage static'`.

6.2.3 Using non-installed libraries with `mmc -make`

Suppose the user wants to link against library `'mypackage'` without installing the library. The source of the library is stored in the directory `'<dir>'` and that the library has been

properly built using ‘`mmc --make libmypackage`’. To link against the library, the following options have to be added to ‘`mmc`’:

```
mmc ... --search-lib-files-dir <dir> \\  
      --init-file <dir>/mypackage.init \\  
      --link-object <dir>/libmypackage.a \\  
...
```

Note that the option ‘`--ml`’ is not used.

You need to make sure the library ‘`libmypackage.a`’ and the main program were compiled in the same grade.

If you need to experiment with more grades, be sure to build the library in all the grades (building several times using ‘`mmc --grade <grade> --make libmypackage`’) and use the ‘`libmypackage.a`’ that is compatible with your main program’s grade:

```
mmc ... --use-grade-subdirs \\  
      --grade <grade> \\  
      --search-lib-files-dir <dir> \\  
      --init-file <dir>/mypackage.init \\  
      --link-object <dir>/Mercury/<grade>*/Mercury/lib/libmypackage.a \\  
...
```

6.3 Building with Mmake

6.3.1 Building libraries with Mmake

Generally Mmake will do most of the work of building libraries automatically. Here’s a sample Mmakefile for creating a library.

```
MAIN_TARGET = libmypackage  
depend: mypackage.depend
```

The Mmake target ‘`libfoo`’ is a built-in target for creating a library whose top-level module is ‘`foo.m`’. The automatically generated Mmake rules for the target ‘`libfoo`’ will create all the files needed to use the library. (You will need to run ‘`mmake foo.depend`’ first to generate the module dependency information.)

Mmake will create static (non-shared) object libraries and, on most platforms, shared object libraries; however, we do not yet support the creation of dynamic link libraries (DLLs) on Windows. Static libraries are created using the standard tools ‘`ar`’ and ‘`ranlib`’. Shared libraries are created using the ‘`--make-shared-lib`’ option to ‘`ml`’. The automatically-generated Make rules for ‘`libmypackage`’ will look something like this:

```
libmypackage: libmypackage.a libmypackage.so \  
$(mypackage.ints) $(mypackage.int3s) \  
$(mypackage.opts) $(mypackage.trans_opts) mypackage.init  
  
libmypackage.a: $(mypackage.os)  
rm -f libmypackage.a  
$(AR) $(ARFLAGS) libmypackage.a $(mypackage.os) $(MLOBS)  
$(RANLIB) $(RANLIBFLAGS) mypackage.a
```

```

libmypackage.so: $(mypackage.pic_os)
$(ML) $(MLFLAGS) --make-shared-lib -o libmypackage.so \
$(mypackage.pic_os) $(MLPICOBJS) $(MLLIBS)

libmypackage.init:
...

clean:
rm -f libmypackage.a libmypackage.so

```

If necessary, you can override the default definitions of the variables such as ‘ML’, ‘MLFLAGS’, ‘MLPICOBJS’, and ‘MLLIBS’ to customize the way shared libraries are built. Similarly ‘AR’, ‘ARFLAGS’, ‘MLOBS’, ‘RANLIB’, and ‘RANLIBFLAGS’ control the way static libraries are built. (The ‘MLOBS’ variable is supposed to contain a list of additional object files to link into the library, while the ‘MLLIBS’ variable should contain a list of ‘-l’ options naming other libraries used by this library. ‘MLPICOBJS’ is described below.)

Note that to use a library, as well as the shared or static object library, you also need the interface files. That’s why the ‘libmypackage’ target builds ‘\$(mypackage.ints)’ and ‘\$(mypackage.int3s)’. If the people using the library are going to use intermodule optimization, you will also need the intermodule optimization interfaces. The ‘libmypackage’ target will build ‘\$(mypackage.opts)’ if ‘--intermodule-optimization’ is specified in your ‘MCFLAGS’ variable (this is recommended). Similarly, if the people using the library are going to use transitive intermodule optimization, you will also need the transitive intermodule optimization interfaces (‘\$(mypackage.trans_opt)’). These will be built if ‘--trans-intermod-opt’ is specified in your ‘MCFLAGS’ variable.

In addition, with certain compilation grades, programs will need to execute some startup code to initialize the library; the ‘mypackage.init’ file contains information about initialization code for the library. The ‘libmypackage’ target will build this file.

On some platforms, shared objects must be created using position independent code (PIC), which requires passing some special options to the C compiler. On these platforms, Mmake will create ‘.pic.o’ files, and ‘\$(mypackage.pic_os)’ will contain a list of the ‘.pic.o’ files for the library whose top-level module is ‘mypackage’. In addition, ‘\$(MLPICOBJS)’ will be set to ‘\$MLOBS’ with all occurrences of ‘.o’ replaced with ‘.pic.o’. On other platforms, position independent code is the default, so ‘\$(mypackage.pic_os)’ will just be the same as ‘\$(mypackage.os)’, which contains a list of the ‘.o’ files for that module, and ‘\$(MLPICOBJS)’ will be the same as ‘\$(MLOBS)’.

6.3.2 Installing libraries with Mmake

‘mmake’ has support for alternative library directory hierarchies. These have the same structure as the ‘prefix/lib/mercury’ tree, including the different subdirectories for different grades and different machine architectures.

In order to support the installation of a library into such a tree, you simply need to specify (e.g. in your ‘Mmakefile’) the path prefix and the list of grades to install:

```

INSTALL_PREFIX = /my/install/dir
LIBGRADES = asm_fast asm_fast.gc.tr.debug

```

This specifies that libraries should be installed in `‘/my/install/dir/lib/mercury’`, in the default grade plus `‘asm_fast’` and `‘asm_fast.gc.tr.debug’`. If `‘INSTALL_PREFIX’` is not specified, `‘mmake’` will attempt to install the library in the same place as the standard Mercury libraries. If `‘LIBGRADES’` is not specified, `‘mmake’` will use the Mercury compiler’s default set of grades, which may or may not correspond to the actual set of grades in which the standard Mercury libraries were installed.

To actually install a library `‘libfoo’`, use the `‘mmake’` target `‘libfoo.install’`. This also installs all the needed interface files, and (if intermodule optimisation is enabled) the relevant intermodule optimisation files.

One can override the list of grades to install for a given library `‘libfoo’` by setting the `‘LIBGRADES-foo’` variable, or add to it by setting `‘EXTRA_LIBGRADES-foo’`.

The command used to install each file is specified by `‘INSTALL’`. If `‘INSTALL’` is not specified, `‘cp’` will be used.

The command used to create directories is specified by `‘INSTALL_MKDIR’`. If `‘INSTALL_MKDIR’` is not specified, `‘mkdir -p’` will be used.

Note that currently it is not possible to set the installation prefix on a library-by-library basis.

6.3.3 Using libraries with Mmake

Once a library is installed, using it is easy. Suppose the user wishes to use the library `‘mypackage’` (installed in the tree rooted at `‘/some/directory/mypackage’`) and the library `‘myotherlib’` (installed in the tree rooted at `‘/some/directory/myotherlib’`). The user need only set the following Mmake variables:

```
EXTRA_LIB_DIRS = /some/directory/mypackage/lib/mercury \
                 /some/directory/myotherlib/lib/mercury
EXTRA_LIBRARIES = mypackage myotherlib
```

When using `‘--intermodule-optimization’` with a library which uses the C interface, it may be necessary to add `‘-I’` options to `‘MGNUCLFLAGS’` so that the C compiler can find any header files used by the library’s C code.

Mmake will ensure that the appropriate directories are searched for the relevant interface files, module initialisation files, compiled libraries, etc.

Beware that the directory name that you must use in `‘EXTRA_LIB_DIRS’` or as the argument of the `‘--mld’` option is not quite the same as the name that was specified in the `‘INSTALL_PREFIX’` when the library was installed — the name needs to have `‘/lib/mercury’` appended.

One can specify extra libraries to be used on a program-by-program basis. For instance, if the program `‘foo’` also uses the library `‘mylib4foo’`, but the other programs governed by the Mmakefile don’t, then one can declare:

```
EXTRA_LIBRARIES-foo = mylib4foo
```

6.4 Libraries and the Java grade

Libraries are handled a little differently for the Java grade. Instead of compiling object code into a static or shared library, the class files are added to a jar (Java ARchive) file of the form `‘library-name.jar’`.

To create or install a Java library, simply specify that you want to use the java grade, either by setting `GRADE=java` in your Mmakefile, or by including `--java` or `--grade java` in your `GRADEFLAGS`, then follow the instructions as above.

Java libraries are installed to the directory `prefix/lib/mercury/lib/java`. To include them in a program, in addition to the instructions above, you will need to include the installed jar file in your `CLASSPATH`, which you can set using `--java-classpath jarfile` in `MCFLAGS`.

6.5 Libraries and the Erlang grade

Since the Erlang implementation does not have library files, the Mercury compiler puts all the `.beam` files for a single Mercury library into a directory named `liblibrary-name.beams`.

To create or install an Erlang "library", specify that you want to use the erlang grade and use `mmc --make`. Mmake does not, at present, support Erlang targets.

7 Debugging

7.1 Quick overview

This section gives a quick and simple guide to getting started with the debugger. The remainder of this chapter contains more detailed documentation.

To use the debugger, you must first compile your program with debugging enabled. You can do this by using one of the `--debug` or `--decl-debug` options when invoking `mmc`, or by including `GRADEFLAGS = --debug` or `GRADEFLAGS = --decl-debug` in your `Mmakefile`.

```
bash$ mmc --debug hello.m
```

Once you've compiled with debugging enabled, you can use the `mdb` command to invoke your program under the debugger:

```
bash$ mdb ./hello arg1 arg2 ...
```

Any arguments (such as `arg1 arg2 ...` in this example) that you pass after the program name will be given as arguments to the program.

The debugger will print a start-up message and will then show you the first trace event, namely the call to `main/2`:

```
1:      1  1 CALL pred hello:main/2-0 (det)
             hello.m:13
```

```
mdb>
```

By hitting enter at the `mdb>` prompt, you can step through the execution of your program to the next trace event:

```
2:      2  2 CALL pred io:write_string/3-0 (det)
             io.m:2837 (hello.m:14)
```

```
mdb>
```

```
Hello, world
```

```

3:      2  2 EXIT pred io:write_string/3-0 (det)
          io.m:2837 (hello.m:14)

mdb>

```

For each trace event, the debugger prints out several pieces of information. The three numbers at the start of the display are the event number, the call sequence number, and the call depth. (You don't really need to pay too much attention to those.) They are followed by the event type (e.g. 'CALL' or 'EXIT'). After that comes the identification of the procedure in which the event occurred, consisting of the module-qualified name of the predicate or function to which the procedure belongs, followed by its arity, mode number and determinism. This may sometimes be followed by a "path" (see [\[Tracing of Mercury programs\]](#), page [\[undefined\]](#)). At the end is the file name and line number of the called procedure and (if available) also the file name and line number of the call.

The most useful `mdb` commands have single-letter abbreviations. The 'alias' command will show these abbreviations:

```

mdb> alias
?      =>  help
EMPTY  =>  step
NUMBER =>  step
P      =>  print *
b      =>  break
c      =>  continue
d      =>  stack
f      =>  finish
g      =>  goto
h      =>  help
p      =>  print
r      =>  retry
s      =>  step
v      =>  vars

```

The 'P' or 'print *' command will display the values of any live variables in scope. The 'f' or 'finish' command can be used if you want to skip over a call. The 'b' or 'break' command can be used to set break-points. The 'd' or 'stack' command will display the call stack. The 'quit' command will exit the debugger.

That should be enough to get you started. But if you have GNU Emacs installed, you should strongly consider using the Emacs interface to 'mdb' — see the following section.

For more information about the available commands, use the '?' or 'help' command, or see [\[Debugger commands\]](#), page [\[undefined\]](#).

7.2 GNU Emacs interface

As well as the command-line debugger, `mdb`, there is also an Emacs interface to this debugger. Note that the Emacs interface only works with GNU Emacs, not with XEmacs.

With the Emacs interface, the debugger will display your source code as you trace through it, marking the line that is currently being executed, and allowing you to easily set breakpoints on particular lines in your source code. You can have separate windows for the

debugger prompt, the source code being executed, and for the output of the program being executed. In addition, most of the mdb commands are accessible via menus.

To start the Emacs interface, you first need to put the following text in the file `‘.emacs’` in your home directory, replacing `“/usr/local/mercury-1.0”` with the directory that your Mercury implementation was installed in.

```
(setq load-path (cons (expand-file-name
  "/usr/local/mercury-1.0/lib/mercury/elisp")
  load-path))
(autoload 'mdb "gud" "Invoke the Mercury debugger" t)
```

Build your program with debugging enabled, as described in [\[Quick overview\]](#), [page \[undefined\]](#) or [\[Preparing a program for debugging\]](#), [page \[undefined\]](#). Then start up Emacs, e.g. using the command `‘emacs’`, and type `M-x mdb` `(RET)`. Emacs will then prompt you for the mdb command to invoke

```
Run mdb (like this): mdb
```

and you should type in the name of the program that you want to debug and any arguments that you want to pass to it:

```
Run mdb (like this): mdb ./hello arg1 arg2 ...
```

Emacs will then create several “buffers”: one for the debugger prompt, one for the input and output of the program being executed, and one or more for the source files. By default, Emacs will split the display into two parts, called “windows”, so that two of these buffers will be visible. You can use the command `C-x o` to switch between windows, and you can use the command `C-x 2` to split a window into two windows. You can use the “Buffers” menu to select which buffer is displayed in each window.

If you’re using X-Windows, then it is a good idea to set the Emacs variable `‘pop-up-frames’` to `‘t’` before starting mdb, since this will cause each buffer to be displayed in a new “frame” (i.e. a new X window). You can set this variable interactively using the `‘set-variable’` command, i.e. `M-x set-variable` `(RET)` `pop-up-frames` `(RET)` `t` `(RET)`. Or you can put `‘(setq pop-up-frames t)’` in the `‘.emacs’` file in your home directory.

For more information on buffers, windows, and frames, see the Emacs documentation.

Another useful Emacs variable is `‘gud-mdb-directories’`. This specifies the list of directories to search for source files. You can use a command such as

```
M-x set-variable (RET)
gud-mdb-directories (RET)
(list "/foo/bar" "../other" "/home/guest") (RET)
```

to set it interactively, or you can put a command like

```
(setq gud-mdb-directories
  (list "/foo/bar" "../other" "/home/guest"))
```

in your `‘.emacs’` file.

At each trace event, the debugger will search for the source file corresponding to that event, first in the same directory as the program, and then in the directories specified by the `‘gud-mdb-directories’` variable. It will display the source file, with the line number corresponding to that trace event marked by an arrow (`‘=>’`) at the start of the line.

Several of the debugger features can be accessed by moving the cursor to the relevant part of the source code and then selecting a command from the menu. You can set a break

point on a line by moving the cursor to the appropriate line in your source code (e.g. with the arrow keys, or by clicking the mouse there), and then selecting the “Set breakpoint on line” command from the “Breakpoints” sub-menu of the “MDB” menu. You can set a breakpoint on a procedure by moving the cursor over the procedure name and then selecting the “Set breakpoint on procedure” command from the same menu. And you can display the value of a variable by moving the cursor over the variable name and then selecting the “Print variable” command from the “Data browsing” sub-menu of the “MDB” menu. Most of the menu commands also have keyboard short-cuts, which are displayed on the menu.

Note that mdb’s ‘`context`’ and ‘`user_event_context`’ commands should not be used if you are using the Emacs interface, otherwise the Emacs interface won’t be able to parse the file names and line numbers that mdb outputs, and so it won’t be able to highlight the correct location in the source code.

7.3 Tracing of Mercury programs

The Mercury debugger is based on a modified version of the box model on which the four-port debuggers of most Prolog systems are based. Such debuggers abstract the execution of a program into a sequence, also called a *trace*, of execution events of various kinds. The four kinds of events supported by most Prolog systems (their *ports*) are

<i>call</i>	A call event occurs just after a procedure has been called, and control has just reached the start of the body of the procedure.
<i>exit</i>	An exit event occurs when a procedure call has succeeded, and control is about to return to its caller.
<i>redo</i>	A redo event occurs when all computations to the right of a procedure call have failed, and control is about to return to this call to try to find alternative solutions.
<i>fail</i>	A fail event occurs when a procedure call has run out of alternatives, and control is about to return to the rightmost computation to its left that still has possibly successful alternatives left.

Mercury also supports these four kinds of events, but not all events can occur for every procedure call. Which events can occur for a procedure call, and in what order, depend on the determinism of the procedure. The possible event sequences for procedures of the various determinisms are as follows.

<i>nondet procedures</i>	a call event, zero or more repeats of (exit event, redo event), and a fail event
<i>multi procedures</i>	a call event, one or more repeats of (exit event, redo event), and a fail event
<i>semidet and cc_nondet procedures</i>	a call event, and either an exit event or a fail event
<i>det and cc_multi procedures</i>	a call event and an exit event
<i>failure procedures</i>	a call event and a fail event

erroneous procedures

a call event

In addition to these four event types, Mercury supports *exception* events. An exception event occurs when an exception has been thrown inside a procedure, and control is about to propagate this exception to the caller. An exception event can replace the final exit or fail event in the event sequences above or, in the case of erroneous procedures, can come after the call event.

Besides the event types call, exit, redo, fail and exception, which describe the *interface* of a call, Mercury also supports several types of events that report on what is happening *internal* to a call. Each of these internal event types has an associated parameter called a path. The internal event types are:

<i>cond</i>	A cond event occurs when execution reaches the start of the condition of an if-then-else. The path associated with the event specifies which if-then-else this is.
<i>then</i>	A then event occurs when execution reaches the start of the then part of an if-then-else. The path associated with the event specifies which if-then-else this is.
<i>else</i>	An else event occurs when execution reaches the start of the else part of an if-then-else. The path associated with the event specifies which if-then-else this is.
<i>disj</i>	A disj event occurs when execution reaches the start of a disjunct in a disjunction. The path associated with the event specifies which disjunct of which disjunction this is.
<i>switch</i>	A switch event occurs when execution reaches the start of one arm of a switch (a disjunction in which each disjunct unifies a bound variable with different function symbol). The path associated with the event specifies which arm of which switch this is.
<i>neg_enter</i>	A neg_enter event occurs when execution reaches the start of a negated goal. The path associated with the event specifies which negation goal this is.
<i>neg_fail</i>	A neg_fail event occurs when a goal inside a negation succeeds, which means that its negation fails. The path associated with the event specifies which negation goal this is.
<i>neg_success</i>	A neg_success event occurs when a goal inside a negation fails, which means that its negation succeeds. The path associated with the event specifies which negation goal this is.

A goal path is a sequence of path components separated by semicolons. Each path component is one of the following:

<i>cnum</i>	The <i>num</i> 'th conjunct of a conjunction.
<i>dnum</i>	The <i>num</i> 'th disjunct of a disjunction.
<i>snum</i>	The <i>num</i> 'th arm of a switch.

<code>?</code>	The condition of an if-then-else.
<code>t</code>	The then part of an if-then-else.
<code>e</code>	The else part of an if-then-else.
<code>~</code>	The goal inside a negation.
<code>q!</code>	The goal inside an existential quantification or other scope that changes the determinism of the goal.
<code>q</code>	The goal inside an existential quantification or other scope that doesn't change the determinism of the goal.

A goal path describes the position of a goal inside the body of a procedure definition. For example, if the procedure body is a disjunction in which each disjunct is a conjunction, then the goal path `'d2;c3;'` denotes the third conjunct within the second disjunct. If the third conjunct within the second disjunct is an atomic goal such as a call or a unification, then this will be the only goal with whose path has `'d2;c3;'` as a prefix. If it is a compound goal, then its components will all have paths that have `'d2;c3;'` as a prefix, e.g. if it is an if-then-else, then its three components will have the paths `'d2;c3;?;'`, `'d2;c3;t;'` and `'d2;c3;e;'`.

Goal paths refer to the internal form of the procedure definition. When debugging is enabled (and the option `'--trace-optimized'` is not given), the compiler will try to keep this form as close as possible to the source form of the procedure, in order to make event paths as useful as possible to the programmer. Due to the compiler's flattening of terms, and its introduction of extra unifications to implement calls in implied modes, the number of conjuncts in a conjunction will frequently differ between the source and internal form of a procedure. This is rarely a problem, however, as long as you know about it. Mode reordering can be a bit more of a problem, but it can be avoided by writing single-mode predicates and functions so that producers come before consumers. The compiler transformation that potentially causes the most trouble in the interpretation of goal paths is the conversion of disjunctions into switches. In most cases, a disjunction is transformed into a single switch, and it is usually easy to guess, just from the events within a switch arm, just which disjunct the switch arm corresponds to. Some cases are more complex; for example, it is possible for a single disjunction can be transformed into several switches, possibly with other, smaller disjunctions inside them. In such cases, making sense of goal paths may require a look at the internal form of the procedure. You can ask the compiler to generate a file with the internal forms of the procedures in a given module by including the options `'-dfinal -Dpaths'` on the command line when compiling that module.

7.4 Preparing a program for debugging

When you compile a Mercury program, you can specify whether you want to be able to run the Mercury debugger on the program or not. If you do, the compiler embeds calls to the Mercury debugging system into the executable code of the program, at the execution points that represent trace events. At each event, the debugging system decides whether to give control back to the executable immediately, or whether to first give control to you, allowing you to examine the state of the computation and issue commands.

Mercury supports two broad ways of preparing a program for debugging. The simpler way is to compile a program in a debugging grade, which you can do directly by specifying a grade that includes the word “debug” or “decldebug” (e.g. ‘asm_fast.gc.debug’, or ‘asm_fast.gc.decldebug’), or indirectly by specifying one of the ‘--debug’ or ‘--decl-debug’ grade options to the compiler, linker, and other tools (in particular `mmc`, `mgnuc`, `ml`, and `c2init`). If you follow this way, and accept the default settings of the various compiler options that control the selection of trace events (which are described below), you will be assured of being able to get control at every execution point that represents a potential trace event, which is very convenient.

The “decldebug” grades improve declarative debugging by allowing the user to track the source of subterms (see [\[Improving the search\]](#), page [\[Improving the search\]](#)). Doing this increases the size of executables, so these grades should only be used when you need the subterm dependency tracking feature of the declarative debugger. Note that declarative debugging, with the exception of the subterm dependency tracking features, also works in the .debug grades.

The two drawbacks of using a debugging grade are the large size of the resulting executables, and the fact that often you discover that you need to debug a big program only after having built it in a non-debugging grade. This is why Mercury also supports another way to prepare a program for debugging, one that does not require the use of a debugging grade. With this way, you can decide, individually for each module, which of four trace levels, ‘none’, ‘shallow’, ‘deep’, and ‘rep’ you want to compile them with:

- ‘none’ A procedure compiled with trace level ‘none’ will never generate any events.
- ‘deep’ A procedure compiled with trace level ‘deep’ will always generate all the events requested by the user. By default, this is all possible events, but you can tell the compiler that you are not interested in some kinds of events via compiler options (see below). However, declarative debugging requires all events to be generated if it is to operate properly, so do not disable the generation of any event types if you want to use declarative debugging. For more details see [\[Declarative debugging\]](#), page [\[Declarative debugging\]](#).
- ‘rep’ This trace level is the same as trace level ‘deep’, except that a representation of the module is stored in the executable along with the usual debugging information. The declarative debugger can use this extra information to help it avoid asking unnecessary questions, so this trace level has the effect of better declarative debugging at the cost of increased executable size. For more details see [\[Declarative debugging\]](#), page [\[Declarative debugging\]](#).
- ‘shallow’ A procedure compiled with trace level ‘shallow’ will generate interface events if it is called from a procedure compiled with trace level ‘deep’, but it will never generate any internal events, and it will not generate any interface events either if it is called from a procedure compiled with trace level ‘shallow’. If it is called from a procedure compiled with trace level ‘none’, the way it will behave is dictated by whether its nearest ancestor whose trace level is not ‘none’ has trace level ‘deep’ or ‘shallow’.

The intended uses of these trace levels are as follows.

<code>'deep'</code>	You should compile a module with trace level <code>'deep'</code> if you suspect there may be a bug in the module, or if you think that being able to examine what happens inside that module can help you locate a bug.
<code>'rep'</code>	You should compile a module with trace level <code>'rep'</code> if you suspect there may be a bug in the module, you wish to use the full power of the declarative debugger, and you are not concerned about the size of the executable.
<code>'shallow'</code>	You should compile a module with trace level <code>'shallow'</code> if you believe the code of the module is reliable and unlikely to have bugs, but you still want to be able to get control at calls to and returns from any predicates and functions defined in the module, and if you want to be able to see the arguments of those calls.
<code>'none'</code>	You should compile a module with trace level <code>'none'</code> only if you are reasonably confident that the module is reliable, and if you believe that knowing what calls other modules make to this module would not significantly benefit you in your debugging.

In general, it is a good idea for most or all modules that can be called from modules compiled with trace level `'deep'` or `'rep'` to be compiled with at least trace level `'shallow'`.

You can control what trace level a module is compiled with by giving one of the following compiler options:

<code>'--trace shallow'</code>	This always sets the trace level to <code>'shallow'</code> .
<code>'--trace deep'</code>	This always sets the trace level to <code>'deep'</code> .
<code>'--trace rep'</code>	This always sets the trace level to <code>'rep'</code> .
<code>'--trace minimum'</code>	In debugging grades, this sets the trace level to <code>'shallow'</code> ; in non-debugging grades, it sets the trace level to <code>'none'</code> .
<code>'--trace default'</code>	In debugging grades, this sets the trace level to <code>'deep'</code> ; in non-debugging grades, it sets the trace level to <code>'none'</code> .

As the name implies, the last alternative is the default, which is why by default you get no debugging capability in non-debugging grades and full debugging capability in debugging grades. The table also shows that in a debugging grade, no module can be compiled with trace level `'none'`.

Important note: If you are not using a debugging grade, but you compile some modules with a trace level other than none, then you must also pass the `'--trace'` (or `'-t'`) option to `c2init` and to the Mercury linker. If you're using Mmake, then you can do this by including `'--trace'` in the `'MLFLAGS'` variable.

If you're using Mmake, then you can also set the compilation options for a single module named *Module* by setting the Mmake variable `'MCFLAGS-Module'`. For example, to compile the file `'foo.m'` with deep tracing, `'bar.m'` with shallow tracing, and everything else with no tracing, you could use the following:

```

MLFLAGS      = --trace
MCFLAGS-foo  = --trace deep
MCFLAGS-bar  = --trace shallow

```

7.5 Tracing optimized code

By default, all trace levels other than ‘none’ turn off all compiler optimizations that can affect the sequence of trace events generated by the program, such as inlining. If you are specifically interested in how the compiler’s optimizations affect the trace event sequence, you can specify the option ‘--trace-optimized’, which tells the compiler that it does not have to disable those optimizations. (A small number of low-level optimizations have not yet been enhanced to work properly in the presence of tracing, so compiler disables these even if ‘--trace-optimized’ is given.)

7.6 Mercury debugger invocation

The executables of Mercury programs by default do not invoke the Mercury debugger even if some or all of their modules were compiled with some form of tracing, and even if the grade of the executable is a debugging grade. This is similar to the behaviour of executables created by the implementations of other languages; for example the executable of a C program compiled with ‘-g’ does not automatically invoke gdb or dbx etc when it is executed.

Unlike those other language implementations, when you invoke the Mercury debugger ‘mdb’, you invoke it not just with the name of an executable but with the command line you want to debug. If something goes wrong when you execute the command

```
prog arg1 arg2 ...
```

and you want to find the cause of the problem, you must execute the command

```
mdb [mdb-options] prog arg1 arg2 ...
```

because you do not get a chance to specify the command line of the program later.

When the debugger starts up, as part of its initialization it executes commands from the following three sources, in order:

1. The file named by the ‘MERCURY_DEBUGGER_INIT’ environment variable. Usually, ‘mdb’ sets this variable to point to a file that provides documentation for all the debugger commands and defines a small set of aliases. However, if ‘MERCURY_DEBUGGER_INIT’ is already defined when ‘mdb’ is invoked, it will leave its value unchanged. You can use this override ability to provide alternate documentation. If the file named by ‘MERCURY_DEBUGGER_INIT’ cannot be read, ‘mdb’ will print a warning, since in that case, that usual online documentation will not be available.
2. The file named ‘.mdbrc’ in your home directory. You can put your usual aliases and settings here.
3. The file named ‘.mdbrc’ in the current working directory. You can put program-specific aliases and settings here.

mdb will ignore any lines starting with the character # in any of the above mentioned files.

mdb accepts the following options from the command line. The options should be given to mdb before the name of the executable to be debugged.

-t *file-name*, **--tty** *file-name*

Redirect all of the I/O for the debugger to the device specified by *file-name*. The I/O for the program being debugged will not be redirected. This option allows the contents of a file to be piped to the program being debugged and not to mdb. For example on Linux the command ‘`mdb -t /dev/tty ./myprog < myinput`’ will cause the contents of myinput to be piped to the program myprog, but mdb will read its input from the terminal.

-w, **--window**, **--mdb-in-window**

Run mdb in a new window, with mdb’s I/O going to that window, but with the program’s I/O going to the current terminal. Note that this will not work on all systems.

--program-in-window

Run the program in a new window, with the program’s I/O going to that window, but with mdb’s I/O going to the current terminal. Note that input and output redirection will not work with the ‘**--program-in-window**’ option. ‘**--program-in-window**’ will work on most UNIX systems running the X Window System, even those for which ‘**--mdb-in-window**’ is not supported.

-c *window-command*, **--window-command** *window-command*

Specify the command used by the ‘**--program-in-window**’ option for executing a command in a new window. The default such command is ‘`xterm -e`’.

7.7 Mercury debugger concepts

The operation of the Mercury debugger ‘mdb’ is based on the following concepts.

break points

The user may associate a break point with some events that occur inside a procedure; the invocation condition of the break point says which events these are. The four possible invocation conditions (also called scopes) are:

- the call event,
- all interface events,
- all events, and
- the event at a specific point in the procedure.

The effect of a break point depends on the state of the break point.

- If the state of the break point is ‘**stop**’, execution will stop and user interaction will start at any event within the procedure that matches the invocation conditions, unless the current debugger command has specifically disabled this behaviour (see the concept ‘**strict commands**’ below).

- If the state of the break point is `'print'`, the debugger will print any event within the procedure that matches the invocation conditions, unless the current debugger command has specifically disabled this behaviour (see the concept `'print level'` below).

Neither of these will happen if the break point is disabled.

Every break point has a print list. Every time execution stops at an event that matches the breakpoint, mdb implicitly executes a print command for each element in the breakpoint's print list. A print list element can be the word `'goal'`, which causes the goal to be printed as if by `'print goal'`; it can be the word `'*'`, which causes all the variables to be printed as if by `'print *'`; or it can be the name or number of a variable, possibly followed (without white space) by term path, which causes the specified variable or part thereof to be printed as if the element were given as an argument to the `'print'` command.

strict commands

When a debugger command steps over some events without user interaction at those events, the *strictness* of the command controls whether the debugger will stop execution and resume user interaction at events to which a break point with state `'stop'` applies. By default, the debugger will stop at such events. However, if the debugger is executing a strict command, it will not stop at an event just because a break point in the stop state applies to it.

If the debugger receives an interrupt (e.g. if the user presses control-C), it will stop at the next event regardless of what command it is executing at the time.

print level When a debugger command steps over some events without user interaction at those events, the *print level* controls under what circumstances the stepped over events will be printed.

- When the print level is `'none'`, none of the stepped over events will be printed.
- When the print level is `'all'`, all the stepped over events will be printed.
- When the print level is `'some'`, the debugger will print the event only if a break point applies to the event.

Regardless of the print level, the debugger will print any event that causes execution to stop and user interaction to start.

default print level

The debugger maintains a default print level. The initial value of this variable is `'some'`, but this value can be overridden by the user.

current environment

Whenever execution stops at an event, the current environment is reset to refer to the stack frame of the call specified by the event. However, the ‘`up`’, ‘`down`’ and ‘`level`’ commands can set the current environment to refer to one of the ancestors of the current call. This will then be the current environment until another of these commands changes the environment yet again or execution continues to another event.

paths in terms

When browsing or printing a term, you can use “`^n`” to refer to the *n*th subterm of that term. If the term’s type has named fields, you can use “`^fname`” to refer to the subterm of the field named ‘`fname`’. You can use several of subterm specifications in a row to refer to subterms deep within the original term. For example, when applied to a list, “`^2`” refers to the tail of the list (the second argument of the list constructor), “`^2^2`” refers to the tail of the tail of the list, and “`^2^2^1`” refers to the head of the tail of the tail, i.e. to the third element of the list. You can think of terms as Unix directories, with constants (function symbols of arity zero) being plain files and function symbols of arity greater than zero being directories themselves. Each subterm specification such as “`^2`” goes one level down in the hierarchy. The exception is the subterm specification “`^..`”, which goes one level up, to the parent of the current directory.

held variables

Normally, the only variables from the program accessible in the debugger are the variables in the current environment at the current program point. However, the user can *hold* variables, causing their values -or selected parts of their values- to stay available for the rest of the debugger session. All the commands that accept variable names also accept the names of held variables; users can ask for a held variable by prefixing the name of the held variable with a dollar sign.

user defined events

Besides the built-in set of events, the Mercury debugger also supports events defined by the user. Each event appears in the source code of the Mercury program as a call prefixed by the keyword ‘`event`’, with each argument of the call giving the value of an event *attribute*. Users can specify the set of user defined events that can appear in a program, and the names, types and order of the attributes of each kind of user defined event, by giving the name of an event set specification file to the compiler when compiling that program. For more details, see [\[User defined events\]](#), page [\[undefined\]](#).

user defined event attributes

Normally, the only variables from the program accessible in the debugger are the variables in the current environment at the current program point. However, if the current event is a user defined event, then the attributes of that event are also available. All the commands that accept variable names also accept the

names of attributes; users can ask for an attribute by prefixing the name of the attribute with an exclamation point.

procedure specification

Some debugger commands, e.g. ‘**break**’, require a parameter that specifies a procedure. The procedure may or may not be a compiler-generated unify, compare, index or init procedure of a type constructor. If it is, the procedure specification has the following components in the following order:

- An optional prefix of the form ‘**unif***’, ‘**comp***’, ‘**indx***’ or ‘**init***’, that specifies whether the procedure belongs to a unify, compare, index or init predicate.
- An optional prefix of the form ‘*module.*’ or ‘*module__*’ that specifies the name of the module that defines the predicate or function to which the procedure belongs.
- The name of the type constructor.
- An optional suffix of the form ‘*/arity*’ that specifies the arity of the type constructor.
- An optional suffix of the form ‘*-modenum*’ that specifies the mode number of the procedure within the predicate or function to which the procedure belongs.

For other procedures, the procedure specification has the following components in the following order:

- An optional prefix of the form ‘**pred***’ or ‘**func***’ that specifies whether the procedure belongs to a predicate or a function.
- An optional prefix of the form ‘*module:*’, ‘*module.*’ or ‘*module__*’ that specifies the name of the module that defines the predicate or function to which the procedure belongs.
- The name of the predicate or function to which the procedure belongs.
- An optional suffix of the form ‘*/arity*’ that specifies the arity of the predicate or function to which the procedure belongs.
- An optional suffix of the form ‘*-modenum*’ that specifies the mode number of the procedure within the predicate or function to which the procedure belongs.

7.8 User defined events

Besides the built-in set of events, the Mercury debugger also supports events defined by the user. The intention is that users can define one kind of event for each semantically important event in the program that isn’t captured by the standard builtin events, and can then generate those events at the appropriate point in the source code. Each event appears in the source code as a call prefixed by the keyword ‘**event**’, with each argument of the call giving the value of an event *attribute*.

Users can specify the set of user defined events that can appear in a program, and the names, types and order of the attributes of each kind of user defined event, by giving the

name of an event set specification file to the compiler when compiling that program as the argument of the ‘`event-set-file-name`’ option. This file should contain a header giving the event set’s name, followed by a sequence of one or more event specifications, like this:

```
event set queens

event nodiag_fail(
  test_failed:  string,
  arg_b:  int,
  arg_d:  int,
  arg_list_len:  int synthesized by list_len_func(sorted_list),
  sorted_list:  list(int) synthesized by list_sort_func(arg_list),
  list_len_func:  function,
  list_sort_func:  function,
  arg_list:  list(int)
)

event safe_test(
  test_list:  listint
)

event noargs
```

The header consists of the keywords ‘`event set`’ and an identifier giving the event set name. Each event specification consists of the keyword ‘`event`’, the name of the event, and, if the event has any attributes, a parenthesized list of those attributes. Each attribute’s specification consists of a name, a colon and information about the attribute.

There are three kinds of attributes.

For ordinary attributes, like ‘`arg_b`’, the information about the attribute is the Mercury type of that attribute.

For function attributes, like ‘`list_sort_func`’, the information about the attribute is just the keyword ‘`function`’.

For synthesized attributes, like ‘`sorted_list`’, the information about the attribute is the type of the attribute, the keywords ‘`synthesized by`’, and a description of the Mercury function call required to synthesize the value of the attribute. The synthesis call consists of the name of a function attribute and a list of the names of one or more argument attributes. Argument attributes cannot be function attributes; they may be either ordinary attributes, or previously synthesized attributes. A synthesized attribute is not allowed to depend on itself directly or indirectly, but there are no restrictions on the positions of synthesized attributes compared to the positions of the function attributes computing them or of the argument attributes of the synthesis functions.

The result types of function attributes are given by the types of the synthesized attributes they compute. The argument types of function attributes (and the number of those arguments) are given by the types of the arguments they are applied to. Each function attribute must be used to compute at least one synthesized attribute, otherwise there would be no way to compute its type. If it is used to compute more than one synthesized attribute, the result and argument types must be consistent.

Each event goal in the program must use the name of one of the events defined here as the predicate name of the call, and the call's arguments must match the types of that event's non-synthesized attributes. Given that *B* and *N* are integers and *L* is a list of integers, these event goals are fine,

```
event nodiag_fail("N - B", B, N, list.length, list.sort, [N | L]),
event safe_test([1, 2, 3])
```

but these goals

```
event nodiag_fail("N - B", B, N, list.sort, list.length, [N | L]),
event nodiag_fail("N - B", B, list.length, N, list.sort, [N | L]),
event safe_test([1], [2])
event safe_test(42)
event nonexistent_event(42)
```

will all generate errors.

The attributes of event calls are always input, and the event goal is always ‘*det*’.

7.9 I/O tabling

In Mercury, predicates that want to do I/O must take a *di/uo* pair of I/O state arguments. Some of these predicates call other predicates to do I/O for them, but some are *I/O primitives*, i.e. they perform the I/O themselves. The Mercury standard library provides a large set of these primitives, and programmers can write their own through the foreign language interface. An I/O action is the execution of one call to an I/O primitive.

In debugging grades, the Mercury implementation has the ability to automatically record, for every I/O action, the identity of the I/O primitive involved in the action and the values of all its arguments. The size of the table storing this information is proportional to the number of *tabled* I/O actions, which are the I/O actions whose details are entered into the table. Therefore the tabling of I/O actions is never turned on automatically; instead, users must ask for I/O tabling to start with the ‘*table_io start*’ command in *mdb*.

The purpose of I/O tabling is to enable transparent retries across I/O actions. (The *mdb* ‘*retry*’ command restores the computation to a state it had earlier, allowing the programmer to explore code that the program has already executed; see its documentation in the [\[Debugger commands\]](#), page [\[undefined\]](#) section below.) In the absence of I/O tabling, retries across I/O actions can have bad consequences. Retry of a goal that reads some input requires that input to be provided twice; retry of a goal that writes some output generates duplicate output. Retry of a goal that opens a file leads to a file descriptor leak; retry of a goal that closes a file can lead to errors (duplicate closes, reads from and writes to closed files).

I/O tabling avoids these problems by making I/O primitives *idempotent*. This means that they will generate their desired effect when they are first executed, but reexecuting them after a retry won't have any further effect. The Mercury implementation achieves this by looking up the action (which is identified by a I/O action number) in the table and returning the output arguments stored in the table for the given action *without* executing the code of the primitive.

Starting I/O tabling when the program starts execution and leaving it enabled for the entire program run will work well for program runs that don't do lots of I/O. For program

runs that *do* lots of I/O, the table can fill up all available memory. In such cases, the programmer may enable I/O tabling with ‘`table_io start`’ just before the program enters the part they wish to debug and in which they wish to be able to perform transparent retries across I/O actions, and turn it off with ‘`table_io stop`’ after execution leaves that part.

The commands ‘`table_io start`’ and ‘`table_io stop`’ can each be given only once during an `mdb` session. They divide the execution of the program into three phases: before ‘`table_io start`’, between ‘`table_io start`’ and ‘`table_io stop`’, and after ‘`table_io stop`’. Retries across I/O will be transparent only in the middle phase.

7.10 Debugger commands

When the debugger (as opposed to the program being debugged) is interacting with the user, the debugger prints a prompt and reads in a line of text, which it will interpret as its next command line. A command line consists of a single command, or several commands separated by semicolons. Each command consists of several words separated by white space. The first word is the name of the command, while any other words give options and/or parameters to the command.

A word may itself contain semicolons or whitespace if it is enclosed in single quotes (‘’). This is useful for commands that have other commands as parameters, for example ‘`view -w 'xterm -e'`’. Characters that have special meaning to ‘`mdb`’ will be treated like ordinary characters if they are escaped with a backslash (\). It is possible to escape single quotes, whitespace, semicolons, newlines and the escape character itself.

Some commands take a number as their first parameter. For such commands, users can type ‘*number command*’ as well as ‘*command number*’. The debugger will treat the former as the latter, even if the number and the command are not separated by white space.

7.10.1 Interactive query commands

```
query module1 module2 ...
cc_query module1 module2 ...
io_query module1 module2 ...
```

These commands allow you to type in queries (goals) interactively in the debugger. When you use one of these commands, the debugger will respond with a query prompt (‘?-’ or ‘`run <--`’), at which you can type in a goal; the debugger will then compile and execute the goal and display the answer(s). You can return from the query prompt to the ‘`mdb>`’ prompt by typing the end-of-file indicator (typically control-D or control-Z), or by typing ‘`quit.`’.

The module names *module1*, *module2*, ... specify which modules will be imported. Note that you can also add new modules to the list of imports directly at the query prompt, by using a command of the form ‘`[module]`’, e.g. ‘`[int]`’. You need to import all the modules that define symbols used in your query. Queries can only use symbols that are exported from a module; entities which are declared in a module’s implementation section only cannot be used.

The three variants differ in what kind of goals they allow. For goals which perform I/O, you need to use `'io_query'`; this lets you type in the goal using DCG syntax. For goals which don't do I/O, but which have determinism `'cc_nondet'` or `'cc_multi'`, you need to use `'cc_query'`; this finds only one solution to the specified goal. For all other goals, you can use plain `'query'`, which finds all the solutions to the goal.

For `'query'` and `'cc_query'`, the debugger will print out all the variables in the goal using `'io.write'`. The goal must bind all of its variables to ground terms, otherwise you will get a mode error.

The current implementation works by compiling the queries on-the-fly and then dynamically linking them into the program being debugged. Thus it may take a little while for your query to be executed. Each query will be written to a file named `'mdb_query.m'` in the current directory, so make sure you don't name your source file `'mdb_query.m'`. Note that dynamic linking may not be supported on some systems; if you are using a system for which dynamic linking is not supported, you will get an error message when you try to run these commands.

You may also need to build your program using shared libraries for interactive queries to work. With Linux on the Intel x86 architecture, the default is for executables to be statically linked, which means that dynamic linking won't work, and hence interactive queries won't work either (the error message is rather obscure: the dynamic linker complains about the symbol `'__data_start'` being undefined). To build with shared libraries, you can use `'MGNUGFLAGS=--pic-reg'` and `'MLFLAGS=--shared'` in your Mmakefile. See the `'README.Linux'` file in the Mercury distribution for more details.

7.10.2 Forward movement commands

`step [-NSans] [num]`

Steps forward *num* events. If this command is given at event *cur*, continues execution until event *cur* + *num*. The default value of *num* is 1.

The options `'-n'` or `'--none'`, `'-s'` or `'--some'`, `'-a'` or `'--all'` specify the print level to use for the duration of the command, while the options `'-S'` or `'--strict'` and `'-N'` or `'--nostrict'` specify the strictness of the command.

By default, this command is not strict, and it uses the default print level.

A command line containing only a number *num* is interpreted as if it were `'step num'`.

An empty command line is interpreted as ‘step 1’.

goto [-NSans] *num*

Continues execution until the program reaches event number *num*. If the current event number is larger than *num*, it reports an error.

The options ‘-n’ or ‘--none’, ‘-s’ or ‘--some’, ‘-a’ or ‘--all’ specify the print level to use for the duration of the command, while the options ‘-S’ or ‘--strict’ and ‘-N’ or ‘--nostrict’ specify the strictness of the command.

By default, this command is strict, and it uses the default print level.

next [-NSans] [*num*]

Continues execution until it reaches the next event of the *num*’th ancestor of the call to which the current event refers. The default value of *num* is zero, which means skipping to the next event of the current call. Reports an error if execution is already at the end of the specified call.

The options ‘-n’ or ‘--none’, ‘-s’ or ‘--some’, ‘-a’ or ‘--all’ specify the print level to use for the duration of the command, while the options ‘-S’ or ‘--strict’ and ‘-N’ or ‘--nostrict’ specify the strictness of the command.

By default, this command is strict, and it uses the default print level.

finish [-NSans] [*num*]

Continues execution until it reaches a final (EXIT, FAIL or EXCP) port of the *num*’th ancestor of the call to which the current event refers. The default value of *num* is zero, which means skipping to the end of the current call. Reports an error if execution is already at the desired port.

The options ‘-n’ or ‘--none’, ‘-s’ or ‘--some’, ‘-a’ or ‘--all’ specify the print level to use for the duration of the command, while the options ‘-S’ or ‘--strict’ and ‘-N’ or ‘--nostrict’ specify the strictness of the command.

By default, this command is strict, and it uses the default print level.

exception [-NSans]

Continues the program until execution reaches an exception event. Reports an error if the current event is already an exception event.

The options ‘-n’ or ‘--none’, ‘-s’ or ‘--some’, ‘-a’ or ‘--all’ specify the print level to use for the duration of the command, while the options ‘-S’ or ‘--strict’ and ‘-N’ or ‘--nostrict’ specify the strictness of the command.

By default, this command is strict, and it uses the default print level.

return [-NSans]

Continues the program until the program finished returning, i.e. until it reaches a port other than EXIT. Reports an error if the current event already refers to such a port.

The options ‘-n’ or ‘--none’, ‘-s’ or ‘--some’, ‘-a’ or ‘--all’ specify the print level to use for the duration of the command, while the options ‘-S’ or ‘--strict’ and ‘-N’ or ‘--nostrict’ specify the strictness of the command.

By default, this command is strict, and it uses the default print level.

user [-NSans]

Continues the program until the next user defined event.

The options ‘-n’ or ‘--none’, ‘-s’ or ‘--some’, ‘-a’ or ‘--all’ specify the print level to use for the duration of the command, while the options ‘-S’ or ‘--strict’ and ‘-N’ or ‘--nostrict’ specify the strictness of the command.

By default, this command is strict, and it uses the default print level.

forward [-NSans]

Continues the program until the program resumes forward execution, i.e. until it reaches a port other than REDO or FAIL. Reports an error if the current event already refers to such a port.

The options ‘-n’ or ‘--none’, ‘-s’ or ‘--some’, ‘-a’ or ‘--all’ specify the print level to use for the duration of the command, while the options ‘-S’ or ‘--strict’ and ‘-N’ or ‘--nostrict’ specify the strictness of the command.

By default, this command is strict, and it uses the default print level.

mindepth [-NSans] *depth*

Continues the program until the program reaches an event whose depth is at least *depth*. Reports an error if the current event already refers to such a port.

The options ‘-n’ or ‘--none’, ‘-s’ or ‘--some’, ‘-a’ or ‘--all’ specify the print level to use for the duration of the command, while the options ‘-S’ or ‘--strict’ and ‘-N’ or ‘--nostrict’ specify the strictness of the command.

By default, this command is strict, and it uses the default print level.

maxdepth [-NSans] *depth*

Continues the program until the program reaches an event whose depth is at most *depth*. Reports an error if the current event already refers to such a port.

The options ‘-n’ or ‘--none’, ‘-s’ or ‘--some’, ‘-a’ or ‘--all’ specify the print level to use for the duration of the command, while the options ‘-S’ or ‘--strict’ and ‘-N’ or ‘--nostrict’ specify the strictness of the command.

By default, this command is strict, and it uses the default print level.

continue [-NSans]

Continues execution until it reaches the end of the program.

The options ‘-n’ or ‘--none’, ‘-s’ or ‘--some’, ‘-a’ or ‘--all’ specify the print level to use for the duration of the command, while the options ‘-S’ or ‘--strict’ and ‘-N’ or ‘--nostrict’ specify the strictness of the command.

By default, this command is not strict. The print level used by the command by default depends on the final strictness level: if the command is strict, it is ‘none’, otherwise it is ‘some’.

7.10.3 Backward movement commands

retry [-fio] [*num*]

If the optional number is not given, restarts execution at the call port of the call corresponding to the current event. If the optional number is given, restarts execution at the call port of the call corresponding to the *num*’th ancestor of the call to which the current event belongs. For example, if *num* is 1, it restarts the parent of the current call.

The command will report an error unless the values of all the input arguments of the selected call are available at the return site at which control would reenter the selected call. (The compiler will keep the values of the input arguments of traced predicates as long as possible, but it cannot keep them beyond the point where they are destructively updated.) The exception is values of type ‘io.state’; the debugger can perform a retry if the only missing value is of type ‘io.state’ (there can be only one io.state at any given time).

Retries over I/O actions are guaranteed to be safe only if the events at which the retry starts and ends are both within the I/O tabled region of the program’s execution. If the retry is not guaranteed to be safe, the debugger will normally ask the user if they really want to do this. The option ‘-f’ or ‘--force’ suppresses the question, telling the debugger that retrying over I/O is OK; the option ‘-o’ or ‘--only-if-safe’ suppresses the question, telling the debugger that retrying over I/O is not OK; the option ‘-i’ or ‘--interactive’ restores the question if a previous option suppressed it.

`track num [termpath]`

Goto the EXIT event of the procedure in which the subterm in argument *num* at term path *termpath* was bound, and display information about where the term was bound.

Note that this command just invokes a script that is equivalent to running the following sequence of commands:

```
dd
browse num
cd termpath
track
info
pd
```

7.10.4 Browsing commands

`vars` Prints the names of all the known variables in the current environment, together with an ordinal number for each variable.

`held_vars` Prints the names of all the held variables.

`print [-fpv] name[termpath]`

`print [-fpv] num[termpath]`

Prints the value of the variable in the current environment with the given name, or with the given ordinal number. If the name or number is followed by a term path such as "²", then only the specified subterm of the given variable is printed. This is a non-interactive version of the `'browse'` command (see below). Various settings which affect the way that terms are printed out (including e.g. the maximum term depth) can be set using the `'format_param'` command.

The options `'-f'` or `'--flat'`, `'-p'` or `'--pretty'`, and `'-v'` or `'--verbose'` specify the format to use for printing.

`print [-fpv] *`

Prints the values of all the known variables in the current environment.

The options `'-f'` or `'--flat'`, `'-p'` or `'--pretty'`, and `'-v'` or `'--verbose'` specify the format to use for printing.

`print [-fpv]`

`print [-fpv] goal`

Prints the goal of the current call in its present state of instantiation.

The options ‘-f’ or ‘--flat’, ‘-p’ or ‘--pretty’, and ‘-v’ or ‘--verbose’ specify the format to use for printing.

print [-fpv] exception

Prints the value of the exception at an EXCP port. Reports an error if the current event does not refer to such a port.

The options ‘-f’ or ‘--flat’, ‘-p’ or ‘--pretty’, and ‘-v’ or ‘--verbose’ specify the format to use for printing.

print [-fpv] action num

Prints a representation of the *num*’th I/O action executed by the program.

The options ‘-f’ or ‘--flat’, ‘-p’ or ‘--pretty’, and ‘-v’ or ‘--verbose’ specify the format to use for printing.

browse [-fpvx] name[termpath]

browse [-fpvx] num[termpath]

Invokes an interactive term browser to browse the value of the variable in the current environment with the given ordinal number or with the given name. If the name or number is followed by a term path such as “^2”, then only the specified subterm of the given variable is given to the browser.

The interactive term browser allows you to selectively examine particular subterms. The depth and size of printed terms may be controlled. The displayed terms may also be clipped to fit within a single screen.

The options ‘-f’ or ‘--flat’, ‘-p’ or ‘--pretty’, and ‘-v’ or ‘--verbose’ specify the format to use for browsing. The ‘-x’ or ‘--xml’ option tells mdb to dump the value of the variable to an XML file and then invoke an XML browser on the file. The XML filename as well as the command to invoke the XML browser can be set using the ‘set’ command. See the documentation for ‘set’ for more details.

For further documentation on the interactive term browser, invoke the ‘browse’ command from within ‘mdb’ and then type ‘help’ at the ‘browser>’ prompt.

browse [-fpvx]

browse [-fpvx] goal

Invokes the interactive term browser to browse the goal of the current call in its present state of instantiation.

The options ‘-f’ or ‘--flat’, ‘-p’ or ‘--pretty’, and ‘-v’ or ‘--verbose’ specify the format to use for browsing. The ‘-x’ or ‘--xml’ option tells mdb to dump the goal to an XML file and then invoke an XML browser on the file. The XML

filename as well as the command to invoke the XML browser can be set using the **set** command. See the documentation for **set** for more details.

browse [-fpvx] exception

Invokes the interactive term browser to browse the value of the exception at an EXCP port. Reports an error if the current event does not refer to such a port.

The options **-f** or **--flat**, **-p** or **--pretty**, and **-v** or **--verbose** specify the format to use for browsing. The **-x** or **--xml** option tells mdb to dump the exception to an XML file and then invoke an XML browser on the file. The XML filename as well as the command to invoke the XML browser can be set using the **set** command. See the documentation for **set** for more details.

browse [-fpvx] action num

Invokes an interactive term browser to browse a representation of the *num*'th I/O action executed by the program.

The options **-f** or **--flat**, **-p** or **--pretty**, and **-v** or **--verbose** specify the format to use for browsing. The **-x** or **--xml** option tells mdb to dump the io action representation to an XML file and then invoke an XML browser on the file. The XML filename as well as the command to invoke the XML browser can be set using the **set** command. See the documentation for **set** for more details.

stack [-d] [-fnumframes] [numlines]

Prints the names of the ancestors of the call specified by the current event. If two or more ancestor calls are for the same procedure, the procedure identification will be printed once with the appropriate multiplicity annotation.

The option **-d** or **--detailed** specifies that for each ancestor call, the call's event number, sequence number and depth should also be printed if the call is to a procedure that is being execution traced.

If the **-f** option, if present, specifies that only the topmost *numframes* stack frames should be printed.

The optional number *numlines*, if present, specifies that only the topmost *numlines* lines should be printed.

This command will report an error if there is no stack trace information available about any ancestor.

up [-d] [num]

Sets the current environment to the stack frame of the *num*'th level ancestor of the current environment (the immediate caller is the first-level ancestor).

If *num* is not specified, the default value is one.

This command will report an error if the current environment doesn't have the required number of ancestors, or if there is no execution trace information about the requested ancestor, or if there is no stack trace information about any of the ancestors between the current environment and the requested ancestor.

The option `'-d'` or `'--detailed'` specifies that for each ancestor call, the call's event number, sequence number and depth should also be printed if the call is to a procedure that is being execution traced.

down `[-d]` [*num*]

Sets the current environment to the stack frame of the *num*'th level descendant of the current environment (the procedure called by the current environment is the first-level descendant).

If *num* is not specified, the default value is one.

This command will report an error if there is no execution trace information about the requested descendant.

The option `'-d'` or `'--detailed'` specifies that for each ancestor call, the call's event number, sequence number and depth should also be printed if the call is to a procedure that is being execution traced.

level `[-d]` [*num*]

Sets the current environment to the stack frame of the *num*'th level ancestor of the call to which the current event belongs. The zero'th ancestor is the call of the event itself.

This command will report an error if the current environment doesn't have the required number of ancestors, or if there is no execution trace information about the requested ancestor, or if there is no stack trace information about any of the ancestors between the current environment and the requested ancestor.

The option `'-d'` or `'--detailed'` specifies that for each ancestor call, the call's event number, sequence number and depth should also be printed if the call is to a procedure that is being execution traced.

current Prints the current event. This is useful if the details of the event, which were printed when control arrived at the event, have since scrolled off the screen.

```
view [-vf2] [-w window-cmd] [-s server-cmd] [-n server-name] [-t timeout]
view -c [-v] [-s server-cmd] [-n server-name]
```

Opens a new window displaying the source code, at the location of the current event. As mdb stops at new events, the window is updated to track through the source code. This requires X11 and a version of ‘vim’ compiled with the client/server option enabled.

The debugger only updates one window at a time. If you try to open a new source window when there is already one open, this command aborts with an error message.

The variant with ‘-c’ (or ‘--close’) does not open a new window but instead attempts to close a currently open source window. The attempt may fail if, for example, the user has modified the source file without saving.

The option ‘-v’ (or ‘--verbose’) prints the underlying system calls before running them, and prints any output the calls produced. This is useful to find out what is wrong if the server does not start.

The option ‘-f’ (or ‘--force’) stops the command from aborting if there is already a window open. Instead it attempts to close that window first.

The option ‘-2’ (or ‘--split-screen’) starts the vim server with two windows, which allows both the callee as well as the caller to be displayed at interface events. The lower window shows what would normally be seen if the split-screen option was not used, which at interface events is the caller. At these events, the upper window shows the callee definition. At internal events, the lower window shows the associated source, and the view in the upper window (which is not interesting at these events) remains unchanged.

The option ‘-w’ (or ‘--window-command’) specifies the command to open a new window. The default is ‘xterm -e’.

The option ‘-s’ (or ‘--server-command’) specifies the command to start the server. The default is ‘vim’.

The option ‘-n’ (or ‘--server-name’) specifies the name of an existing server. Instead of starting up a new server, mdb will attempt to connect to the existing one.

The option ‘-t’ (or ‘--timeout’) specifies the maximum number of seconds to wait for the server to start.

hold *name*[*termpath*] [*heldname*]

Holds on to the variable *name* of the current event, or the part of the specified by *termpath*, even after execution leaves the current event. The held value will stay accessible via the name *\$heldname*. If *heldname* is not specified, it defaults to *name*. There must not already be a held variable named *heldname*.

diff [-s *start*] [-m *max*] *name1*[*termpath1*] *name2*[*termpath2*]

Prints a list of some of the term paths at which the (specified parts of) the specified terms differ. Normally this command prints the term paths of the first 20 differences.

The option ‘-s’ (or ‘--start’), if present, specifies how many of the initial differences to skip.

The option ‘-m’ (or ‘--max’), if present, specifies how many differences to print.

dump [-x] *goal filename*

Writes the goal of the current call in its present state of instantiation to the specified file. The option ‘-x’ (or ‘--xml’) causes the output to be in XML.

dump [-x] *exception filename*

Writes the value of the exception at an EXCP port to the specified file. Reports an error if the current event does not refer to such a port. The option ‘-x’ (or ‘--xml’) causes the output to be in XML.

dump [-x] *name filename*

dump [-x] *num filename*

Writes the value of the variable in the current environment with the given ordinal number or with the given name to the specified file. The option ‘-x’ (or ‘--xml’) causes the output to be in XML.

open *term* Save *term* to a temporary file and open the file in an editor. The environment variable EDITOR is consulted to determine what editor to use. If this environment variable is not set then ‘vi’ is used. *term* may be any term that can be saved to a file with the ‘save_to_file’ command.

grep *pattern term*

Save the given term to a temporary file and perform a grep on the file using *pattern*. *term* may be any term that can be saved to a file with the ‘save_to_file’ command. The unix ‘grep’ command must be available from the shell for this command to work.

list [*num*]

Lists the source code text for the current environment, including *num* preceding and following lines. If *num* is not provided then the default of two is used.

7.10.5 Breakpoint commands

break [-PS] [-Eignore-count] [-Iignore-count] [-n] [-pprint-spec]*
filename:linenumber

Puts a break point on the specified line of the specified source file, if there is an event or a call at that position. If the filename is omitted, it defaults to the filename from the context of the current event.

The options ‘-P’ or ‘--print’, and ‘-S’ or ‘--stop’ specify the action to be taken at the break point.

The options ‘-Eignore-count’ and ‘--ignore-entry ignore-count’ tell the debugger to ignore the breakpoint until after *ignore-count* occurrences of a call event that matches the breakpoint. The options ‘-Iignore-count’ and ‘--ignore-interface ignore-count’ tell the debugger to ignore the breakpoint until after *ignore-count* occurrences of interface events that match the breakpoint.

Each occurrence of the options ‘-pprintspec’ and ‘--print-list *printspec*’ tells the debugger to include the specified entity in the breakpoint’s print list.

Normally, if a variable with the given name or number doesn’t exist when execution reaches the breakpoint, mdb will issue a warning. The option ‘-n’ or ‘--no-warn’, if present, suppresses this warning. This can be useful if e.g. the name is the name of an output variable, which of course won’t be present at call events.

By default, the action of the break point is ‘stop’, the ignore count is zero, and the print list is empty.

break [-AOPSaei] [-Eignore-count] [-Iignore-count] [-n] [-pprint-spec]* *proc-spec*
 Puts a break point on the specified procedure.

The options ‘-A’ or ‘--select-all’, and ‘-O’ or ‘--select-one’ select the action to be taken if the specification matches more than one procedure. If you have specified option ‘-A’ or ‘--select-all’, mdb will put a breakpoint on all matched procedures, whereas if you have specified option ‘-O’ or ‘--select-one’, mdb will report an error. By default, mdb will ask you whether you want to put a breakpoint on all matched procedures or just one, and if so, which one.

The options ‘-P’ or ‘--print’, and ‘-S’ or ‘--stop’ specify the action to be taken at the break point.

The options ‘-a’ or ‘--all’, ‘-e’ or ‘--entry’, and ‘-i’ or ‘--interface’ specify the invocation conditions of the breakpoint. If none of these options are specified, the default is the one indicated by the current scope (see the ‘scope’ command below). The initial scope is ‘interface’.

The options ‘-Eignore-count’ and ‘--ignore-entry ignore-count’ tell the debugger to ignore the breakpoint until after *ignore-count* occurrences of a call event that matches the breakpoint. The options ‘-Iignore-count’ and ‘--ignore-interface ignore-count’ tell the debugger to ignore the breakpoint until after *ignore-count* occurrences of interface events that match the breakpoint.

Each occurrence of the options ‘-pprintspec’ and ‘--print-list printspec’ tells the debugger to include the specified entity in the breakpoint’s print list.

Normally, if a variable with the given name or number doesn’t exist when execution reaches the breakpoint, mdb will issue a warning. The option ‘-n’ or ‘--no-warn’, if present, suppresses this warning. This can be useful if e.g. the name is the name of an output variable, which of course won’t be present at call events.

By default, the action of the break point is ‘stop’, its invocation condition is ‘interface’, the ignore count is zero, and the print list is empty.

```
break [-OPS] [-Eignore-count] [-Iignore-count] [-n] [-pprint-spec]* proc-spec
portname
```

Puts a break point on one or more events of the specified type in the specified procedure. Port names should be specified as they are printed at events, e.g. ‘CALL’, ‘EXIT’, ‘DISJ’, ‘SWTC’, etc.

The option ‘-0’ or ‘--select-one’ selects the action to be taken if the specification matches more than one procedure. If you have specified option ‘-0’ or ‘--select-one’, mdb will report an error; otherwise, mdb will ask you which of the matched procedures you want to select.

If there is only one event of the given type in the specified procedure, mdb will put the breakpoint on it; otherwise, it will ask you whether you want to put a breakpoint on all matched events or just one, and if so, which one.

The options ‘-P’ or ‘--print’, and ‘-S’ or ‘--stop’ specify the action to be taken at the break point.

The options ‘-Eignore-count’ and ‘--ignore-entry ignore-count’ tell the debugger to ignore the breakpoint until after *ignore-count* occurrences of a call event that matches the breakpoint. The options ‘-Iignore-count’

and ‘`--ignore-interface ignore-count`’ tell the debugger to ignore the breakpoint until after *ignore-count* occurrences of interface events that match the breakpoint.

Each occurrence of the options ‘`-pprintspec`’ and ‘`--print-list printspec`’ tells the debugger to include the specified entity in the breakpoint’s print list.

Normally, if a variable with the given name or number doesn’t exist when execution reaches the breakpoint, mdb will issue a warning. The option ‘`-n`’ or ‘`--no-warn`’, if present, suppresses this warning. This can be useful if e.g. the name is the name of an output variable, which of course won’t be present at call events.

By default, the action of the break point is ‘`stop`’, the ignore count is zero, and the print list is empty.

break [-PS] [-Eignore-count] [-Iignore-count] [-n] [-pprint-spec]* here

Puts a break point on the procedure referred to by the current event, with the invocation condition being the event at the current location in the procedure body.

The options ‘`-P`’ or ‘`--print`’, and ‘`-S`’ or ‘`--stop`’ specify the action to be taken at the break point.

The options ‘`-Eignore-count`’ and ‘`--ignore-entry ignore-count`’ tell the debugger to ignore the breakpoint until after *ignore-count* occurrences of a call event that matches the breakpoint. The options ‘`-Iignore-count`’ and ‘`--ignore-interface ignore-count`’ tell the debugger to ignore the breakpoint until after *ignore-count* occurrences of interface events that match the breakpoint.

Each occurrence of the options ‘`-pprintspec`’ and ‘`--print-list printspec`’ tells the debugger to include the specified entity in the breakpoint’s print list.

Normally, if a variable with the given name or number doesn’t exist when execution reaches the breakpoint, mdb will issue a warning. The option ‘`-n`’ or ‘`--no-warn`’, if present, suppresses this warning. This can be useful if e.g. the name is the name of an output variable, which of course won’t be present at call events.

By default, the action of the break point is ‘`stop`’, the ignore count is zero, and the print list is empty.

```
break [-PS] [-Xignore-count] [-n] [-pprint-spec]* user_event [user-event-set]
user-event-name
```

Puts a break point on all user events named *user-event-name*, or, if *user-event-set* is specified as well, on the user event named *user-event-name* in that event set.

The options ‘-P’ or ‘--print’, and ‘-S’ or ‘--stop’ specify the action to be taken at the break point.

The options ‘-Xignore-count’ and ‘--ignore ignore-count’ tell the debugger to ignore the breakpoint until after *ignore-count* occurrences of an event that matches the breakpoint.

Each occurrence of the options ‘-pprintspec’ and ‘--print-list *printspect*’ tells the debugger to include the specified entity in the breakpoint’s print list.

Normally, if a variable with the given name or number doesn’t exist when execution reaches the breakpoint, mdb will issue a warning. The option ‘-n’ or ‘--no-warn’, if present, suppresses this warning. This can be useful if e.g. the name is the name of an output variable, which of course won’t be present at call events.

By default, the action of the break point is ‘stop’, the ignore count is zero, and the print list is empty.

```
break [-PS] [-Xignore-count] [-n] [-pprint-spec]* user_event_set [user-event-set]
```

Puts a break point either on all user events in all event sets, or, if *user-event-set* is specified, on all user events in the event set of the given name.

The options ‘-P’ or ‘--print’, and ‘-S’ or ‘--stop’ specify the action to be taken at the break point.

The options ‘-Xignore-count’ and ‘--ignore ignore-count’ tell the debugger to ignore the breakpoint until after *ignore-count* occurrences of an event that matches the breakpoint.

Each occurrence of the options ‘-pprintspec’ and ‘--print-list *printspect*’ tells the debugger to include the specified entity in the breakpoint’s print list.

Normally, if a variable with the given name or number doesn’t exist when execution reaches the breakpoint, mdb will issue a warning. The option ‘-n’ or ‘--no-warn’, if present, suppresses this warning. This can be useful if e.g. the name is the name of an output variable, which of course won’t be present at call events.

By default, the action of the break point is ‘**stop**’, the ignore count is zero, and the print list is empty.

break info

Lists the details, status and print lists of all break points.

condition [-bbreak-num] [-p] [-v] *varname*[*pathspec*] *op term*

Attaches a condition to the most recent breakpoint, or, if the ‘-b’ or ‘--break-num’ is given, to the breakpoint whose number is given as the argument. Execution won’t stop at the breakpoint if the condition is false.

The condition is a match between a variable live at the breakpoint, or a part thereof, and *term*. It is ok for *term* to contain spaces. The term from the program to be matched is specified by *varname*; if it is followed by *pathspec* (without a space), it specifies that the match is to be against the specified part of *varname*.

There are two kinds of values allowed for *op*. If *op* is ‘=’ or ‘==’, the condition is true if the term specified by *varname* (and *pathspec*, if present) matches *term*. If *op* is ‘!=’ or ‘\!=’, the condition is true if the term specified by *varname* (and *pathspec*, if present) doesn’t match *term*. *term* may contain integers and strings (as long as the strings don’t contain double quotes), but floats and characters aren’t supported (yet), and neither is any special syntax for operators. Operators can be specified in prefix form by quoting them with escaped single quotes, as in ‘\'+\'(1, 2)’. Lists can be specified using the usual syntax. *term* also may not contain variables, with one exception: any occurrence of ‘_’ in *term* matches any term.

If execution reaches a breakpoint and the condition cannot be evaluated, execution will normally stop at that breakpoint with a message to that effect. If the ‘-p’ or ‘--dont-require-path’ option is given, execution won’t stop at breakpoints at which the specified part of the specified variable doesn’t exist. If the ‘-v’ or ‘--dont-require-var’ option is given, execution won’t stop at breakpoints at which the specified variable itself doesn’t exist. The ‘-v’ or ‘--dont-require-var’ option is implicitly assumed if the specified breakpoint is on all user events.

ignore [-Eignore-count] [-Iignore-count] *num*

The options ‘-Eignore-count’ and ‘--ignore-entry ignore-count’ tell the debugger to ignore the breakpoint until after *ignore-count* occurrences of a call event that matches the breakpoint with the specified number. The options ‘-Iignore-count’ and ‘--ignore-interface ignore-count’ tell the debugger to ignore the breakpoint until after *ignore-count* occurrences of interface events that match the breakpoint with the specified number. If neither option is given, the default is to ignore one call event that matches the breakpoint with

the specified number. Reports an error if there is no break point with the specified number.

ignore [-E*ignore-count*] [-I*ignore-count*]

The options ‘-E*ignore-count*’ and ‘--ignore-entry *ignore-count*’ tell the debugger to ignore the breakpoint until after *ignore-count* occurrences of a call event that matches the most recently added breakpoint. The options ‘-I*ignore-count*’ and ‘--ignore-interface *ignore-count*’ tell the debugger to ignore the breakpoint until after *ignore-count* occurrences of interface events that match the most recently added breakpoint. If neither option is given, the default is to ignore one call event that matches the most recently added breakpoint. Reports an error if the most recently added breakpoint has since been deleted.

break_print [-fpv] [-e] [-n] [-b *num*] *print-spec**

Adds the specified print list elements (there may be more than one) to the print list of the breakpoint numbered *num* (if the ‘-b’ or ‘--break-num’ option is given), or to the print list of the most recent breakpoint (if it is not given).

Normally, if a variable with the given name or number doesn’t exist when execution reaches the breakpoint, mdb will issue a warning. The option ‘-n’ or ‘--no-warn’, if present, suppresses this warning. This can be useful if e.g. the name is the name of an output variable, which of course won’t be present at call events.

Normally, the specified elements will be added at the start of the breakpoint’s print list. The option ‘-e’ or ‘--end’, if present, causes them to be added at the end.

By default, the specified elements will be printed with format "flat". The options ‘-f’ or ‘--flat’, ‘-p’ or ‘--pretty’, and ‘-v’ or ‘--verbose’, if given, explicitly specify the format to use.

break_print [-b *num*] none

Clears the print list of the breakpoint numbered *num* (if the ‘-b’ or ‘--break-num’ option is given), or the print list of the most recent breakpoint (if it is not given).

disable *num*

Disables the break point with the given number. Reports an error if there is no break point with that number.

disable * Disables all break points.

disable Disables the most recently added breakpoint. Reports an error if the most recently added breakpoint has since been deleted.

- enable** *num*
Enables the break point with the given number. Reports an error if there is no break point with that number.
- enable *** Enables all break points.
- enable** Enables the most recently added breakpoint. Reports an error if the most recently added breakpoint has since been deleted.
- delete** *num*
Deletes the break point with the given number. Reports an error if there is no break point with that number.
- delete *** Deletes all break points.
- delete** Deletes the most recently added breakpoint. Reports an error if the most recently added breakpoint has already been deleted.
- modules** Lists all the debuggable modules (i.e. modules that have debugging information).
- procedures** *module*
Lists all the procedures in the debuggable module *module*.
- register** [-q]
Registers all debuggable modules with the debugger. Has no effect if this registration has already been done. The debugger will perform this registration when creating breakpoints and when listing debuggable modules and/or procedures. The command will print a message to this effect unless the ‘-q’ or ‘--quiet’ option is given.

7.10.6 I/O tabling commands

- table_io** Reports which phase of I/O tabling we are in at the moment.
- table_io start**
Tells the debugger to start tabling I/O actions.
- table_io stop**
Tells the debugger to stop tabling I/O actions.
- table_io stats**
Reports statistics about I/O tabling.

7.10.7 Parameter commands

`mmc_options option1 option2 ...`

This command sets the options that will be passed to ‘mmc’ to compile your query when you use one of the query commands: ‘query’, ‘cc_query’, or ‘io_query’. For example, if a query results in a compile error, it may sometimes be helpful to use ‘mmc_options --verbose-error-messages’.

`printlevel none`

Sets the default print level to ‘none’.

`printlevel some`

Sets the default print level to ‘some’.

`printlevel all`

Sets the default print level to ‘all’.

`printlevel`

Reports the current default print level.

`scroll on` Turns on user control over the scrolling of sequences of event reports. This means that every screenful of event reports will be followed by a ‘--more--’ prompt. You may type an empty line, which allows the debugger to continue to print the next screenful of event reports. By typing a line that starts with ‘a’, ‘s’ or ‘n’, you can override the print level of the current command, setting it to ‘all’, ‘some’ or ‘none’ respectively. By typing a line that starts with ‘q’, you can abort the current debugger command and get back control at the next event.

`scroll off`

Turns off user control over the scrolling of sequences of event reports.

`scroll size`

Sets the scroll window size to *size*, which tells scroll control to stop and print a ‘--more--’ prompt after every *size* – 1 events. The default value of *size* is the value of the ‘LINES’ environment variable, which should correspond to the number of lines available on the terminal.

`scroll` Reports whether user scroll control is enabled and what the window size is.

`stack_default_limit size`

Set the default number of lines printed by the ‘stack’ and ‘nondet_stack’ commands to *size*. If *size* is zero, the limit is disabled.

`goal_paths on`

Turns on printing of goal paths at events.

`goal_paths off`

Turns off printing of goal paths at events.

`goal_paths`

Reports whether goal paths are printed at events.

`scope all` Sets the default scope of new breakpoints to “all”, i.e. by default, new breakpoints on procedures will stop at all events in the procedure.

`scope interface`

Sets the default scope of new breakpoints to “interface”, i.e. by default, new breakpoints on procedures will stop at all interface events in the procedure.

`scope entry`

Sets the default scope of new breakpoints to “entry”, i.e. by default, new breakpoints on procedures will stop only at events representing calls to the procedure.

`scope` Reports the current default scope of new breakpoints.

`echo on` Turns on the echoing of commands.

`echo off` Turns off the echoing of commands.

`echo` Reports whether commands are being echoed or not.

`context none`

When reporting events or ancestor levels, does not print contexts (filename/line number pairs).

`context before`

When reporting events or ancestor levels, prints contexts (filename/line number pairs) before the identification of the event or call to which they refer, on the same line. With long fully qualified predicate and function names, this may make the line wrap around.

`context after`

When reporting events or ancestor levels, prints contexts (filename/line number pairs) after the identification of the event or call to which they refer, on the same line. With long fully qualified predicate and function names, this may make the line wrap around.

context prevline

When reporting events or ancestor levels, prints contexts (filename/line number pairs) on a separate line before the identification of the event or call to which they refer.

context nextline

When reporting events or ancestor levels, prints contexts (filename/line number pairs) on a separate line after the identification of the event or call to which they refer.

context Reports where contexts are being printed.

user_event_context none

When reporting user-defined events, does not print either filename/line number pairs or procedure ids.

user_event_context file

When reporting user-defined events, prints only filename/line number pairs, not procedure ids.

user_event_context proc

When reporting user-defined events, prints only procedure ids, not filename/line number pairs.

user_event_context full

When reporting user-defined events, prints both filename/line number pairs and procedure ids.

user_event_context

Reports what parts of the context are being printed at user events.

list_context_lines num

Sets the number of lines to be printed by the ‘list’ command printed before and after the target context.

list_context_lines

Prints the number of lines to be printed by the ‘list’ command printed before and after the target context.

list_path dir1 dir2 ...

The ‘list’ command searches a list of directories when looking for a source code file. The ‘list_path’ command sets the search path to the given list of directories.

list_path

When invoked without arguments, the ‘**list_path**’ command prints the search path consulted by the ‘**list**’ command.

push_list_dir *dir1 dir2 ...*

Pushes the given directories on to the search path consulted by the ‘**list**’ command.

pop_list_dir

Pops the leftmost (most recently pushed) directory from the search path consulted by the ‘**list**’ command.

fail_trace_counts *filename*

The declarative debugger can exploit information about the failing and passing test cases to ask better questions. This command tells the ‘**dice**’ command that *filename* contains execution trace counts from failing test cases. The ‘**dice**’ command will use this file unless this is overridden with its ‘**--fail-trace-counts**’ option.

fail_trace_counts

Prints the name of the file containing execution trace counts from failing test cases, if this has already been set.

pass_trace_counts *filename*

The declarative debugger can exploit information about the failing and passing test cases to ask better questions. This command tells the ‘**dice**’ command that *filename* contains execution trace counts from passing test cases. The ‘**dice**’ command will use this file unless this is overridden with its ‘**--pass-trace-counts**’ option.

pass_trace_counts

Prints the name of the file containing execution trace counts from passing test cases, if this has already been set.

max_io_actions *num*

Set the maximum number of I/O actions to print in questions from the declarative debugger to *num*.

max_io_actions

Prints the maximum number of I/O actions to print in questions from the declarative debugger.

xml_browser_cmd *command*

Set the shell command used to launch an XML browser to *command*. If you want a stylesheet to be applied to the XML before the browser is

invoked, then you should do that in this command using the appropriate program (such as `xsltproc`, which comes with `libxslt` and is available from <http://xmlsoft.org/XSLT/>). By default if `xsltproc` and `mozilla` (or `firefox`) are available, `xsltproc` is invoked to apply the `xul_tree.xsl` stylesheet in `extras/xml_stylesheets`, then `mozilla` is invoked on the resulting XUL file.

You can use the apostrophe character (') to quote the command string, for example `"xml_browser_cmd 'firefox file:///tmp/mdbtmp.xml'"`.

`xml_browser_cmd`

Prints the shell command used to launch an XML browser, if this has been set.

`xml_tmp_filename filename`

Tells the debugger to dump XML into the named file before invoking the XML browser. The command named as the argument of `'xml_browser_cmd'` will usually refer to this file.

`xml_tmp_filename`

Prints the temporary filename used for XML browsing, if this has been set.

`format [-APB] format`

Sets the default format of the browser to *format*, which should be one of `'flat'`, `'pretty'` and `'verbose'`.

The browser maintains separate configuration parameters for the three commands `'print *'`, `'print var'`, and `'browse var'`. A `'format'` command applies to all three, unless it specifies one or more of the options `'-A'` or `'--print-all'`, `'-P'` or `'--print'`, and `'-B'` or `'--browse'`, in which case it will set only the selected command's default format.

`format_param [-APBfpv] param value`

Sets one of the parameters of the browser to the given value. The parameter *param* must be one of `'depth'`, `'size'`, `'width'` and `'lines'`.

- `'depth'` is the maximum depth to which subterms will be displayed. Subterms at the depth limit may be abbreviated as functor/arity, or (in lists) may be replaced by an ellipsis (`'...'`). The principal functor of any term has depth zero. For subterms which are not lists, the depth of any argument of the functor is one greater than the depth of the functor. For subterms which are lists, the depth of each element of the list is one greater than the depth of the list.
- `'size'` is the suggested maximum number of functors to display. Beyond this limit, subterms may be abbreviated as functor/arity, or (in lists) may be replaced by an ellipsis (`'...'`). For the purposes of this parameter, the

size of a list is one greater than the sum of the sizes of the elements in the list.

- ‘**width**’ is the width of the screen in characters.
- ‘**lines**’ is the preferred maximum number of lines of one term to display.

The browser maintains separate configuration parameters for the three commands ‘**print ***’, ‘**print var**’, and ‘**browse var**’. A ‘**format_param**’ command applies to all three, unless it specifies one or more of the options ‘**-A**’ or ‘**--print-all**’, ‘**-P**’ or ‘**--print**’, and ‘**-B**’ or ‘**--browse**’, in which case it will set only the selected command’s parameters.

The browser also maintains separate configuration parameters for the different output formats: flat, pretty and verbose. A ‘**format_param**’ command applies to all of these, unless it specifies one or more of the options ‘**-f**’ or ‘**--flat**’, ‘**-p**’ or ‘**--pretty**’, and ‘**-v**’ or ‘**--verbose**’, in which case it will set only the selected format’s parameter.

alias *name* *command* [*command-parameter* ...]

Introduces *name* as an alias for the given command with the given parameters. Whenever a command line has *name* as its first word, the debugger will substitute the given command and parameters for this word before executing the command line.

If *name* is the upper-case word ‘**EMPTY**’, the debugger will substitute the given command and parameters whenever the user types in an empty command line.

If *name* is the upper-case word ‘**NUMBER**’, the debugger will insert the given command and parameters before the command line whenever the user types in a command line that consists of a single number.

unalias *name*

Removes any existing alias for *name*.

7.10.8 Help commands

document_category *slot* *category*

Create a new category of help items, named *category*. The summary text for the category is given by the lines following this command, up to but not including a line containing only the lower-case word ‘**end**’. The list of category summaries printed in response to the command ‘**help**’ is ordered on the integer *slot* numbers of the categories involved.

document *category slot item*

Create a new help item named *item* in the help category *category*. The text for the help item is given by the lines following this command, up to but not including a line containing only the lower-case word ‘end’. The list of items printed in response to the command ‘**help category**’ is ordered on the integer *slot* numbers of the items involved.

help *category item*

Prints help text about the item *item* in category *category*.

help word Prints help text about *word*, which may be the name of a help category or a help item.

help Prints summary information about all the available help categories.

7.10.9 Declarative debugging mdb commands

The following commands relate to the declarative debugger. See [\[Declarative debugging\]](#), page [\[undefined\]](#) for details.

dd [-r] [-R] [-nnodes] [-ssearch-mode] [-ppassfile] [-ffailfile]

Starts declarative debugging using the current event as the initial symptom.

When searching for bugs the declarative debugger needs to keep portions of the execution trace in memory. If it requires a new portion of the trace then it needs to rerun the program. The ‘-nnodes’ or ‘--nodes *nodes*’ option tells the declarative debugger how much of the execution trace to gather when it reruns the program. A higher value for *nodes* requires more memory, but improves the performance of the declarative debugger for long running programs since it will not have to rerun the program as often.

The ‘-ssearch-mode’ or ‘--search-mode *search-mode*’ option tells the declarative debugger which search mode to use. Valid search modes are ‘top_down’ (or ‘td’), ‘divide_and_query’ (or ‘dq’) and ‘suspicion_divide_and_query’ (or ‘sdq’). ‘top_down’ is the default when this option is not given.

Use the ‘-r’ or ‘--resume’ option to continue your previous declarative debugging session. If the ‘--resume’ option is given and there were no previous declarative debugging sessions then the option will be ignored. A ‘dd --resume’ command can be issued at any event. The ‘--search-mode’ option may be used with the ‘--resume’ option to change the search mode of a previously started declarative debugging session.

Use the ‘-R’ or ‘--reset-knowledge-base’ option to reset the declarative debugger’s knowledge base. The declarative debugger will forget any previous answers that have been supplied. It will ask previous questions again if it needs to. This option does not affect what predicates or modules are trusted.

The arguments supplied to the ‘--pass-trace-counts’ (or ‘-p’) and ‘--fail-trace-counts’ (or ‘-f’) options are either trace count files or files containing a list of trace count files. The supplied trace counts are used to assign a suspicion to each event based on which parts of program were executed in the failing test case(s), but not the passing test case(s). This is used to guide the declarative debugger when the suspicion-divide-and-query search mode is used. If the suspicion-divide-and-query search mode is specified then either both the ‘-p’ and ‘-f’ options must be given, or the ‘fail_trace_counts’ and ‘pass_trace_counts’ configuration parameters must be set (using the ‘set’ command).

trust *module-name|proc-spec*

Tells the declarative debugger to trust the given module, predicate or function.

Individual predicates or functions can be trusted by just giving the predicate or function name. If there is more than one predicate or function with the given name then a list of alternatives will be shown.

The entire Mercury standard library is trusted by default and can be untrusted in the usual manner using the ‘untrust’ command. To restore trusted status to the Mercury standard library issue the command ‘trust standard library’ or just ‘trust std lib’.

See also ‘trusted’ and ‘untrust’.

trusted Lists all the trusted modules, predicates and functions. See also ‘trust’ and ‘untrust’.

untrust *num*

Removes the object from the list of trusted objects. *num* should correspond with the number shown in the list produced by issuing a ‘trusted’ command. See also ‘trust’ and ‘trusted’.

7.10.10 Experimental commands

histogram_all *filename*

Prints (to file *filename*) a histogram that counts all events at various depths since the start of the program. This histogram is available only in some experimental versions of the Mercury runtime system.

histogram_exp filename

Prints (to file *filename*) a histogram that counts all events at various depths since the start of the program or since the histogram was last cleared. This histogram is available only in some experimental versions of the Mercury runtime system.

clear_histogram

Clears the histogram printed by ‘**histogram_exp**’, i.e. sets the counts for all depths to zero.

dice [-pfilename] [-ffilename] [-nnum] [-s[pPfFsS]+] [-o filename] [-m module]

Display a program dice on the screen.

A dice is a comparison between some successful test runs of the program and a failing test run. Before using the ‘**dice**’ command one or more passing execution summaries and one failing execution summary need to be generated. This can be done by compiling the program with deep tracing enabled (either by compiling in a `.debug` or `.decldebug` grade or with the ‘**--trace deep**’ or ‘**--trace rep**’ compiler options) and then running the program under `mtc`. This will generate a file with the prefix ‘`.mercury_trace_counts`’ and a unique suffix, that contains a summary of the program’s execution. This summary is called a slice. Copy the generated slice to a new file for each test case, to end up with a failing slice, say ‘`fail`’, and some passing slices, say ‘`pass1`’, ‘`pass2`’, ‘`pass3`’, etc. Union the passing slices with a command such as ‘`mtc_union -p passes pass1 pass2 pass3`’.

The ‘**dice**’ command can use these files to display a table of statistics comparing the passing test runs to the failing run. Here is an example of a dice displayed in an `mdb` session:

```

mdb> dice -f fail -p passes -s S -n 4
Procedure      Path/Port  File:Line Pass (3) Fail Suspicion
pred s.mrg/3-0 <s2;c2;e;> s.m:74      0 (0)    1      1.00
pred s.mrg/3-0 <s2;c2;t;> s.m:67     10 (3)    4      0.29
pred s.mrg/3-0 CALL      s.m:64     18 (3)    7      0.28
pred s.mrg/3-0 EXIT      s.m:64     18 (3)    7      0.28

```

This example tells us that the ‘`else`’ in ‘`s.m`’ on line 74 was executed once in the failing test run, but never in the passing test runs, so this would be a good place to start looking for a bug.

Each row in the table contains statistics about the execution of a separate goal in the program. Six columns are displayed:

- ‘**Procedure**’: The procedure in which the goal appears.

- ‘Path/Port’: The goal path and/or port of the goal. For atomic goals, statistics about the CALL event and the corresponding EXIT, FAIL or EXCP event are displayed on separate rows. For other types of goals the goal path is displayed, except for NEGE, NEGS and NEGF events where the goal path and port are displayed.
- ‘File:Line’: The file name and line number of the goal. This can be used to set a breakpoint on the goal.
- ‘Pass (total passing test runs)’: The total number of times the goal was executed in all the passing test runs. This is followed by a number in parentheses which indicates the number of test runs the goal was executed in. The heading of this column also has a number in parentheses which is the total number of passing test cases. In the example above we can see that 3 passing tests were run.
- ‘Fail’: The number of times the goal was executed in the failing test run.
- ‘Suspicion’: A number between 0 and 1 which gives an indication of how likely a particular goal is to be buggy. The is calculated as $\text{Suspicion} = F / (P + F)$ where F is the number of times the goal was executed in the failing test run and P is the number of times the goal was executed in passing test runs.

The name of the file containing the failing slice can be specified with the ‘-f’ or ‘--fail-trace-counts’ option or with a separate ‘set fail_trace_count filename’ command.

The name of the file containing the union of the passing slices can be given with the ‘-p’ or ‘--pass-trace-counts’ option. Alternatively a separate ‘set pass_trace_counts filename’ command can be given. See [\[Trace counts\]](#), page [\[Trace counts\]](#) for more information about trace counts.

The table can be sorted on the Pass, Fail or Suspicion columns, or a combination of these. This can be done with the ‘-s’ or ‘--sort’ option. The argument of this option is a string made up of any combination of the letters ‘pPfFsS’. The letters in the string indicate how the table should be sorted:

- ‘p’: Pass ascending
- ‘P’: Pass descending
- ‘f’: Fail ascending
- ‘F’: Fail descending
- ‘s’: Suspicion ascending
- ‘S’: Suspicion descending

For example the string "SF" means sort the table by suspicion, descending, and if any two suspicions are the same, then by number of executions in the failing test case, descending.

The option ‘-n’ or ‘--top’ can be used to limit the number lines displayed. Only the top *num* lines, with respect to the ordering specified by the ‘-s’ option, will be displayed. By default the table is limited to 50 lines.

If the ‘-o’ or ‘--output-to-file’ option is given then the output will be written to the specified file instead of being displayed on the screen. Note that the file will be overwritten without warning if it already exists.

The ‘-m’ or ‘--module’ option limits the output to the given module and its submodules, if any.

7.10.11 Miscellaneous commands

source [-i] *filename* [*args*]

Executes the commands in the file named *filename*. Optionally a list of at most nine arguments can be given. Occurrences of the strings "\$1" to "\$9" in the sourced file will be replaced by the corresponding arguments given in the source command before the commands in the sourced file are executed.

Lines that start with a hash (#) character are ignored. Hash characters can be used to place comments in your mdb scripts.

The option ‘-i’ or ‘--ignore-errors’ tells ‘mdb’ not to complain if the named file does not exist or is not readable.

save *filename*

Saves the persistent state of the debugger (aliases, print level, scroll controls, set of breakpoints, browser parameters, set of objects trusted by the declarative debugger, etc) to the specified file. The state is saved in the form of mdb commands, so that sourcing the file will recreate the saved state. Note that this command does not save transient state, such as the current event. There is also a small part of the persistent state (breakpoints established with a ‘break here’ command) that cannot be saved.

quit [-y] Quits the debugger and aborts the execution of the program. If the option ‘-y’ is not present, asks for confirmation first. Any answer starting with ‘y’, or end-of-file, is considered confirmation.

End-of-file on the debugger’s input is considered a quit command.

7.10.12 Developer commands

The following commands are intended for use by the developers of the Mercury implementation.

- var_details**
Prints all the information the debugger has about all the variables at the current program point.
- flag**
Prints the values of all the runtime low-level debugging flags.
- flag *flagname***
Prints the value of the specified runtime low-level debugging flag.
- flag *flagname* on**
Sets the specified runtime low-level debugging flag to true.
- flag *flagname* off**
Sets the specified runtime low-level debugging flag to false.
- subgoal *n*** In minimal model grades, prints the details of the specified subgoal. In other grades, it reports an error.
- consumer *n***
In minimal model grades, prints the details of the specified consumer. In other grades, it reports an error.
- gen_stack**
In minimal model grades, prints the contents of the frames on the generator stack. In other grades, it reports an error.
- cut_stack**
In minimal model grades, prints the contents of the frames on the cut stack. In other grades, it reports an error.
- pneg_stack**
In minimal model grades, prints the contents of the frames on the possible negated context stack. In other grades, it reports an error.
- mm_stacks**
In minimal model grades, prints the contents of the frames on the generator stack, the cut stack and the possible negated context stack. In other grades, it reports an error.
- nondet_stack [-d] [-f*numframes*] [*numlines*]**
Prints the contents of the frames on the nondet stack. By default, it prints only the fixed slots in each nondet stack frame, but if the ‘-d’ or ‘--detailed’ option is given, it will also print the names and values of the live variables in them.

The ‘-f’ option, if present, specifies that only the topmost *numframes* stack frames should be printed.

The optional number *numlines*, if present, specifies that only the topmost *numlines* lines should be printed.

stack_regs

Prints the contents of the virtual machine registers that point to the det and nondet stacks.

all_regs Prints the contents of all the virtual machine registers.

debug_vars

Prints the values of the variables used by the debugger to record event numbers, call sequence numbers and call depths.

stats [-f *filename*] *subject*

Prints statistics about the given subject to standard output, unless the ‘-f’ or ‘--filename’ option is given, in which case it prints the statistic to *filename*.

subject can be ‘procs’, which asks for statistics about proc layout structures in the program.

subject can be ‘labels’, which asks for statistics about label layout structures in the program.

subject can be ‘var_names’, which asks for statistics about the space occupied by variable names in the layout structures in the program.

subject can be ‘io_tabling’, which asks for statistics about the number of times each predicate appears in the I/O action table.

print_optionals

Reports whether optionally-printed values such as typeinfos that are usually of interest only to implementors are being printed or not.

print_optionals on

Tells the debugger to print optionally-printed values.

print_optionals off

Tells the debugger not to print optionally-printed values.

unhide_events

Reports whether events that are normally hidden (that are usually of interest only to implementors) are being exposed or not.

unhide_events on

Tells the debugger to expose events that are normally hidden.

unhide_events off

Tells the debugger to hide events that are normally hidden.

table proc [*num1* ...]

Tells the debugger to print the call table of the named procedure, together with the saved answer (if any) for each call. Reports an error if the named procedure isn't tabled.

For now, this command is supported only for procedures whose arguments are all either integers, floats or strings.

If the user specifies one or more integers on the command line, the output is restricted to the entries in the call table in which the *nth* argument is equal to the *nth* number on the command line.

type_ctor [-fr] *modulename typectorname arity*

Tests whether there is a type constructor defined in the given module, with the given name, and with the given arity. If there isn't, it prints a message to that effect. If there is, it echoes the identity of the type constructor.

If the option '-r' or '--print-rep' option is given, it also prints the name of the type representation scheme used by the type constructor (known as its 'type_ctor_rep' in the implementation).

If the option '-f' or '--print-functors' option is given, it also prints the names and arities of function symbols defined by type constructor.

all_type_ctors [-fr] [*modulename*]

If the user specifies a module name, lists all the type constructors defined in the given module. If the user doesn't specify a module name, lists all the type constructors defined in the whole program.

If the option '-r' or '--print-rep' option is given, it also prints the name of the type representation scheme of each type constructor (known as its 'type_ctor_rep' in the implementation).

If the option '-f' or '--print-functors' option is given, it also prints the names and arities of function symbols defined by each type constructor.

`class_decl [-im] modulename typeclassname arity`

Tests whether there is a type class defined in the given module, with the given name, and with the given arity. If there isn't, it prints a message to that effect. If there is, it echoes the identity of the type class.

If the option '-m' or '--print-methods' option is given, it also lists all the methods of the type class.

If the option '-i' or '--print-instance' option is given, it also lists all the instances of the type class.

`all_class_decls [-im] [modulename]`

If the user specifies a module name, lists all the type classes defined in the given module. If the user doesn't specify a module name, lists all the type classes defined in the whole program.

If the option '-m' or '--print-methods' option is given, it also lists all the methods of each type class.

If the option '-i' or '--print-instance' option is given, it also lists all the instances of each type class.

`all_procedures [-su] [-m modulename] filename`

In the absence of the '-m' or '--module' option, puts a list of all the debuggable procedures in the program into the named file. In the presence of the '-m' or '--module' option, puts a list of all the debuggable procedures in the names module into the named file.

If the '-s' or '--separate' option is given, the various components of procedure names are separated by spaces.

If the '-u' or '--uci' option is given, the list will include the procedures of compiler generated unify, compare, index and initialization predicates. Normally, the list includes the procedures of only user defined predicates.

`ambiguity [-o filename] [-ptf] [modulename ...]`

Print ambiguous procedure, type constructor and/or function symbol names. A procedure name is ambiguous if a predicate or function is defined with that name in more than one module or with more than one arity. A type constructor name is ambiguous if a type constructor is defined with that name in more than one module or with more than one arity. A function symbol name is ambiguous if a function symbol is defined with that name in more than one module or with more than one arity.

If any module names are given, then only those modules are consulted, (any ambiguities involving predicates, functions and type constructors in non-listed modules are ignored). The module names have to be fully qualified, if a module *child* is a submodule of module *parent*, the module name list must include *parent.child*; listing just *child* won't work, since that is not a fully qualified module name.

If the `'-o'` or `'--outputfile'` option is given, the output goes to the file named as the argument of the option; otherwise, it goes to standard output.

If one or more of the `'-p'`, `'-t'`, `'-f'` options or their long equivalents, `'--types'`, or `'--functors'`, this command prints ambiguities only for the indicated kinds of constructs. The default is to print all ambiguities.

`trail_details`

Prints out low-level details of the state of the trail. In other grades, it reports an error.

7.11 Declarative debugging

The debugger incorporates a declarative debugger which can be accessed from its command line. Starting from an event that exhibits a bug, e.g. an event giving a wrong answer, the declarative debugger can find a bug which explains that behaviour using knowledge of the intended interpretation of the program only.

Note that this is a work in progress, so there are some limitations in the implementation.

7.11.1 Overview

The declarative debugger tries to find a bug in your program by asking questions about the correctness of calls executed in your program.

Because pure Mercury code does not have any side effects, the declarative debugger can make inferences such as “if a call produces incorrect output from correct input, then there must be a bug in the code executed by one of the descendents of the call”.

The declarative debugger is therefore able to automate much of the ‘detective work’ that must be done manually when using the procedural debugger.

7.11.2 Concepts

Every CALL event corresponds to an atomic goal, the one printed by the `"print"` command at that event. This atom has the actual arguments in the input argument positions and distinct free variables in the output argument positions (including the return value for functions). We refer to this as the *call atom* of the event.

The same view can be taken of EXIT events, although in this case the outputs as well as the inputs will be bound. We refer to this as the *exit atom* of the event. The exit atom is always an instance of the call atom for the corresponding CALL event.

Using these concepts, it is possible to interpret the events at which control leaves a procedure as assertions about the semantics of the program. These assertions may be true or false, depending on whether or not the program's actual semantics are consistent with its intended semantics.

- | | |
|------|--|
| EXIT | The assertion corresponding to an EXIT event is that the exit atom is valid in the intended interpretation. In other words, the procedure generates correct outputs for the given inputs. |
| FAIL | Every FAIL event has a matching CALL event, and a (possibly empty) set of matching EXIT events between the call and fail. The assertion corresponding to a FAIL event is that every instance of the call atom which is true in the intended interpretation is an instance of one of the exit atoms. In other words, the procedure generates the complete set of answers for the given inputs. (Note that this does not imply that all exit atoms represent correct answers; some exit atoms may in fact be wrong, but the truth of the assertion is not affected by this.) |
| EXCP | Every EXCP event is associated with an exception term, and has a matching CALL event. The assertion corresponding to an EXCP event is that the call atom can abnormally terminate with the given exception. In other words, the thrown exception was expected for that call. |

If one of these assertions is wrong, then we consider the event to represent incorrect behaviour of the program. If the user encounters an event for which the assertion is wrong, then they can request the declarative debugger to diagnose the incorrect behaviour by giving the 'dd' command to the procedural debugger at that event.

7.11.3 Oracle questions

Once the 'dd' command has been given, the declarative debugger asks the user a series of questions about the truth of various assertions in the intended interpretation. The first question in this series will be about the validity of the event for which the 'dd' command was given. The answer to this question will nearly always be "no", since the user has just implied the assertion is false by giving the 'dd' command. Later questions will be about other events in the execution of the program, not all of them necessarily of the same kind as the first.

The user is expected to act as an "oracle" and provide answers to these questions based on their knowledge of the intended interpretation. The debugger provides some help here: previous answers are remembered and used where possible, so questions are not repeated unnecessarily. Commands are available to provide answers, as well as to browse the arguments more closely or to change the order in which the questions are asked. See the next section for details of the commands that are available.

When seeking to determine the validity of the assertion corresponding to an EXIT event, the declarative debugger prints the exit atom followed by the question 'Valid?' for the user to answer. The atom is printed using the same mechanism that the debugger uses to print values, which means some arguments may be abbreviated if they are too large.

When seeking to determine the validity of the assertion corresponding to a FAIL event, the declarative debugger prints the call atom, prefixed by ‘Call’, followed by each of the exit atoms (indented, and on multiple lines if need be), and prints the question ‘Complete?’ (or ‘Unsatisfiable?’ if there are no solutions) for the user to answer. Note that the user is not required to provide any missing instance in the case that the answer is no. (A limitation of the current implementation is that it is difficult to browse a specific exit atom. This will hopefully be addressed in the near future.)

When seeking to determine the validity of the assertion corresponding to an EXCP event, the declarative debugger prints the call atom followed by the exception that was thrown, and prints the question ‘Expected?’ for the user to answer.

In addition to asserting whether a call behaved correctly or not the user may also assert that a call should never have occurred in the first place, because its inputs violated some precondition of the call. For example if an unsorted list is passed to a predicate that is only designed to work with sorted lists. Such calls should be deemed *inadmissible* by the user. This tells the declarative debugger that either the call was given the wrong input by its caller or whatever generated the input is incorrect.

In some circumstances the declarative debugger provides a default answer to the question. If this is the case, the default answer will be shown in square brackets immediately after the question, and simply pressing return is equivalent to giving that answer.

7.11.4 Commands

At the above mentioned prompts, the following commands may be given. Most commands can be abbreviated by their first letter.

It is also legal to press return without specifying a command. If there is a default answer (see [\[Oracle questions\]](#), page [\[undefined\]](#)), pressing return is equivalent to giving that answer. If there is no default answer, pressing return is equivalent to the skip command.

yes Answer ‘yes’ to the current question.

no Answer ‘no’ to the current question.

inadmissible
 Answer that the call is inadmissible.

trust Answer that the predicate or function the question is about does not contain any bugs. However predicates or functions called by this predicate/function may contain bugs. The debugger will not ask you further questions about the predicate or function in the current question.

trust module
 Answer that the module the current question relates to does not contain any bugs. No more questions about any predicates or functions from this module will be asked.

skip Skip this question and ask a different one if possible.

undo Undo the most recent answer or mode change.

mode [**top-down** | **divide-and-query** | **binary**]

Change the current search mode. The search modes may be abbreviated to ‘td’, ‘dq’ and ‘b’ respectively.

browse [--xml] [*n*]

Start the interactive term browser and browse the *n*th argument before answering. If the argument number is omitted then browse the whole call as if it were a data term. While browsing a ‘track’ command may be issued to find the point at which the current subterm was bound (see [\[Improving the search\]](#), page [\[undefined\]](#)). To return to the declarative debugger question issue a ‘quit’ command from within the interactive term browser. For more information on the use of the interactive term browser see the ‘browse’ command in [\[Browsing commands\]](#), page [\[undefined\]](#) or type ‘help’ from within the interactive query browser.

Giving the ‘--xml’ or ‘-x’ option causes the term to be displayed in an XML browser.

browse io [--xml] *n*

Browse the *n*th IO action.

print [*n*] Print the *n*th argument of the current question. If no argument is given then display the current question.

print io *n*

Print the *n*th IO action.

print io *n-m*

Print the *n*th to *m*th IO actions (inclusive).

print io limits

Print the values for which ‘print *n*’ makes sense.

print io Print some I/O actions, starting just after the last action printed (if there was one) or at the first available action (if there was not).

format *format*

Set the default format to *format*, which should be one of ‘flat’, ‘verbose’ and ‘pretty’.

depth *num*

Set the maximum depth to which terms are printed to *num*.

depth io *num*

Set the maximum depth to which I/O actions are printed to *num*. I/O actions are printed using the browser's '**print ***' command so the '**depth io**' command updates the configuration parameters for the browser's '**print ***' command.

size *num* Set the maximum number of function symbols to be printed in terms to *num*.

size io *num*

Set the maximum number of function symbols to be printed in I/O actions to *num*. I/O actions are printed using the browser's '**print ***' command so the '**size io**' command updates the configuration parameters for the browser's '**print ***' command.

width *num*

Set the number of columns in which terms are to be printed to *num*.

width io *num*

Set the number of columns in which I/O actions are to be printed to *num*. I/O actions are printed using the browser's '**print ***' command so the '**width io**' command updates the configuration parameters for the browser's '**print ***' command.

lines *num*

Set the maximum number of lines in terms to be printed to *num*.

lines io *num*

Set the maximum number of lines in I/O actions to be printed to *num*. I/O actions are printed using the browser's '**print ***' command so the '**lines io**' command updates the configuration parameters for the browser's '**print ***' command.

actions *num*

Set the maximum number of I/O actions to be printed in questions to *num*.

params Print the current values of browser parameters.

track [-a] [*term-path*]

The '**track**' command can only be given from within the interactive term browser and tells the declarative debugger to find the point at which the current subterm was bound. If no argument is given the current subterm is taken to be incorrect. If a *term-path* is given then the subterm at *term-path* relative to the current subterm will be considered incorrect. The declarative debugger will

ask about the call that bound the given subterm next. To find out the location of the unification that bound the subterm, issue an ‘**info**’ command when asked about the call that bound the subterm. The declarative debugger can use one of two algorithms to find the point at which the subterm was bound. The first algorithm uses some heuristics to find the subterm more quickly than the second algorithm. It is possible, though unlikely, for the first algorithm to find the wrong call. The first algorithm is the default. To tell the declarative debugger to use the second, more accurate but slower algorithm, give the ‘**-a**’ or ‘**--accurate**’ option to the ‘**track**’ command.

mark [-a] [*term-path*]

The ‘**mark**’ command has the same effect as the ‘**track**’ command except that it also asserts that the atom is inadmissible or erroneous, depending on whether the subterm is input or output respectively.

pd Commence procedural debugging from the current point. This command is notionally the inverse of the ‘**dd**’ command in the procedural debugger. The session can be resumed with a ‘**dd --resume**’ command.

quit End the declarative debugging session and return to the event at which the ‘**dd**’ command was given. The session can be resumed with a ‘**dd --resume**’ command.

info List the filename and line number of the predicate the current question is about as well as the filename and line number where the predicate was called (if this information is available). Also print some information about the state of the bug search, such as the current search mode, how many events are yet to be eliminated and the reason for asking the current question.

help [*command*]

Summarize the list of available commands or give help on a specific command.

7.11.5 Diagnoses

If the oracle keeps providing answers to the asked questions, then the declarative debugger will eventually locate a bug. A “bug”, for our purposes, is an assertion about some call which is false, but for which the assertions about every child of that call are not false (i.e. they are either correct or inadmissible). There are four different classes of bugs that this debugger can diagnose, one associated with each kind of assertion.

Assertions about EXIT events lead to a kind of bug we call an “incorrect contour”. This is a contour (an execution path through the body of a clause) which results in a wrong answer for that clause. When the debugger diagnoses a bug of this kind, it displays the exit atoms in the contour. The resulting incorrect exit atom is displayed last. The program event associated with this bug, which we call the “bug event”, is the exit event at the end of the contour.

Assertions about FAIL events lead to a kind of bug we call a “partially uncovered atom”. This is a call atom which has some instance which is valid, but which is not covered by any

of the applicable clauses. When the debugger diagnoses a bug of this kind, it displays the call atom; it does not, however, provide an actual instance that satisfies the above condition. The bug event in this case is the fail event reached after all the solutions were exhausted.

Assertions about EXCP events lead to a kind of bug we call an “unhandled exception”. This is a contour which throws an exception that needs to be handled but which is not handled. When the debugger diagnoses a bug of this kind, it displays the call atom followed by the exception which was not handled. The bug event in this case is the exception event for the call in question.

If the assertion made by an EXIT, FAIL or EXCP event is false and one or more of the children of the call that resulted in the incorrect EXIT, FAIL or EXCP event is inadmissible, while all the other calls are correct, then an “inadmissible call” bug has been found. This is a call that behaved incorrectly (by producing the incorrect output, failing or throwing an exception) because it passed unexpected input to one of its children. The guilty call is displayed as well as the inadmissible child.

After the diagnosis is displayed, the user is asked to confirm that the event located by the declarative debugger does in fact represent a bug. The user can answer ‘yes’ or ‘y’ to confirm the bug, ‘no’ or ‘n’ to reject the bug, or ‘abort’ or ‘a’ to abort the diagnosis.

If the user confirms the diagnosis, they are returned to the procedural debugger at the event which was found to be the bug event. This gives the user an opportunity, if they need it, to investigate (procedurally) the events in the neighbourhood of the bug.

If the user rejects the diagnosis, which implies that some of their earlier answers may have been mistakes, diagnosis is resumed from some earlier point determined by the debugger. The user may now be asked questions they have already answered, with the previous answer they gave being the default, or they may be asked entirely new questions.

If the user aborts the diagnosis, they are returned to the event at which the ‘dd’ command was given.

7.11.6 Search Modes

The declarative debugger can operate in one of several modes when searching for a bug. Different search modes will result in different sequences of questions being asked by the declarative debugger. The user can specify which mode to use by giving the ‘--search-mode’ option to the ‘dd’ command (see [\[Declarative debugging mdb commands\]](#), page [\[Declarative debugging mdb commands\]](#)) or with the ‘mode’ declarative debugger command (see [\[Declarative debugging commands\]](#), page [\[Declarative debugging commands\]](#)).

7.11.6.1 Top-down Mode

Using this mode the declarative debugger will ask about the children of the last question the user answered ‘no’ to. The child calls will be asked about in the order they were executed. This makes the search more predictable from the user’s point of view as the questions will more or less follow the program execution. The drawback of top-down search is that it may require a lot of questions to be answered before a bug is found, especially with deeply recursive programs.

This search mode is used by default when no other mode is specified.

7.11.6.2 Divide and Query Mode

With this search mode the declarative debugger attempts to halve the size of the search space with each question. In many cases this will result in the bug being found after $O(\log(N))$ questions where N is the number of events between the event where the ‘dd’ command was given and the corresponding ‘CALL’ event. This makes the search feasible for long running programs where top-down search would require an unreasonably large number of questions to be answered. However, the questions may appear to come from unrelated parts of the program which can make them harder to answer.

7.11.6.3 Suspicion Divide and Query Mode

In this search mode the declarative debugger assigns a suspicion level to each event based on which parts of the program were executed in failing test cases, but not in passing test cases. It then attempts to divide the search space into two areas of equal suspicion with each question. This tends to result in questions about parts of the program executed in a failing test case, but not in passing test cases.

7.11.6.4 Binary Search Mode

The user may ask the declarative debugger to do a binary search along the path in the call tree between the current question and the question that the user last answered ‘no’ to. This is useful, for example, when a recursive predicate is producing incorrect output, but the base case is correct.

7.11.7 Improving the search

The number of questions asked by the declarative debugger before it pinpoints the location of a bug can be reduced by giving it extra information. The kind of extra information that can be given and how to convey this information are explained in this section.

7.11.7.1 Tracking suspicious subterms

An incorrect subterm can be tracked to the call that bound the subterm from within the interactive term browser (see [\[Declarative debugging commands\]](#), page [\(undefined\)](#)).

After issuing a ‘track’ command, the next question asked by the declarative debugger will be about the call that bound the incorrect subterm, unless that call was eliminated as a possible bug because of an answer to a previous question or the call that bound the subterm was not traced.

For example consider the following fragment of a program that calculates payments for a loan:

```
:- type payment
    ---> payment(
                date    :: date,
                amount  :: float
            ).
```

```
:- type date ---> date(int, int, int). % date(day, month, year).

:- pred get_payment(loan::in, int::in, payment::out) is det.

get_payment(Loan, PaymentNo, Payment) :-
    get_payment_amount(Loan, PaymentNo, Amount),
    get_payment_date(Loan, PaymentNo, Date),
    Payment = payment(Date, Amount).
```

Suppose that `get_payment` produces an incorrect result and the declarative debugger asks:

```
get_payment(loan(...), 10, payment(date(9, 10, 1977), 10.000000000000)).
Valid?
```

Then if we know that this is the right payment amount for the given loan, but the date is incorrect, we can track the `date(...)` subterm and the debugger will then ask us about `get_payment_date`:

```
get_payment(loan(...), 10, payment(date(9, 10, 1977), 10.000000000000)).
Valid? browse
browser> cd 3/1
browser> ls
date(9, 10, 1977)
browser> track
get_payment_date(loan(...), 10, date(9, 10, 1977)).
Valid?
```

Thus irrelevant questions about `get_payment_amount` are avoided.

If, say, the date was only wrong in the year part, then we could also have tracked the year subterm in which case the next question would have been about the call that constructed the year part of the date.

This feature is also useful when using the procedural debugger. For example, suppose that you come across a ‘CALL’ event and you would like to know the source of a particular input to the call. To find out you could first go to the final event by issuing a ‘finish’ command. Invoke the declarative debugger with a ‘dd’ command and then track the input term you are interested in. The next question will be about the call that bound the term. Issue a ‘pd’ command at this point to return to the procedural debugger. It will now show the final event of the call that bound the term.

Note that this feature is only available if the executable is compiled in a `.decldebug` grade or with the ‘`--trace rep`’ option. If a module is compiled with the ‘`--trace rep`’ option but other modules in the program are not then you will not be able to track subterms through those other modules.

7.11.7.2 Trusting predicates, functions and modules

The declarative debugger can also be told to assume that certain predicates, functions or entire modules do not contain any bugs. The declarative debugger will never ask questions about trusted predicates or functions. It is a good idea to trust standard library modules imported by a program being debugged.

The declarative debugger can be told which predicates/functions it can trust before the ‘`dd`’ command is given. This is done using the ‘`trust`’, ‘`trusted`’ and ‘`untrust`’ commands at the `mdb` prompt (see [\[Declarative debugging mdb commands\]](#), page [\[undefined\]](#) for details on how to use these commands).

Trust commands may be placed in the ‘`.mdbrc`’ file which contains default settings for `mdb` (see [\[Mercury debugger invocation\]](#), page [\[undefined\]](#)). Trusted predicates will also be exported with a ‘`save`’ command (see [\[Miscellaneous commands\]](#), page [\[undefined\]](#)).

During the declarative debugging session the user may tell the declarative debugger to trust the predicate or function in the current question. Alternatively the user may tell the declarative debugger to trust all the predicates and functions in the same module as the predicate or function in the current question. See the ‘`trust`’ command in [\[Declarative debugging commands\]](#), page [\[undefined\]](#).

7.11.7.3 When different search modes are used

If a search mode is given when invoking the declarative debugger then that search mode will be used, unless (a) a subterm is tracked during the session, or (b) the user has not answered ‘`no`’ to any questions yet, in which case top-down search is used until ‘`no`’ is answered to at least one question.

If no search mode is specified with the ‘`dd`’ command then the search mode depends on if the ‘`--resume`’ option is given. If it is then the previous search mode will be used, otherwise top-down search will be used.

You can check the search mode used to find a particular question by issuing an ‘`info`’ command at the question prompt in the declarative debugger. You can also change the search mode from within the declarative debugger with the ‘`mode`’ command.

7.12 Trace counts

A program with debugging enabled may be run in a special mode that causes it to write out to a *trace count file* a record of how many times each *debugger event* in the program was executed during that run.

Trace counts are useful for determining what parts of a failing program are being run and possibly causing the failure; this is called *slicing*. Slices from failing and passing runs can be compared to see which parts of the program are being executed during failing runs, but not during passing runs; this is called *dicing*.

7.12.1 Generating trace counts

To generate a slice for a program run, first compile the program with deep tracing enabled (either by using the ‘`--trace deep`’ option or by compiling the program in a debugging grade). Then invoke the program with the ‘`mtc`’ script, passing any required arguments after the program name.

For example:

```
mtc ./myprog arg1 arg2
```

The program will run as usual, except that when it terminates it will write the number of times each debugger event was executed to a trace count file.

‘`mtc`’ accepts an ‘`-f`’ or ‘`--output-file`’ option. The argument to this option is the filename to use for the generated trace count file. If this option is not given, then the trace count will be written to a file with the prefix ‘`.mercury_trace_counts`’ and a unique suffix.

Ordinarily, the generated trace count file will list only the debugger events that were actually executed during this run. However, it will list all debugger events, even unexecuted ones, if ‘`mtc`’ is given the ‘`-c`’ or ‘`--coverage-test`’ option.

‘`mtc`’ also supports two more options intended for coverage testing: ‘`-s`’ or ‘`--summary-file`’, and ‘`--summary-count`’. These each set an option in the ‘`MERCURY_OPTIONS`’ environment variable, ‘`--trace-count-summary-file`’ and ‘`--trace-count-summary-max`’ respectively. For the documentation of these ‘`mtc`’ options, see the documentation of ‘`MERCURY_OPTIONS`’ environment variable.

Trace count files can be manipulated with the ‘`mtc_union`’ and ‘`mtc_diff`’ tools, and they can be analysed by the ‘`mslice`’ and ‘`mdice`’ tools. They can also be used to help direct a declarative debugging search (see [\[Search Modes\]](#), page [\[undefined\]](#)).

7.12.2 Combining trace counts

The ‘`mtc_union`’ tool can be used to combine several trace count files into one trace count file. You need to use this when you have many trace count files you wish to analyse with ‘`mslice`’ or ‘`mdice`’.

‘`mtc_union`’ is invoked by issuing a command of the form:

```
mtc_union [-v] -o output_file file1 file2 ...
```

‘`file1`’, ‘`file2`’, etc. are the trace count files that should be combined. The new trace count file will be written to ‘`output_file`’. This file will preserve the count of the test cases that contributed to its contents, even if some of ‘`file1`’, ‘`file2`’, etc themselves were created by ‘`mtc_union`’. If the ‘`-v`’ or ‘`--verbose`’ option is specified then a progress message will be displayed as each file is read and its contents merged into the union.

The ‘`mtc_diff`’ tool can be used to subtract one trace count file from another.

‘`mtc_diff`’ is invoked by issuing a command of the form:

```
mtc_diff -o output_file file1 file2
```

‘`file1`’ and ‘`file2`’ must both be trace counts files. The output, written to ‘`output_file`’, will contain the difference between the trace counts in ‘`file1`’ and ‘`file2`’.

for every event that occurs in ‘file1’. Unlike ‘mtc_union’, ‘mtc_diff’ does not preserve the count of the test cases that contributed to its contents in any useful way.

7.12.3 Slicing

Once a slice has been generated it can be viewed in various ways using the mslice tool. The output of the mslice tool will look something like the following:

Procedure	Path/Port	File:Line	Count (1)
pred mrg.merge/3-0	CALL	mrg.m:60	14 (1)
pred mrg.merge/3-0	EXIT	mrg.m:60	14 (1)
pred mrg.msort_n/4-0	CALL	mrg.m:33	12 (1)
pred mrg.msort_n/4-0	EXIT	mrg.m:33	12 (1)
pred mrg.msort_n/4-0	<?;>	mrg.m:35	12 (1)

Each row corresponds to a label in the program. The meanings of the columns are as follows:

- ‘Procedure’: This column displays the procedure that the label relates to.
- ‘Path/Port’: For interface events this column displays the event port, while for internal events it displays the goal path. (See [\[Tracing of Mercury programs\]](#), [page](#) [\[Tracing of Mercury programs\]](#) for an explanation of interface and internal events.)
- ‘File:Line’: This column displays the context of the event.
- ‘Count’: This column displays how many times the event was executed. The number in parentheses for each event row says in how many runs the event was executed. The number in parentheses in the heading row (after the word "Count") indicates how many runs were represented in the trace counts file analysed by the mslice tool.

The mslice tool is invoked using a command of the form:

```
mslice [-s sortspec] [-l N] [-m module] [-n N] [-p N] [-f N] file
```

where ‘file’ is a trace count file, generated either directly by a program run or indirectly by the ‘mtc_union’ or ‘mtc_diff’ tools.

The ‘-s’ or ‘--sort’ option specifies how the output should be sorted. ‘sortspec’ should be a string made up of any combination of the letters ‘cCtT’. Each letter specifies a column and direction to sort on:

- ‘c’: Count ascending
- ‘C’: Count descending
- ‘t’: Number of runs ascending
- ‘T’: Number of runs descending

For example the option ‘-s cT’ will sort the output table by the Count column in ascending order. If the counts for two or more events are the same, then those events will be sorted by number of runs in descending order.

The default is to sort descending on the Count column.

The ‘-l’ or ‘--limit’ option limits the output to ‘N’ lines.

The ‘-m’ or ‘--module’ option limits the output to events only from the given module.

The ‘-n’ or ‘--max-name-column’ option’s argument gives the maximum width of the column containing predicate names. If the argument is zero, there is no maximum width.

The ‘-p’ or ‘--max-path-column’ option’s argument gives the maximum width of the column containing ports and goal paths. If the argument is zero, there is no maximum width.

The ‘-f’ or ‘--max-file-column’ option’s argument gives the maximum width of the column containing file names and line numbers. If the argument is zero, there is no maximum width.

7.12.4 Dicing

A dice is a comparison between passing and failing runs of a program.

Dice are created using the ‘mdice’ tool. To use the ‘mdice’ tool, one must first generate a set of trace count files for passing runs and a set of trace count files for failing runs using the ‘mtc’ tool ([\[Generating trace counts\]](#), [page \[undefined\]\(#\)](#)). Once this has been done, and the union of each set computed using ‘mtc_union’, ‘mdice’ can be used to display a table of statistics that compares the passing runs to the failing runs.

Here is an example of the output of the ‘mdice’ tool:

Procedure	Path/Port	File:Line	Pass (3)	Fail	Suspicion
pred s.mrg/3-0	<s2;c2;e;>	s.m:74	0 (0)	1	1.00
pred s.mrg/3-0	<s2;c2;t;>	s.m:67	10 (3)	4	0.29
pred s.mrg/3-0	CALL	s.m:64	18 (3)	7	0.28
pred s.mrg/3-0	EXIT	s.m:64	18 (3)	7	0.28

This example tells us that the ‘else’ in ‘s.m’ on line 74 was executed once in the failing test run, but never during the passing test runs, so this would be a good place to start looking for a bug.

Each row corresponds to an event in the program. The meanings of the columns are as follows:

- ‘Procedure’: This column displays the procedure the event relates to.
- ‘Path/Port’: For interface events this column displays the event port, while for internal events it displays the goal path. (See [\[Tracing of Mercury programs\]](#), [page \[undefined\]\(#\)](#) for an explanation of interface and internal events.)

- ‘File:Line’: This column displays the context of the event.
- ‘Pass (total passing test runs)’: This column displays the total number of times the event was executed in all the passing test runs. This is followed by a number in parentheses which indicates the number of test runs the event was executed in. The heading of this column also has a number in parentheses which is the total number of passing test cases.
- ‘Fail’: This column displays the number of times the goal was executed in the failing test run(s).
- ‘Suspicion’: This column displays a number between 0 and 1 which gives an indication of how likely a particular goal is to contain a bug. The suspicion is calculated as $\text{Suspicion} = F / (P + F)$ where F is the number of times the goal was executed in failing runs and P is the number of times the goal was executed in passing runs.

The ‘mdice’ tool is invoked with a command of the form:

```
mdice [-s sortspec] [-l N] [-m module] [-n N] [-p N] [-f N] passfile failfile
```

‘passfile’ is a trace count file, generated either directly by a passing program run or as the union of the trace count files of passing program runs. ‘failfile’ is a trace count file, generated either directly by a failing program run or as the union of the trace count files of failing program runs.

The table can be sorted on the Pass, Fail or Suspicion columns, or a combination of these. This can be done with the ‘-s’ or ‘--sort’ option. The argument of this option is a string made up of any combination of the letters ‘pPfFsS’. The letters in the string indicate how the table should be sorted:

- ‘p’: Pass ascending
- ‘P’: Pass descending
- ‘f’: Fail ascending
- ‘F’: Fail descending
- ‘s’: Suspicion ascending
- ‘S’: Suspicion descending

For example the string "SF" means sort the table by suspicion in descending order, and if any two suspicions are the same, then by number of executions in the failing run(s), also in descending order.

The default is to sort descending on the Suspicion column.

The option ‘-l’ or ‘--limit’ can be used to limit the number of lines displayed.

The ‘-m’ or ‘--module’ option limits the output to the given module and any submodules.

The ‘-n’ or ‘--max-name-column’ option’s argument gives the maximum width of the column containing predicate names. If the argument is zero, there is no maximum width.

The ‘-p’ or ‘--max-path-column’ option’s argument gives the maximum width of the column containing ports and goal paths. If the argument is zero, there is no maximum width.

The ‘-f’ or ‘--max-file-column’ option’s argument gives the maximum width of the column containing file names and line numbers. If the argument is zero, there is no maximum width.

7.12.5 Coverage testing

Coverage testing is the process of finding out which parts of the code of a program are not executed during any test case, so that new test cases can be designed specifically to exercise those parts.

The first step in coverage testing a Mercury program is compiling that program with execution tracing enabled, either by using the ‘--trace deep’ option or by compiling the program in a debugging grade. The second step is to execute that program on all its test cases with coverage testing enabled. This can be done either by running the program with ‘mtc --coverage-test’, or by including one of the corresponding options (‘--coverage-test’ or ‘--coverage-test-if-exec=programname’) in the value of the ‘MERCURY_OPTIONS’ environment variable. These runs generate a set of trace counts files that can be given to the Mercury test coverage tool, the ‘mcov’ program.

The ‘mcov’ tool is invoked with a command of the form:

```
mcov [-d] [-v] [-o output_file] tracecountfile1 ...
```

The arguments consist of one or more trace count files. The output will normally be a list of all the procedures in the program that were not executed in any of the runs that generated these trace count files. The output will go to standard output unless this is overridden by the ‘-o’ or ‘--output-file’ option.

If the ‘-d’ or ‘--detailed’ option is specified, then the output will list all the *events* in the program that were not executed in any of these runs. This option can thus show the unexecuted parts of the executed procedures.

If the ‘-v’ or ‘--verbose’ option is specified, then a progress message will be displayed as each file is read.

8 Profiling

8.1 Profiling introduction

To obtain the best trade-off between productivity and efficiency, programmers should not spend too much time optimizing their code until they know which parts of the code are really taking up most of the time. Only once the code has been profiled should the programmer consider making optimizations that would improve efficiency at the expense of readability or ease of maintenance. A good profiler is therefore a tool that should be part of every software engineer’s toolkit.

Mercury programs can be analyzed using two distinct profilers. The Mercury profiler ‘`mprof`’ is a conventional call-graph profiler (or graph profiler for short) in the style of ‘`gprof`’. The Mercury deep profiler ‘`mdprof`’ is a new kind of profiler that associates a lot more context with each measurement. ‘`mprof`’ can be used to profile either time or space, but not both at the same time; ‘`mdprof`’ can profile both time and space at the same time.

The parallel execution of Mercury programs can be analyzed with a third profiler called ‘`threadscope`’. ‘`threadscope`’ allows programmers to visualise CPU utilisation for work, garbage collection, and idle time. This enables programmers to see the effect of parallelization decisions such as task granularity. The ‘`threadscope`’ tool is not included with the Melbourne Mercury Compiler. See <http://research.microsoft.com/en-us/projects/threadscope/>, Threadscope: Performance Tuning Parallel Haskell Programs.

8.2 Building profiled applications

To enable profiling, your program must be built with profiling enabled. The three different profilers require different support, and thus you must choose which one to enable when you build your program.

- To build your program with time profiling enabled for ‘`mprof`’, pass the ‘`-p`’ (‘`--profiling`’) option to ‘`mmc`’ (and also to ‘`mgnuc`’ and ‘`ml`’, if you invoke them separately).
- To build your program with memory profiling enabled for ‘`mprof`’, pass the ‘`--memory-profiling`’ option to ‘`mmc`’, ‘`mgnuc`’ and ‘`ml`’.
- To build your program with deep profiling enabled (for ‘`mdprof`’), pass the ‘`--deep-profiling`’ option to ‘`mmc`’, ‘`mgnuc`’ and ‘`ml`’.
- To build your program with threadscope profiling enabled (for ‘`threadscope`’). pass the ‘`--parallel`’ and ‘`--threadscope`’ options to ‘`mmc`’, ‘`mgnuc`’ and ‘`ml`’.

If you are using Mmake, then you pass these options to all the relevant programs by setting the ‘`GRADEFLAGS`’ variable in your Mmakefile, e.g. by adding the line ‘`GRADEFLAGS--profiling`’. (For more information about the different grades, see [\[Compilation model options\]](#), page [\[undefined\]](#).)

Enabling ‘`mprof`’ or ‘`mdprof`’ profiling has several effects. First, it causes the compiler to generate slightly modified code, which counts the number of times each predicate or function is called, and for every call, records the caller and callee. With deep profiling, there are other modifications as well, the most important impact of which is the loss of tail-recursion for groups of mutually tail-recursive predicates (self-tail-recursive predicates stay tail-recursive). Second, your program will be linked with versions of the library and

runtime that were compiled with the same kind of profiling enabled. Third, if you enable graph profiling, the compiler will generate for each source file the static call graph for that file in `'module.prof'`.

Enabling `'threadscope'` profiling causes the compiler to build the program against a different runtime system. This runtime system logs events relevant to parallel execution. `'threadscope'` support is not compatible with all processors, see `'README.ThreadScope'` for more information.

8.3 Creating profiles

Once you have created a profiled executable, you can gather profiling information by running the profiled executable on some test data that is representative of the intended uses of the program. The profiling version of your program will collect profiling information during execution, and save this information at the end of execution, provided execution terminates normally and not via an abort.

Executables compiled with `'--profiling'` save profiling data in the files `'Prof.Counts'`, `'Prof.Decls'`, and `'Prof.CallPair'`. (`'Prof.Decl'` contains the names of the procedures and their associated addresses, `'Prof.CallPair'` records the number of times each procedure was called by each different caller, and `'Prof.Counts'` records the number of times that execution was in each procedure when a profiling interrupt occurred.) Executables compiled with `'--memory-profiling'` will use two of those files (`'Prof.Decls'` and `'Prof.CallPair'`) and a two others: `'Prof.MemoryWords'` and `'Prof.MemoryCells'`. Executables compiled with `'--deep-profiling'` save profiling data in a single file, `'Deep.data'`. Executables compiled with the `'--threadscope'` option write profiling data to a file whose name is that of the program being profiled with the extension `'.eventlog'`. For example, the profile for the program `'my_program'` would be written to the file `'my_program.eventlog'`.

It is also possible to combine `'mprof'` profiling results from multiple runs of your program. You can do by running your program several times, and typing `'mprof_merge_counts'` after each run. It is not (yet) possible to combine `'mdprof'` profiling results from multiple runs of your program.

Due to a known timing-related bug in our code, you may occasionally get segmentation violations when running your program with `'mprof'` profiling enabled. If this happens, just run it again — the problem occurs only very rarely. The same vulnerability does not occur with `'mdprof'` profiling.

With the `'mprof'` and `'mdprof'` profilers, you can control whether time profiling measures real (elapsed) time, user time plus system time, or user time only, by including the options `'-Tr'`, `'-Tp'`, or `'-Tv'` respectively in the environment variable `MERCURY_OPTIONS` when you run the program to be profiled. Currently, the `'-Tp'` and `'-Tv'` options don't work on Windows, so on Windows you must explicitly specify `'-Tr'`.

The default is user time plus system time, which counts all time spent executing the process, including time spent by the operating system working on behalf of the process, but not including time that the process was suspended (e.g. due to time slicing, or while waiting for input). When measuring real time, profiling counts even periods during which the process was suspended. When measuring user time only, profiling does not count time inside the operating system at all.

8.4 Using mprof for time profiling

To display the graph profile information gathered from one or more profiling runs, just type `mprof` or `mprof -c`. (For programs built with `--high-level-code`, you need to also pass the `--no-demangle` option to `mprof` as well.) Note that `mprof` can take quite a while to execute (especially with `-c`), and will usually produce quite a lot of output, so you will usually want to redirect the output into a file with a command such as `mprof > mprof.out`.

The output of `mprof -c` consists of three major sections. These are named the call graph profile, the flat profile and the alphabetic listing. The output of `mprof` contains the flat profile and the alphabetic listing only.

The call graph profile presents the local call graph of each procedure. For each procedure it shows the parents (callers) and children (callees) of that procedure, and shows the execution time and call counts for each parent and child. It is sorted on the total amount of time spent in the procedure and all of its descendents (i.e. all of the procedures that it calls, directly or indirectly.)

The flat profile presents the just execution time spent in each procedure. It does not count the time spent in descendents of a procedure.

The alphabetic listing just lists the procedures in alphabetical order, along with their index number in the call graph profile, so that you can quickly find the entry for a particular procedure in the call graph profile.

The profiler works by interrupting the program at frequent intervals, and each time recording the currently active procedure and its caller. It uses these counts to determine the proportion of the total time spent in each procedure. This means that the figures calculated for these times are only a statistical approximation to the real values, and so they should be treated with some caution. In particular, if the profiler's assumption that calls to a procedure from different callers have roughly similar costs is not true, the graph profile can be quite misleading.

The time spent in a procedure and its descendents is calculated by propagating the times up the call graph, assuming that each call to a procedure from a particular caller takes the same amount of time. This assumption is usually reasonable, but again the results should be treated with caution. (The deep profiler does not make such an assumption, and hence its output is significantly more reliable.)

Note that any time spent in a C function (e.g. time spent in `GC_malloc()`, which does memory allocation and garbage collection) is credited to the Mercury procedure that called that C function.

Here is a small portion of the call graph profile from an example program.

index	%time	self	descendents	called/total	parents	index
				called+self	name	
				called/total	children	
					<spontaneous>	
[1]	100.0	0.00	0.75	0	call_engine_label	[1]
		0.00	0.75	1/1	do_interpreter	[3]

[2]	100.0	0.00	0.75	1/1	do_interpreter [3]
		0.00	0.75	1	io.run/0(0) [2]
		0.00	0.00	1/1	io.init_state/2(0) [11]
		0.00	0.74	1/1	main/2(0) [4]

[3]	100.0	0.00	0.75	1/1	call_engine_label [1]
		0.00	0.75	1	do_interpreter [3]
		0.00	0.75	1/1	io.run/0(0) [2]

[4]	99.9	0.00	0.74	1/1	io.run/0(0) [2]
		0.00	0.74	1	main/2(0) [4]
		0.00	0.74	1/1	sort/2(0) [5]
		0.00	0.00	1/1	print_list/3(0) [16]
		0.00	0.00	1/10	io.write_string/3(0) [18]

[5]	99.9	0.00	0.74	1/1	main/2(0) [4]
		0.00	0.74	1	sort/2(0) [5]
		0.05	0.65	1/1	list.perm/2(0) [6]
		0.00	0.09	40320/40320	sorted/1(0) [10]

[6]	86.6			8	list.perm/2(0) [6]
		0.05	0.65	1/1	sort/2(0) [5]
		0.05	0.65	1+8	list.perm/2(0) [6]
		0.00	0.60	5914/5914	list.insert/3(2) [7]
			8	list.perm/2(0) [6]	

[7]	80.0	0.00	0.60	5914/5914	list.perm/2(0) [6]
		0.00	0.60	5914	list.insert/3(2) [7]
		0.60	0.60	5914/5914	list.delete/3(3) [8]

[8]	80.0			40319	list.delete/3(3) [8]
		0.60	0.60	5914/5914	list.insert/3(2) [7]
		0.60	0.60	5914+40319	list.delete/3(3) [8]
				40319	list.delete/3(3) [8]

		0.00	0.00	3/69283	tree234.set/4(0) [15]
		0.09	0.09	69280/69283	sorted/1(0) [10]
[9]	13.3	0.10	0.10	69283	compare/3(0) [9]
		0.00	0.00	3/3	--Compare--io--stream/0(0) [20]
		0.00	0.00	69280/69280	builtin_compare_int/3(0) [27]

		0.00	0.09	40320/40320	sort/2(0) [5]
[10]	13.3	0.00	0.09	40320	sorted/1(0) [10]
		0.09	0.09	69280/69283	compare/3(0) [9]

The first entry is ‘call_engine_label’ and its parent is ‘<spontaneous>’, meaning that it is the root of the call graph. (The first three entries, ‘call_engine_label’, ‘do_interpreter’, and ‘io.run/0’ are all part of the Mercury runtime; ‘main/2’ is the entry point to the user’s program.)

Each entry of the call graph profile consists of three sections, the parent procedures, the current procedure and the children procedures.

Reading across from the left, for the current procedure the fields are:

- The unique index number for the current procedure. (The index numbers are used only to make it easier to find a particular entry in the call graph.)
- The percentage of total execution time spent in the current procedure and all its descendents. As noted above, this is only a statistical approximation.
- The “self” time: the time spent executing code that is part of current procedure. As noted above, this is only a statistical approximation.
- The descendent time: the time spent in the current procedure and all its descendents. As noted above, this is only a statistical approximation.
- The number of times a procedure is called. If a procedure is (directly) recursive, this column will contain the number of calls from other procedures, a plus sign, and then the number of recursive calls. These numbers are exact, not approximate.
- The name of the procedure followed by its index number.

The predicate or function names are not just followed by their arity but also by their mode in brackets. A mode of zero corresponds to the first mode declaration of that predicate in the source code. For example, ‘list.delete/3(3)’ corresponds to the ‘(out, out, in)’ mode of ‘list.delete/3’.

Now for the parent and child procedures the self and descendent time have slightly different meanings. For the parent procedures the self and descendent time represent the proportion of the current procedure’s self and descendent time due to that parent. These times are obtained using the assumption that each call contributes equally to the total time of the current procedure.

8.5 Using mprof for memory profiling

To create a memory profile, you can invoke `mprof` with the `-m` (`--profile memory-words`) option. This will profile the amount of memory allocated, measured in units of words. (A word is 4 bytes on a 32-bit architecture, and 8 bytes on a 64-bit architecture.)

Alternatively, you can use `mprof`'s `-M` (`--profile memory-cells`) option. This will profile memory in units of “cells”. A cell is a group of words allocated together in a single allocation, to hold a single object. Selecting this option this will therefore profile the number of memory allocations, while ignoring the size of each memory allocation.

With memory profiling, just as with time profiling, you can use the `-c` (`--call-graph`) option to display call graph profiles in addition to flat profiles.

Note that Mercury's memory profiler will only tell you about allocation, not about deallocation (garbage collection). It can tell you how much memory was allocated by each procedure, but it won't tell you how long the memory was live for, or how much of that memory was garbage-collected. This is also true for `mdprof`.

8.6 Using mdprof

The user interface of the deep profiler is a browser. To display the information contained in a deep profiling data file (which will be called `Deep.data` unless you renamed it), start up your browser and give it a URL of the form `http://server.domain.name/cgi-bin/mdprof.cgi?/full/path/name/Deep.data`. The `server.domain.name` part should be the name of a machine with the following qualifications: it should have a web server running on it, and it should have the `mdprof.cgi` program installed in the web server's CGI program directory. (On many Linux systems, this directory is `/usr/lib/cgi-bin`.) The `/full/path/name/Deep.data` part should be the full path name of the deep profiling data file whose data you wish to explore. The name of this file must not have percent signs in it, and it must end in the suffix `.data`.

When you start up `mdprof` using the command above, you will see a list of the usual places where you may want to start looking at the profile. Each place is represented by a link. Clicking on and following that link will give you a web page that contains both the profile information you asked for and other links, some of which present the same information in a different form and some of which lead to further information. You explore the profile by clicking on links and looking at the resulting pages.

The deep profiler can generate several kinds of pages.

The menu page

The menu page gives summary information about the profile, and the usual starting points for exploration.

Clique pages

Clique pages are the most fundamental pages of the deep profiler. Each clique page presents performance information about a clique, which is either a single procedure or a group of mutually recursive procedures, in a given ancestor

context, which in turn is a list of other cliques starting with the caller of the entry point of the clique and ending with the clique of the ‘**main**’ predicate.

Each clique page lists the closest ancestor cliques, and then the procedures of the clique. It gives the cost of each call site in each procedure, as well as the cost of each procedure in total. These costs will be just those incurred in the given ancestor context; the costs incurred by these call sites and procedures in other ancestor contexts will be shown on other clique pages.

Procedure pages

Procedure pages give the total cost of a procedure and its call sites in all ancestor contexts.

Module pages

Module pages give the total cost of all the procedures of a module.

Module getters and setters pages

These pages identifies the getter and setter procedures in a module. Getters and setters are simply predicates and functions that contain ‘**_get_**’ and ‘**_set_**’ respectively in their names; they are usually used to access fields of data structures.

Program modules page

The program modules page gives the list of the program’s modules.

Top procedure pages

Top procedure pages identify the procedures that are most expensive as measured by various criteria.

Procedure caller pages

A procedure caller page lists the call sites, procedures, modules or cliques that call the given procedure.

When exploring a procedure’s callers, you often want only the ancestors that are at or above a certain level of abstraction. Effectively you want to draw a line through the procedures of the program, such that you are interested in the procedures on or above the line but those below the line. Since we want to exclude procedures below the line from procedure caller pages, we call this line an *exclusion contour*.

You can tell the deep profiler where you want to draw this line by giving it a ‘**exclusion contour file**’. The name of this file should be the same as the name of the deep profiling data file, but with the suffix ‘**.data**’ replaced with ‘**.contour**’. This file should consist of a sequence of lines, and each line should contain two words. The first word should be either ‘**all**’ or ‘**internal**’; the second should the name of a module. If the first word is ‘**all**’, then all procedures in the named module are below the exclusion contour; if the first word is ‘**internal**’, then all internal (non-exported) procedures in the named module are below the exclusion contour. Here is an example of an exclusion contour file.

```
all bag
all list
all map
internal set
```

8.7 Using threadscope

The ThreadScope tools are not distributed with Mercury. For information about how to install them please see the `README.ThreadScope` file included in the Mercury distribution.

ThreadScope provides two programs that can be used to view profiles in `.eventlog` files. The first, `show-ghc-events`, lists the ThreadScope events sorted from the earliest to the latest, while the second, `threadscope` provides a graphical display for browsing the profile.

Both programs accept the name of a `.eventlog` file on the command line. The `threadscope` program also provides a menu from which users can choose a file to open.

8.8 Profiling and shared libraries

On some operating systems, Mercury's profiling doesn't work properly with shared libraries. The symptom is errors (`map.lookup failed`) or warnings from `mprof`. On some systems, the problem occurs because the C implementation fails to conform to the semantics specified by the ISO C standard for programs that use shared libraries. For other systems, we have not been able to analyze the cause of the failure (but we suspect that the cause may be the same as on those systems where we have been able to analyze it).

If you get errors or warnings from `mprof`, and your program is dynamically linked, try rebuilding your application statically linked, e.g. by using `MLFLAGS=--static` in your Mmakefile. Another work-around that sometimes works is to set the environment variable `LD_BIND_NOW` to a non-null value before running the program.

9 Invocation

This section contains a brief description of all the options available for `mmc`, the Mercury compiler. Sometimes this list is a little out-of-date; use `mmc --help` to get the most up-to-date list.

9.1 Invocation overview

`mmc` is invoked as

```
mmc [options] arguments
```

Arguments can be either module names or file names. Arguments ending in `.m` are assumed to be file names, while other arguments are assumed to be module names, with `.'` (rather than `__` or `:`) as module qualifier. If you specify a module name such as `foo.bar.baz`, the compiler will look for the source in files `foo.bar.baz.m`, `bar.baz.m`, and `baz.m`, in that order.

Options are either short (single-letter) options preceded by a single `-`, or long options preceded by `--`. Options are case-sensitive. We call options that do not take arguments *flags*. Single-letter flags may be grouped with a single `-`, e.g. `-vVc`. Single-letter flags may be negated by appending another trailing `-`, e.g. `-v-`. Long flags may be negated by preceding them with `no-`, e.g. `--no-verbose`.

9.2 Warning options

`-w`

`--inhibit-warnings`

Disable all warning messages.

`--halt-at-warn`

This option causes the compiler to treat all warnings as if they were errors. This means that if any warning is issued, the compiler will not generate code — instead, it will return a non-zero exit status.

`--halt-at-syntax-error`

This option causes the compiler to halt immediately after syntax checking and not do any semantic checking if it finds any syntax errors in the program.

`--inhibit-accumulator-warnings`

Don't warn about argument order rearrangement caused by `--introduce-accumulators`.

`--no-warn-singleton-variables`

Don't warn about variables which only occur once.

`--no-warn-missing-det-decls`

For predicates that are local to a module (those that are not exported), don't issue a warning if the `'pred'` or `'mode'` declaration does not have a determinism annotation. Use this option if you want the compiler to perform automatic determinism inference for non-exported predicates.

`--no-warn-det-decls-too-lax`

Don't warn about determinism declarations which could have been stricter.

`--no-warn-inferred-erroneous`

Don't warn about procedures whose determinism is inferred erroneous but whose determinism declarations are laxer.

`--no-warn-insts-without-matching-type`

Don't warn about insts that are not consistent with any types in scope.

`--warn-unused-imports`

Warn about modules that are imported but not used.

`--no-warn-nothing-exported`

Don't warn about modules whose interface sections have no exported predicates, functions, insts, modes or types.

- `--warn-unused-args`
Warn about predicate or function arguments which are not used.
- `--warn-interface-imports`
Warn about modules imported in the interface which are not used in the interface.
- `--warn-missing-opt-files`
Warn about `.opt` files that cannot be opened.
- `--warn-missing-trans-opt-files`
Warn about `.trans_opt` files that cannot be opened.
- `--no-warn-non-contiguous-clauses`
Do not generate a warning if the clauses of a predicate or function are not contiguous.
- `--warn-non-contiguous-foreign-procs`
Generate a warning if the clauses and `foreign_procs` of a predicate or function are not contiguous.
- `--warn-non-stratification`
Warn about possible non-stratification of the predicates and/or functions in the module. Non-stratification occurs when a predicate or function can call itself negatively through some path along its call graph.
- `--no-warn-simple-code`
Disable warnings about constructs which are so simple that they are likely to be programming errors.
- `--warn-duplicate-calls`
Warn about multiple calls to a predicate with the same input arguments.
- `--no-warn-missing-module-name`
Disable warnings for modules that do not start with a `:- module` declaration.
- `--no-warn-wrong-module-name`
Disable warnings for modules whose `:- module` declaration does not match the module's file name.
- `--no-warn-smart-recompilation`
Disable warnings from the smart recompilation system.
- `--no-warn-undefined-options-variables`
Warn about references to undefined variables in options files with `--make`.

--warn-non-tail-recursion

Warn about any directly recursive calls that are not tail recursive. This option also requires ‘--high-level-code’.

--no-warn-target-code

Disable warnings from the compiler used to process the target code (e.g. gcc).

--no-warn-up-to-date

Don’t warn if targets specified on the command line with ‘--make’ are already up to date.

--no-warn-stubs

Disable warnings about procedures for which there are no clauses. Note that this option only has any effect if the ‘--allow-stubs’ option (see [\(undefined\)](#) [Language semantics options], page [\(undefined\)](#)) is enabled.

--warn-dead-procs

Warn about procedures which are never called.

--no-warn-table-with-inline

Disable warnings about tabled procedures that also have a ‘pragma inline’ declaration.

--no-warn-non-term-special-preds

Do not warn about types that have user-defined equality or comparison predicates that cannot be proved to terminate. This option is only enabled when termination analysis is enabled. (See [\(undefined\)](#) [Termination analysis options], page [\(undefined\)](#) for further details).

--no-warn-known-bad-format-calls

Do not warn about calls to string.format or io.format that the compiler knows for sure contain mismatches between the format string and the supplied values.

--warn-unknown-format-calls

Warn about calls to string.format or io.format for which the compiler cannot tell whether there are any mismatches between the format string and the supplied values.

--no-warn-obsolete

Do not warn about calls to predicates or functions that have been marked as obsolete.

--inform-ite-instead-of-switch

Generate informational messages for if-then-elses that could be replaced by switches.

`--no-warn-unresolved-polymorphism`

Do not warn about unresolved polymorphism.

`--warn-suspicious-foreign-procs`

Warn about possible errors in the bodies of foreign procedures. When enabled, the compiler attempts to determine whether the success indicator for a foreign procedure is correctly set, and whether the foreign procedure body contains operations that are not allowed (for example, `return` statements in a C foreign procedure). Note that since the compiler's ability to parse foreign language code is limited some warnings reported by this option may be spurious and some actual errors may not be detected at all.

9.3 Verbosity options

`-v`

`--verbose`

Output progress messages at each stage in the compilation.

`-V`

`--very-verbose`

Output very verbose progress messages.

`-E`

`--verbose-error-messages`

Explain error messages. Asks the compiler to give you a more detailed explanation of any errors it finds in your program.

`--no-verbose-make`

Disable messages about the progress of builds using the `'--make'` option.

`--output-compile-error-lines n`

With `'--make'`, output the first *n* lines of the `'.err'` file after compiling a module (default: 15).

`--verbose-commands`

Output each external command before it is run. Note that some commands will only be printed with `'--verbose'`.

`--verbose-recompilation`

When using `'--smart-recompilation'`, output messages explaining why a module needs to be recompiled.

`--find-all-recompilation-reasons`

Find all the reasons why a module needs to be recompiled, not just the first. Implies `'--verbose-recompilation'`.

-S

--statistics

Output messages about the compiler's time/space usage. At the moment this option implies '**--no-trad-passes**', so you get information at the boundaries between phases of the compiler.

--proc-size-statistics *filename*

Append information about the size of each procedure in the module in terms of goals and variables to the end of the named file.

-N

--debug-modes

Output debugging traces of the mode checking.

--debug-modes-verbose

Output detailed debugging traces of the mode checking.

--debug-modes-pred-id *predid*

With '**--debug-modes**', restrict the debugging traces to the mode checking of the predicate or function with the specified pred id.

--debug-det

--debug-determinism

Output detailed debugging traces of determinism analysis.

--debug-opt

Output detailed debugging traces of the optimization process.

--debug-opt-pred-id *predid*

Output detailed debugging traces of the optimization process only for the predicate/function with the specified pred id. May be given more than once.

--debug-opt-pred-name *name*

Output detailed debugging traces of the optimization process only for the predicate/function with the specified name. May be given more than once.

--debug-pd

Output detailed debugging traces of the partial deduction and deforestation process.

--debug-liveness <n>

Output detailed debugging traces of the liveness analysis of the predicate with the given predicate id.

`--debug-make`

Output detailed debugging traces of the ‘`–make`’ option.

`--debug-intermodule-analysis`

Output detailed debugging traces of the ‘`–intermodule-analysis`’ option.

`--debug-indirect-reuse`

Output detailed debugging traces of the indirect reuse pass of ‘`–structure-reuse`’ option.

9.4 Output options

These options are mutually exclusive. If more than one of these options is specified, only the first in this list will apply. If none of these options are specified, the default action is to compile and link the modules named on the command line to produce an executable.

`-f`

`--generate-source-file-mapping`

Output the module name to file name mapping for the list of source files given as non-option arguments to ‘`mmc`’ to ‘`Mercury.modules`’. This must be done before ‘`mmc --generate-dependencies`’ if there are any modules for which the file name does not match the module name. If there are no such modules the mapping need not be generated.

`-M`

`--generate-dependencies`

Output “Make”-style dependencies for the module and all of its dependencies to ‘`module.dep`’, ‘`module.dv`’ and the relevant ‘`.d`’ files.

`--generate-dependency-file`

Output ‘Make’-style dependencies for the module to ‘`module.d`’.

`--generate-module-order`

Output the strongly connected components of the module dependency graph in top-down order to ‘`module.order`’. Implies ‘`--generate-dependencies`’.

`--generate-standalone-interface` *basename*

Output a stand-alone interface. *basename* is used as the basename of any files generated for the stand-alone interface. (See see [\[Stand-alone Interfaces\]](#), page [\[undefined\]](#) for further details.)

`--generate-mmc-deps`

`--generate-mmc-make-module-dependencies`

Generate dependencies for use by ‘`mmc --make`’ even when using Mmake. This is recommended when building a library for installation.

```

-i
--make-int
--make-interface
    Write the module interface to 'module.int'. Also write the short interface to
    'module.int2'.

--make-short-int
--make-short-interface
    Write the unqualified version of the short interface to 'module.int3'.

--make-priv-int
--make-private-interface
    Write the module's private interface (used for compiling nested sub-modules)
    to 'module.int0'.

--make-opt-int
--make-optimization-interface
    Write information used for inter-module optimization to 'module.opt'.

--make-trans-opt
--make-transitive-optimization-interface
    Write the 'module.trans_opt' file. This file is used to store information used
    for inter-module optimization. The information is read in when the compiler
    is invoked with the '--transitive-intermodule-optimization' option. The
    file is called the "transitive" optimization interface file because a '.trans_opt'
    file may depend on other '.trans_opt' and '.opt' files. In contrast, a '.opt'
    file can only hold information derived directly from the corresponding '.m' file.

--make-xml-documentation
    Output an XML representation of all the declarations in the module into the
    '<module>.xml' file. This XML file can then be transformed via a XSL trans-
    form into another documentation format.

-P
--pretty-print
--convert-to-mercury
    Convert to Mercury. Output to file 'module.ugly'. This option acts as a
    Mercury ugly-printer. (It would be a pretty-printer, except that comments are
    stripped and nested if-then-elses are indented too much — so the result is rather
    ugly.)

--typecheck-only
    Just check the syntax and type-correctness of the code. Don't invoke the mode
    analysis and later passes of the compiler. When converting Prolog code to

```

Mercury, it can sometimes be useful to get the types right first and worry about modes second; this option supports that approach.

`-e`

`--errorcheck-only`

Check the module for errors, but do not generate any code.

`-C`

`--target-code-only`

Generate target code (i.e. C in `'module.c'`, assembler in `'module.s'` or `'module.pic.s'`, IL in `'module.il'`, Java in `'module.java'` or Erlang in `'module.erl'`), but not object code.

`-c`

`--compile-only`

Generate C code in `'module.c'` and object code in `'module.o'` but do not attempt to link the named modules.

`--output-grade-string`

Compute from the rest of the option settings the canonical grade string and print it on the standard output.

`--output-link-command`

Print the command used to link executables to the standard output.

`--output-shared-lib-link-command`

Print the command used to link shared libraries to the standard output.

`--output-libgrades`

Print the list of compilation grades in which a library to be installed should be built to the standard output.

`--output-cc`

Print the name of the C compiler to the standard output.

`--output-cflags`

Print the flags with which the C compiler will be invoked to the standard output.

`--output-library-link-flags`

Print the flags that are passed to the linker in order to link against the current set of libraries. This includes the standard library as well as any other libraries specified via the `'--ml'` option. The flags are printed to the standard output.

9.5 Auxiliary output options

`--smart-recompilation`

When compiling, write program dependency information to be used to avoid unnecessary recompilations if an imported module's interface changes in a way which does not invalidate the compiled code. `--smart-recompilation` does not yet work with `--intermodule-optimization`.

`--no-assume-gmake`

When generating `.d`, `.dep` and `.dv` files, generate Makefile fragments that use only the features of standard make; do not assume the availability of GNU Make extensions. This can make these files significantly larger.

`--trace-level level`

Generate code that includes the specified level of execution tracing. The *level* should be one of `none`, `shallow`, `deep`, `rep` and `default`. See [\[Debugging\]](#), page [\[undefined\]](#).

`--trace-optimized`

Do not disable optimizations that can change the trace.

`--profile-optimized`

Do not disable optimizations that can distort deep profiles.

`--no-delay-death`

When the trace level is `deep`, the compiler normally preserves the values of variables as long as possible, even beyond the point of their last use, in order to make them accessible from as many debugger events as possible. However, it will not do this if this option is given.

`--stack-trace-higher-order`

Enable stack traces through predicates and functions with higher-order arguments, even if stack tracing is not supported in general.

`--generate-bytecode`

Output a bytecode form of the module for use by an experimental debugger.

`--auto-comments`

Output comments in the `module.c` file. This is primarily useful for trying to understand how the generated C code relates to the source code, e.g. in order to debug the compiler. The code may be easier to understand if you also use the `--no-llds-optimize` option.

`-n-`

`--no-line-numbers`

Do not put source line numbers in the generated code. The generated code may be in C (the usual case) or in Mercury (with `--convert-to-mercury`).

`--max-error-line-width <n>`

Set the maximum width of an error message line to `<n>` characters (unless a long single word forces the line over this limit).

--show-dependency-graph

Write out the dependency graph to *module.dependency_graph*.

--imports-graph

Write out the imports graph to *module.imports_graph*. The imports graph contains the directed graph module A imports module B. The resulting file can be processed by the graphviz tools.

-d stage**--dump-hlds stage**

Dump the HLDS (a high-level intermediate representation) after the specified stage number or stage name to '*module.hlds_dump.num-name*'. Stage numbers range from 1 to 599; not all stage numbers are valid. If a stage number is followed by a plus sign, all stages after the given stage will be dumped as well. The special stage name '*all*' causes the dumping of all stages. Multiple dump options accumulate.

--dump-hlds-options options

With '**--dump-hlds**', include extra detail in the dump. Each type of detail is included in the dump if its corresponding letter occurs in the option argument. These details are: a - argument modes in unifications, b - builtin flags on calls, c - contexts of goals and types, d - determinism of goals, e - created, removed, carried, allocated into, and used regions, f - follow_vars sets of goals, g - goal feature lists, i - variables whose instantiation changes, l - pred/mode ids and unify contexts of called predicates, m - mode information about clauses, n - nonlocal variables of goals, p - pre-birth, post-birth, pre-death and post-death sets of goals, r - resume points of goals, s - store maps of goals, t - results of termination analysis, u - unification categories and other implementation details of unifications, v - variable numbers in variable names, z - purity annotations on impure and semipure goals A - argument passing information, B - mode constraint information, C - clause information, D - instmap deltas of goals (meaningful only with i), E - deep profiling information, G - compile-time garbage collection information, I - imported predicates, M - mode and inst information, P - path information, R - live forward use, live backward use and reuse possibilities, S - information about structure sharing, T - type and typeclass information, U - unify and compare predicates, Z - information about globals structs representing call and answer tables.

--dump-hlds-pred-id predid

With '**--dump-hlds**', restrict the output to the HLDS of the predicate or function with the specified pred id. May be given more than once.

--dump-hlds-pred-name name

With '**--dump-hlds**', restrict the output to the HLDS of the predicate or function with the specified name. May be given more than once.

- dump-hlds-file-suffix**
Append the given suffix to the names of the files created by the ‘--dump-hlds’ option.
- dump-same-hlds**
Create a file for a HLDS stage even if the file notes only that this stage is identical to the previously dumped HLDS stage.
- dump-mlds *stage***
Dump the MLDS (a C-like intermediate representation) after the specified stage number or stage name. The MLDS is converted to a C source file/header file pair, which is dumped to ‘*module.c_dump.num-name*’ and ‘*module.h_dump.num-name*’. Stage numbers range from 1 to 99; not all stage numbers are valid. The special stage name ‘all’ causes the dumping of all stages. Multiple dump options accumulate.
- verbose-dump-mlds *stage***
Dump the internal compiler representation of the MLDS, after the specified stage number or stage name, to ‘*module.mlds_dump.num-name*’.
- mode-constraints**
Perform constraint based mode analysis on the given modules. At the moment, the only effect of this is to include more information in HDLS dumps, to allow the constraint based mode analysis algorithm to be debugged.
- simple-mode-constraints**
Ask for the simplified variant of constraint based mode analysis, in which there is only one constraint variable per program variable, rather than one constraint variable per node in the inst graph of a program variable. This option is ignored unless `--mode-constraints` is also given.
- benchmark-modes**
Output information about the performance of the constraint based mode analysis algorithm.
- benchmark-modes-repeat *num***
Specifies the number of times the mode analysis algorithm should run. More repetitions may smooth out fluctuations due to background load or clock granularity. This option is ignored unless `--benchmark-modes` is also given.

9.6 Language semantics options

See the Mercury language reference manual for detailed explanations of these options.

- `--no-reorder-conj`
Execute conjunctions left-to-right except where the modes imply that reordering is unavoidable.
- `--no-reorder-disj`
Execute disjunctions strictly left-to-right.
- `--no-fully-strict`
Allow infinite loops or goals with determinism erroneous to be optimised away.
- `--allow-stubs`
Allow procedures to have no clauses. Any calls to such procedures will raise an exception at run-time. This option is sometimes useful during program development. (See also the documentation for the ‘`--warn-stubs`’ option in [\(undefined\)](#) [Warning options], page [\(undefined\)](#).)
- `--infer-all`
An abbreviation for ‘`--infer-types --infer-modes --infer-det`’.
- `--infer-types`
If there is no type declaration for a predicate or function, try to infer the type, rather than just reporting an error.
- `--infer-modes`
If there is no mode declaration for a predicate, try to infer the modes, rather than just reporting an error.
- `--no-infer-det`
`--no-infer-determinism`
If there is no determinism declaration for a procedure, don’t try to infer the determinism, just report an error.
- `--type-inference-iteration-limit n`
Perform at most n passes of type inference (default: 60).
- `--mode-inference-iteration-limit n`
Perform at most n passes of mode inference (default: 30).

9.7 Termination analysis options

For detailed explanations, see the “Termination analysis” section of the “Implementation-dependent extensions” chapter in the Mercury Language Reference Manual.

```
--enable-term
--enable-termination
```

Enable termination analysis. Termination analysis analyses each mode of each predicate to see whether it terminates. The ‘`terminates`’, ‘`does_not_terminate`’ and ‘`check_termination`’ pragmas have no effect unless termination analysis is enabled. When using termination, ‘`--intermodule-optimization`’ should be enabled, as it greatly improves the accuracy of the analysis.

```
--chk-term
--check-term
--check-termination
```

Enable termination analysis, and emit warnings for some predicates or functions that cannot be proved to terminate. In many cases in which the compiler is unable to prove termination, the problem is either a lack of information about the termination properties of other predicates, or the fact that the program used language constructs (such as higher order calls) which cannot be analysed. In these cases the compiler does not emit a warning of non-termination, as it is likely to be spurious.

```
--verb-chk-term
--verb-check-term
--verbose-check-termination
```

Enable termination analysis, and emit warnings for all predicates or functions that cannot be proved to terminate.

```
--term-single-arg limit
--termination-single-argument-analysis limit
```

When performing termination analysis, try analyzing recursion on single arguments in strongly connected components of the call graph that have up to *limit* procedures. Setting this limit to zero disables single argument analysis.

```
--termination-norm norm
```

The norm defines how termination analysis measures the size of a memory cell. The ‘`simple`’ norm says that size is always one. The ‘`total`’ norm says that it is the number of words in the cell. The ‘`num-data-elems`’ norm says that it is the number of words in the cell that contain something other than pointers to cells of the same type.

```
--term-err-limit limit
--termination-error-limit limit
```

Print at most *n* reasons for any single termination error.

```
--term-path-limit limit
--termination-path-limit limit
```

Perform termination analysis only on predicates with at most *n* paths.

9.8 Compilation model options

The following compilation options affect the generated code in such a way that the entire program must be compiled with the same setting of these options, and it must be linked to a version of the Mercury library which has been compiled with the same setting. (Attempting to link object files compiled with different settings of these options will generally result in an error at link time, typically of the form ‘undefined symbol MR_grade_...’ or ‘symbol MR_runtime_grade multiply defined’.)

The options below must be passed to ‘mgnuc’, ‘c2init’ and ‘ml’ as well as to ‘mmc’. If you are using Mmake, then you should specify these options in the ‘GRADEFLAGS’ variable rather than specifying them in ‘MCFLAGS’, ‘MGNUCFLAGS’ and ‘MLFLAGS’.

9.8.1 Grades and grade components

`-s grade`

`--grade grade`

Select the compilation model. The *grade* should be a ‘.’ separated list of the grade options to set. The grade options may be given in any order. The available options each belong to a set of mutually exclusive alternatives governing a single aspect of the compilation model. The set of aspects and their alternatives are:

What target language to use, what data representation to use, and (for C) what combination of GNU C extensions to use:

‘none’, ‘reg’, ‘jump’, ‘asm_jump’, ‘fast’, ‘asm_fast’, ‘hl’, ‘hlc’, ‘il’, ‘java’ and ‘erlang’ (the default is system dependent).

What garbage collection strategy to use:

‘gc’, and ‘agc’ (the default is no garbage collection).

What kind of profiling to use:

‘prof’, ‘memprof’, and ‘profdeep’ (the default is no profiling).

Whether to enable the trail:

‘tr’ and ‘trseg’ (the default is no trailing).

Whether to use single-precision representation of floating point values:

‘spf’ (the default is to use double-precision floats)

Whether to use dynamically sized stacks that are composed of

small segments: ‘stseg’ (the default is to use fixed size stacks)

What debugging features to enable:

‘debug’ and ‘decldebug’ (the default is no debugging features).

Whether to use a thread-safe version of the runtime environment:

‘par’ (the default is a non-thread-safe environment).

Whether to include support for profiling of parallel programs:

‘threadscope’ (the default is no support for profiling parallel programs).

The default grade is system-dependent; it is chosen at installation time by ‘configure’, the auto-configuration script, but can be overridden with the environment variable ‘MERCURY_DEFAULT_GRADE’ if desired. Depending on your particular installation, only a subset of these possible grades will have been installed. Attempting to use a grade which has not been installed will result in an error at link time. (The error message will typically be something like ‘ld: can’t find library for -lmercury’.)

The tables below show the options that are selected by each base grade and grade modifier; they are followed by descriptions of those options.

<i>Grade</i>	<i>Options implied.</i>
‘none’	--target c --no-gcc-global-registers --no-gcc-nonlocal-gotos --no-asm-labels.
‘reg’	--target c --gcc-global-registers --no-gcc-nonlocal-gotos --no-asm-labels.
‘jump’	--target c --no-gcc-global-registers --gcc-nonlocal-gotos --no-asm-labels.
‘fast’	--target c --gcc-global-registers --gcc-nonlocal-gotos --no-asm-labels.
‘asm_jump’	--target c --no-gcc-global-registers --gcc-nonlocal-gotos --asm-labels.
‘asm_fast’	--target c --gcc-global-registers --gcc-nonlocal-gotos --asm-labels.
‘hlc’	--target c --high-level-code.
‘hl’	--target c --high-level-code --high-level-data.
‘il’	--target il --high-level-code --high-level-data.
‘java’	--target java --high-level-code --high-level-data.
‘erlang’	--target erlang.
‘.gc’	--gc Boehm.
‘.mps’	--gc mps
‘.agc’	--gc accurate.
‘.prof’	--profiling.
‘.memprof’	--memory-profiling.
‘.profdeep’	--deep-profiling.
‘.tr’	--use-trail --no-trail-segments.

```

'.trseg'    --use-trail --trail-segments.
'.tsw'      --record-term-sizes-as-words.
'.tsc'      --record-term-sizes-as-cells.
'.spf'      --single-prec-float
'.stseg'    --stack-segments
'.debug'    --debug.
'.decldebug'
            --decl-debug.
'.par'      --parallel.
'.threadscope'
            --threadscope.

```

9.8.2 Target options

`--target c` (grades: none, reg, jump, fast, asm_jump, asm_fast, hl, hlc)

`--target asm` (grades: hlc)

`--il, --target il` (grades: il)

`--java, --target java` (grades: java)

`--erlang, --target erlang` (grades: erlang)

Specify the target language used for compilation: C, assembler, IL, Java or Erlang. C means ANSI/ISO C, optionally with GNU C extensions (see below). IL means the Intermediate Language of the .NET Common Language Runtime. (IL is sometimes also known as "CIL" or "MSIL".) Targets other than C and Erlang imply `--high-level-code`.

`--il-only`

An abbreviation for `--target il --target-code-only`. Generate IL assembler code in `'module.il'`, but do not invoke ilasm to produce IL object code.

`--dotnet-library-version version-number`

The version number for the mscorlib assembly distributed with the Microsoft .NET SDK.

`--no-support-ms-clr`

Don't use Microsoft CLR specific workarounds in the generated code.

`--support-rotor-clr`

Use ROTOR CLR specific workarounds in the generated code.

`--compile-to-c`

`--compile-to-C`

An abbreviation for `--target c --target-code-only`. Generate C code in `'module.c'`, but do not invoke the C compiler to generate object code.

--java-only

An abbreviation for ‘**--target java --target-code-only**’. Generate Java code in ‘*module.java*’, but do not invoke the Java compiler to produce Java bytecode.

--erlang-only

An abbreviation for ‘**--target erlang --target-code-only**’. Generate Erlang code in ‘*module.erl*’, but do not invoke the Erlang compiler to produce Erlang bytecode.

9.8.3 LLDS back-end compilation model options

--gcc-global-registers (grades: reg, fast, asm_fast)**--no-gcc-global-registers** (grades: none, jump, asm_jump)

Specify whether or not to use GNU C’s global register variables extension. This option is ignored if the ‘**--high-level-code**’ option is enabled.

--gcc-non-local-gotos (grades: jump, fast, asm_jump, asm_fast)**--no-gcc-non-local-gotos** (grades: none, reg)

Specify whether or not to use GNU C’s “labels as values” extension. This option is ignored if the ‘**--high-level-code**’ option is enabled.

--asm-labels (grades: asm_jump, asm_fast)**--no-asm-labels** (grades: none, reg, jump, fast)

Specify whether or not to use GNU C’s asm extensions for inline assembler labels. This option is ignored if the ‘**--high-level-code**’ option is enabled.

--pic-reg (grades: any grade containing ‘.pic.reg’)

[For Unix with intel x86 architecture only.] Select a register usage convention that is compatible with position-independent code (gcc’s ‘-fpic’ option). This is necessary when using shared libraries on Intel x86 systems running Unix. On other systems it has no effect. This option is also ignored if the ‘**--high-level-code**’ option is enabled.

--stack-segments (grades: any grade containing ‘.stseg’)

Specify whether use dynamically sized stacks that are composed of small segments. This can help to avoid stack exhaustion at the cost of increased execution time. This option is ignored if the ‘**--high-level-code**’ option is enabled.

9.8.4 MLDS back-end compilation model option

-H, --high-level-code (grades: hl, hlc, il, java)

Use an alternative back-end that generates high-level code rather than the very low-level code that is generated by our original back-end.

--high-level-data (grades: hl, il, java)

Use an alternative, higher-level data representation that uses structs or classes, rather than treating all objects as arrays.

9.8.5 Optional features compilation model options

--debug (grades: any grade containing `‘.debug’`)

Enables the inclusion in the executable of code and data structures that allow the program to be debugged with `‘mdb’` (see [\[Debugging\]](#), page [\[undefined\]](#)). This option is only supported by the low-level C back-end.

--decl-debug (grades: any grade containing `‘.decldebug’`)

Enables the inclusion in the executable of code and data structures that allow subterm dependency tracking in the declarative debugger. This option is only supported by the low-level C back-end.

--profiling, --time-profiling (grades: any grade containing `‘.prof’`)

Enable time profiling. Insert profiling hooks in the generated code, and also output some profiling information (the static call graph) to the file `‘module.prof’`. See [\[Profiling\]](#), page [\[undefined\]](#). This option is only supported by the C back-ends.

--memory-profiling (grades: any grade containing `‘.memprof’`)

Enable memory profiling. Insert memory profiling hooks in the generated code, and also output some profiling information (the static call graph) to the file `‘module.prof’`. See [\[Using mprof for memory profiling\]](#), page [\[undefined\]](#). This option is only supported by the C back-ends.

--deep-profiling (grades: any grade containing `‘.profdeep’`)

Enable deep profiling by inserting the appropriate hooks in the generated code. This option is only supported by the low-level C back-end.

--coverage-profiling

Enable coverage profiling by inserting coverage points, these are used to measure how execution moves through a procedure. This option only effects programs built in deep profiling grades.

--profile-for-implicit-parallelism

Enable options in the deep profiler that generate suitable profiles for use with profile directed implicit parallelism. This option only effects programs built in deep profiling grades.

--record-term-sizes-as-words (grades: any grade containing `‘.tsw’`)

Record the sizes of terms, using one word as the unit of memory.

`--record-term-sizes-as-cells` (grades: any grade containing `‘.tsc’`)
 Record the sizes of terms, using one cell as the unit of memory.

`--experimental-complexity filename`
 Enable experimental complexity analysis for the predicates listed in the given file. This option is supported for the C back-end, with `‘--no-high-level-code’`. For now, this option itself is for developers only.

`--gc {none, Boehm, mps, accurate, automatic}`
`--garbage-collection {none, Boehm, mps, accurate, automatic}`
 Specify which method of garbage collection to use. Grades containing `‘java’`, `‘il’` or `‘erlang’` use `‘--gc automatic’`, grades containing `‘.gc’` use `‘--gc Boehm’`, grades containing `‘.mps’` use `‘--gc mps’`, other grades use `‘--gc none’`. `‘conservative’` or `‘Boehm’` is Hans Boehm et al’s conservative garbage collector. `‘accurate’` is our own type-accurate copying collector. It requires `‘--high-level-code’`. `‘mps’` is another conservative collector based on Ravenbrook Limited’s MPS (Memory Pool System) kit. `‘automatic’` means the target language provides it. This is the case for the IL, Java and Erlang back-ends, which always use the underlying implementation’s garbage collector.

`--use-trail` (grades: any grade containing `‘.tr’` or `‘.trseg’`)
 Enable use of a trail. This is necessary for interfacing with constraint solvers, or for backtrackable destructive update. This option is only supported by the C back-ends.

`--trail-segments` (grades: any grade containing `‘.trseg’`)
 Enable the use of a dynamically sized trail that is composed of small segments. This can help to avoid trail exhaustion at the cost of increased execution time. This option is only supported by the C back-ends.

`--parallel`
 In low-level C grades this enables support for parallel execution. Parallel execution can be achieved by using either the parallel conjunction operator or the concurrency support in the `‘thread’` module of the standard library. See [section “Goals” in *The Mercury Language Reference Manual*](#), and See [section “thread” in *The Mercury Library Reference Manual*](#).

In high-level C grades this enables support for concurrency, which is accessible via the `‘thread’` module in the standard library. The runtime uses POSIX threads to achieve this, therefore it may also support parallel execution of threads.

The Java, IL and Erlang grades support concurrency without this option. Parallel execution may also be available depending on the target’s runtime.

--threadscope

Enable support for threadscope profiling. This enables runtime support for profiling the parallel execution of programs, See [\[Using threadscope\]](#), page [\[Using threadscope\]](#).

--maybe-thread-safe {yes, no}

Specify how to treat the ‘maybe_thread_safe’ foreign code attribute. ‘yes’ means that a foreign procedure with the ‘maybe_thread_safe’ option is treated as though it has a ‘thread_safe’ attribute. ‘no’ means that the foreign procedure is treated as though it has a ‘not_thread_safe’ attribute. The default is ‘no’.

--single-prec-float (grades: any grade containing ‘.spf’)

Use single precision floats so that, on 32-bit machines floating point values don’t need to be boxed. Double precision floats are used by default. This option is not yet supported for the IL or Java back-ends. This option will not be supported for the Erlang back-end.

9.8.6 Developer compilation model options

Of the options listed below, the ‘--num-tag-bits’ option may be useful for cross-compilation, but apart from that these options are all experimental and are intended for use by developers of the Mercury implementation rather than by ordinary Mercury programmers.

--tags {none, low, high}

(This option is not intended for general use.)

Specify whether to use the low bits or the high bits of each word as tag bits (default: low).

--num-tag-bits *n*

(This option is not intended for general use.)

Use *n* tag bits. This option is required if you specify ‘--tags high’. With ‘--tags low’, the default number of tag bits to use is determined by the auto-configuration script.

--num-reserved-addresses *n*

(This option is not intended for general use.)

Treat the integer values from 0 up to *n* − 1 as reserved addresses that can be used to represent nullary constructors (constants) of discriminated union types.

--num-reserved-objects *n*

(This option is not intended for general use.)

Allocate up to *n* − 1 global objects for representing nullary constructors (constants) of discriminated union types.

Note that reserved objects will only be used if reserved addresses (see `--num-reserved-addresses`) are not available, since the latter are more efficient.

`--no-type-layout`

(This option is not intended for general use.)

Don't output `base_type.layout` structures or references to them. This option will generate smaller executables, but will not allow the use of code that uses the layout information (e.g. `'functor'`, `'arg'`). Using such code will result in undefined behaviour at runtime. The C code also needs to be compiled with `'-DNO_TYPE_LAYOUT'`.

9.9 Code generation options

`--low-level-debug`

Enables various low-level debugging stuff that was in the distant past used to debug the Mercury compiler's low-level code generation. This option is not likely to be useful to anyone except the Mercury implementors. It causes the generated code to become very big and very inefficient, and slows down compilation a lot.

`--pic` Generate position independent code. This option is only used by the `'--target asm'` back-end. The generated assembler will be written to `'module.pic.s'` rather than to `'module.s'`.

`--no-trad-passes`

The default `'--trad-passes'` completely processes each predicate before going on to the next predicate. This option tells the compiler to complete each phase of code generation on all predicates before going on the next phase on all predicates.

`--no-reclaim-heap-on-nondet-failure`

Don't reclaim heap on backtracking in nondet code.

`--no-reclaim-heap-on-semidet-failure`

Don't reclaim heap on backtracking in semidet code.

`--no-reclaim-heap-on-failure`

Combines the effect of the two options above.

`max-jump-table-size n`

The maximum number of entries a jump table can have. The special value 0 indicates the table size is unlimited. This option can be useful to avoid exceeding fixed limits imposed by some C compilers.

%

```

%

"--compare-specialization n
    % % Generate quadratic instead of linear compare predicates for types % with
    up to n function symbols. Higher values of n lead to faster % but also bigger
    compare predicates.
%

%

--no-should-pretest-equality
    % % % If specified, do not add a test for the two values being equal as words
    % to the starts of potentially expensive unify and compare predicates.

--fact-table-max-array-size size
    Specify the maximum number of elements in a single 'pragma fact_table' data
    array (default: 1024). The data for fact tables is placed into multiple C arrays,
    each with a maximum size given by this option. The reason for doing this is
    that most C compilers have trouble compiling very large arrays.

--fact-table-hash-percent-full percentage
    Specify how full the 'pragma fact_table' hash tables should be allowed to get.
    Given as an integer percentage (valid range: 1 to 100, default: 90). A lower
    value means that the compiler will use larger tables, but there will generally be
    less hash collisions, so it may result in faster lookups.

```

9.9.1 Code generation target options

The following options allow the Mercury compiler to optimize the generated C code based on the characteristics of the expected target architecture. The default values of these options will be whatever is appropriate for the host architecture that the Mercury compiler was installed on, so normally there is no need to set these options manually. They might come in handy if you are cross-compiling. But even when cross-compiling, it's probably not worth bothering to set these unless efficiency is absolutely paramount.

```

--have-delay-slot
    (This option is not intended for general use.)
    Assume that branch instructions have a delay slot.

--num-real-r-regs n
    (This option is not intended for general use.)
    Assume r1 up to rn are real general purpose registers.

--num-real-f-regs n
    (This option is not intended for general use.)
    Assume f1 up to fn are real floating point registers.

```

`--num-real-r-temps n`

(This option is not intended for general use.)

Assume that *n* non-float temporaries will fit into real machine registers.

`--num-real-f-temps n`

(This option is not intended for general use.)

Assume that *n* float temporaries will fit into real machine registers.

9.10 Optimization options

9.10.1 Overall optimization options

`-O n`

`--opt-level n`

`--optimization-level n`

Set optimization level to *n*. Optimization levels range from -1 to 6. Optimization level -1 disables all optimizations, while optimization level 6 enables all optimizations except for the cross-module optimizations listed below. Some experimental optimizations (for example constraint propagation) are not be enabled at any optimization level.

In general, there is a trade-off between compilation speed and the speed of the generated code. When developing, you should normally use optimization level 0, which aims to minimize compilation time. It enables only those optimizations that in fact usually *reduce* compilation time. The default optimization level is level 2, which delivers reasonably good optimization in reasonable time. Optimization levels higher than that give better optimization, but take longer, and are subject to the law of diminishing returns. The difference in the quality of the generated code between optimization level 5 and optimization level 6 is very small, but using level 6 may increase compilation time and memory requirements dramatically.

Note that if you want the compiler to perform cross-module optimizations, then you must enable them separately; the cross-module optimizations are not enabled by any ‘-O’ level, because they affect the compilation process in ways that require special treatment by ‘`mmake`’.

`--opt-space`

`--optimize-space`

Turn on optimizations that reduce code size and turn off optimizations that significantly increase code size.

`--intermod-opt`

`--intermodule-optimization`

Perform inlining and higher-order specialization of the code for predicates or functions imported from other modules.

`--trans-intermod-opt`

`--transitive-intermodule-optimization`

Use the information stored in '*module.trans_opt*' files to make intermodule optimizations. The '*module.trans_opt*' files are different to the '*module.opt*' files as '*.trans_opt*' files may depend on other '*.trans_opt*' files, whereas each '*.opt*' file may only depend on the corresponding '*.m*' file. Note that '`--transitive-intermodule-optimization`' does not work with '`mmc --make`'.

`--no-read-opt-files-transitively`

Only read the inter-module optimization information for directly imported modules, not the transitive closure of the imports.

`--use-opt-files`

Perform inter-module optimization using any '*.opt*' files which are already built, e.g. those for the standard library, but do not build any others.

`--use-trans-opt-files`

Perform inter-module optimization using any '*.trans_opt*' files which are already built, e.g. those for the standard library, but do not build any others.

`--intermodule-analysis`

Perform analyses such as termination analysis and unused argument elimination across module boundaries. This option is not yet fully implemented.

`--analysis-repeat n`

The maximum number of times to repeat analyses of suboptimal modules with '`--intermodule-analyses`' (default: 0). This option only works with '`mmc --make`'.

9.10.2 High-level (HLDS -> HLDS) optimization options

These optimizations are high-level transformations on our HLDS (high-level data structure).

`--no-inlining`

Disable all forms of inlining.

`--no-inline-simple`

Disable the inlining of simple procedures.

`--no-inline-builtins`

Generate builtins (e.g. arithmetic operators) as calls to out of line procedures. This is done by default when debugging, as without this option the execution of builtins is not traced.

`--no-inline-single-use`

Disable the inlining of procedures called only once.

- `--inline-compound-threshold threshold`
 Inline a procedure if its size (measured roughly in terms of the number of connectives in its internal form), multiplied by the number of times it is called, is below the given threshold.
- `--inline-simple-threshold threshold`
 Inline a procedure if its size is less than the given threshold.
- `--intermod-inline-simple-threshold threshold`
 Similar to `--inline-simple-threshold`, except used to determine which predicates should be included in `.opt` files. Note that changing this between writing the `.opt` file and compiling to C may cause link errors, and too high a value may result in reduced performance.
- `--inline-vars-threshold threshold`
 Don't inline a call if it would result in a procedure containing more than *threshold* variables. Procedures containing large numbers of variables can cause slow compilation.
- `--loop-invariants`
 Optimize loop invariants by moving computations within a loop that are the same on every iteration to the outside so they are only calculated once. (This is a separate optimization to `--optimize-rl-invariants`.)
- `--no-common-struct`
 Disable optimization of common term structures.
- `--constraint-propagation`
 Enable the constraint propagation transformation, which attempts to transform the code so that goals which can fail are executed as early as possible.
- `--local-constraint-propagation`
 Enable the constraint propagation transformation, but only rearrange goals within each procedure. Specialized versions of procedures will not be created.
- `--no-follow-code`
 Don't migrate builtin goals into branched goals.
- `--optimize-unused-args`
 Remove unused predicate arguments. The compiler will generate more efficient code for polymorphic predicates.
- `--intermod-unused-args`
 Perform unused argument removal across module boundaries. This option implies `--optimize-unused-args` and `--intermodule-optimization`.
- `--unneeded-code`
 Remove goals from computation paths where their outputs are not needed, provided the language semantics options allow the deletion or movement of the goal.

- `--unneeded-code-copy-limit limit`
Gives the maximum number of places to which a goal may be copied when removing it from computation paths on which its outputs are not needed. A value of zero forbids goal movement and allows only goal deletion; a value of one prevents any increase in the size of the code.
- `--optimize-higher-order`
Specialize calls to higher-order predicates where the higher-order arguments are known.
- `--type-specialization`
Specialize calls to polymorphic predicates where the polymorphic types are known.
- `--user-guided-type-specialization`
Enable specialization of polymorphic predicates for which there are ‘:- pragma type_spec’ declarations. See the “Type specialization” section in the “Pragmas” chapter of the Mercury Language Reference Manual for more details.
- `--higher-order-size-limit limit`
Set the maximum goal size of specialized versions created by ‘--optimize-higher-order’ and ‘--type-specialization’. Goal size is measured as the number of calls, unifications and branched goals.
- `--higher-order-arg-limit limit`
Set the maximum size of higher-order arguments to be specialized by ‘--optimize-higher-order’ and ‘--type-specialization’.
- `--optimize-constant-propagation`
Evaluate constant expressions at compile time.
- `--introduce-accumulators`
Attempt to introduce accumulating variables into procedures, so as to make the procedure tail recursive.
- `--optimize-constructor-last-call`
Enable the optimization of “last” calls that are followed by constructor application.
- `--optimize-dead-procs`
Enable dead procedure elimination.
- `--excess-assign`
Remove excess assignment unifications.

--no-optimize-format-calls

Do not attempt to interpret the format string in calls to `string.format` and related predicates at compile time; always leave this to be done at runtime.

--optimize-duplicate-calls

Optimize away multiple calls to a predicate with the same input arguments.

--delay-constructs

Reorder goals to move construction unifications after primitive goals that can fail.

--optimize-saved-vars

Minimize the number of variables that have to be saved across calls.

--deforestation

Enable deforestation. Deforestation is a program transformation whose aim is to avoid the construction of intermediate data structures and to avoid repeated traversals over data structures within a conjunction.

--deforestation-depth-limit *limit*

Specify a depth limit to prevent infinite loops in the deforestation algorithm. A value of -1 specifies no depth limit. The default is 4.

--deforestation-vars-threshold *threshold*

Specify a rough limit on the number of variables in a procedure created by deforestation. A value of -1 specifies no limit. The default is 200.

--deforestation-size-threshold *threshold*

Specify a rough limit on the size of a goal to be optimized by deforestation. A value of -1 specifies no limit. The default is 15.

--analyse-exceptions

Try to identify those procedures that cannot throw an exception. This information can be used by some optimization passes.

--analyse-trail-usage

Enable trail usage analysis. Identify those procedures that will not modify the trail. This information is used to reduce the overhead of trailing.

--analyse-mm-tabling

Identify those goals that do not call procedures that are evaluated using minimal model tabling. This information is used to reduce the overhead of minimal model tabling.

- control-granularity**
Don't try to generate more parallelism than the machine can handle, which is specified using `-parallelism-target`.
- distance-granularity *distance_value***
Control the granularity of parallel execution using the specified distance value.
- parallelism-target *num_cpus***
Specifies the number of CPUs of the target machine, for use by `-control-granularity` option.
- implicit-parallelism**
Introduce parallel conjunctions where it could be worthwhile (implicit parallelism) using deep profiling feedback information generated by `mdprof_feedback`. The profiling feedback file can be specified using the `-feedback-file` option.

9.10.3 MLDS backend (MLDS -> MLDS) optimization options

These optimizations are applied to the medium level intermediate code.

- no-mlds-optimize**
Disable the MLDS -> MLDS optimization passes.
- no-optimize-tailcalls**
Treat tailcalls as ordinary calls rather than optimizing by turning self-tailcalls into loops.
- no-optimize-initializations**
Leave initializations of local variables as assignment statements, rather than converting such assignments statements into initializers.
- eliminate-local-variables**
Eliminate local variables with known values, where possible, by replacing occurrences of such variables with their values.
- no-generate-trail-ops-inline**
Do not generate trailing operations inline, but instead insert calls to the versions of these operations in the standard library.

9.10.4 Erlang (HLDS -> ELDS) optimization options

- erlang-switch-on-strings-as-atoms**
Enable a workaround for slow HiPE compilation of large string switches by converting the string to an atom at runtime and switching on atoms. Do not enable if the string switched on could be longer than 255 characters at runtime.

9.10.5 Medium-level (HLDS -> LLDS) optimization options

These optimizations are applied during the process of generating low-level intermediate code from our high-level data structure.

--no-static-ground-terms

Disable the optimization of constructing constant ground terms at compile time and storing them as static constants. Note that auxiliary data structures created by the compiler for purposes such as debugging will still be created as static constants.

--no-smart-indexing

Generate switches as a simple if-then-else chains; disable string hashing and integer table-lookup indexing.

--dense-switch-req-density *percentage*

The jump table generated for an atomic switch must have at least this percentage of full slots (default: 25).

--dense-switch-size *size*

The jump table generated for an atomic switch must have at least this many entries (default: 4).

--lookup-switch-req-density *percentage*

The lookup tables generated for an atomic switch in which all the outputs are constant terms must have at least this percentage of full slots (default: 25).

--lookup-switch-size *size*

The lookup tables generated for an atomic switch in which all the outputs are constant terms must have at least this many entries (default: 4).

--string-switch-size *size*

The hash table generated for a string switch must have at least this many entries (default: 8).

--tag-switch-size *size*

The number of alternatives in a tag switch must be at least this number (default: 3).

--try-switch-size *size*

The number of alternatives in a try-chain switch must be at least this number (default: 3).

--binary-switch-size *size*

The number of alternatives in a binary search switch must be at least this number (default: 4).

`--no-use-atomic-cells`

Don't use the atomic variants of the Boehm gc allocator calls, even when this would otherwise be possible.

`--no-middle-rec`

Disable the middle recursion optimization.

`--no-simple-neg`

Don't generate simplified code for simple negations.

9.10.6 Low-level (LLDS -> LLDS) optimization options

These optimizations are transformations that are applied to our low-level intermediate code before emitting C code.

`--no-common-data`

Disable optimization of common data structures.

`--no-layout-common-data`

Disable optimization of common subsequences in layout structures.

`--no-llds-optimize`

Disable the low-level optimization passes.

`--no-optimize-peep`

Disable local peephole optimizations.

`--no-optimize-jumps`

Disable elimination of jumps to jumps.

`--no-optimize-fulljumps`

Disable elimination of jumps to ordinary code.

`--pessimize-tailcalls`

Disable the optimization of tailcalls.

`--checked-nondet-tailcalls`

Convert nondet calls into tail calls whenever possible, even when this requires a runtime check. This option tries to minimize stack consumption, possibly at the expense of speed.

`--no-use-local-vars`

Disable the transformation to use local variables in C code blocks wherever possible.

`--no-optimize-labels`

Disable elimination of dead labels and code.

- `--optimize-dups`
Enable elimination of duplicate code within procedures.
- `--optimize-proc-dups`
Enable elimination of duplicate procedures.
- `--no-optimize-frames`
Disable stack frame optimizations.
- `--no-optimize-delay-slot`
Disable branch delay slot optimizations.
- `--optimize-reassign`
Optimize away assignments to locations that already hold the assigned value.
- `--optimize-repeat n`
Iterate most optimizations at most *n* times (default: 3).

9.10.7 Output-level (LLDS -> C) optimization options

These optimizations are applied during the process of generating C intermediate code from our low-level data structure.

- `--no-emit-c-loops`
Use only gotos — don't emit C loop constructs.
- `--use-macro-for-redo-fail`
Emit the fail or redo macro instead of a branch to the fail or redo code in the runtime system.
- `--procs-per-c-function n`
Don't put the code for more than *n* Mercury procedures in a single C function. The default value of *n* is one. Increasing *n* can produce slightly more efficient code, but makes compilation slower. Setting *n* to the special value zero has the effect of putting all the procedures in a single function, which produces the most efficient code but tends to severely stress the C compiler.
- `--no-local-thread-engine-base`
Don't copy the thread-local Mercury engine base address into local variables. This option only affects low-level parallel grades not using the global register variables extension.

9.11 Build system options

- m
- make Treat the non-option arguments to ‘mmc’ as files to make, rather than source files. Create the specified files, if they are not already up-to-date. (Note that this option also enables ‘--use-subdirs’.)

- r
- rebuild Same as ‘--make’, but always rebuild the target files even if they are up to date.

- k
- keep-going With ‘--make’ keep going as far as possible even if an error is detected.

- j <n>
- jobs <n> With ‘--make’, attempt to perform up to ‘<n>’ jobs concurrently for some tasks. In low-level C grades with parallelism support, the number of threads is also limited by the ‘-P’ option in the ‘MERCURY_OPTIONS’ environment variable (see [\[Environment\]](#), page [\[Environment\]](#)).

- track-flags With ‘--make’, keep track of the options used when compiling each module. If an option for a module is added or removed, ‘mmc --make’ will then know to recompile the module even if the timestamp on the file itself has not changed. Warning, verbosity and build system options are not tracked.

- pre-link-command *command* Specify a command to run before linking with ‘mmc --make’. This can be used to compile C source files which rely on header files generated by the Mercury compiler. The command will be passed the names of all of the source files in the program or library, with the source file containing the main module given first.

- extra-init-command *command* Specify a command to produce extra entries in the ‘.init’ file for a library. The command will be passed the names of all of the source files in the program or library, with the source file containing the main module given first.

- install-prefix *dir* Specify the directory under which to install Mercury libraries.

- install-command *command* Specify the command to use to install the files in Mercury libraries. The given command will be invoked as *command source target* to install each file in a Mercury library. The default command is ‘cp’.

- install-command-dir-option** *option*
Specify the flag to pass to the install command to install a directory. The given command will be invoked as *command option source target* to install each directory. The default option is ‘-r’.
- libgrade** *grade*
Add *grade* to the list of compilation grades in which a library to be installed should be built.
- no-libgrade**
Clear the list of compilation grades in which a library to be installed should be built. The main use of this is to avoid building and installing the default set of grades.
- libgrades-include-component** *component*
--libgrades-include *component*
Remove grades that do not contain the specified component from the set of library grades to be installed. (This option does not work with Mmake, only ‘mmc --make’.)
- libgrades-exclude-component** *component*
--libgrade-exclude *component*
Remove grades that contain the specified component from the set of library grades to be installed. (This option does not work with Mmake, only ‘mmc --make’.)
- lib-linkage** {shared,static}
Specify whether libraries should be installed for shared or static linking. This option can be specified multiple times. By default libraries will be installed for both shared and static linking.
- flags** *file*
--flags-file *file*
Take options from the specified file, and handle them as if they were specified on the command line.
- options-file** *file*
Add *file* to the list of options files to be processed. If *file* is ‘-’, an options file will be read from the standard input. By default the file ‘Mercury.options’ in the current directory will be read. See [\(undefined\)](#) [Using Mmake], page [\(undefined\)](#) for a description of the syntax of options files.
- config-file** *file*
Read the Mercury compiler’s configuration information from *file*. If the ‘--config-file’ option is not set, a default configuration will be used, unless

‘`--no-mercury-stdlib-dir`’ is passed to ‘`mmc`’. The configuration file is just an options file (see [Using Mmake](#), page [119](#)).

`--options-search-directory dir`

Add *dir* to the list of directories to be searched for options files.

`--mercury-configuration-directory dir`

`--mercury-config-dir dir`

Search *dir* for Mercury system’s configuration files.

`-I dir`

`--search-directory dir`

Append *dir* to the list of directories to be searched for imported modules.

`--intermod-directory dir`

Append *dir* to the list of directories to be searched for ‘`.opt`’ files.

`--use-search-directories-for-intermod`

Append the arguments of all `-I` options to the list of directories to be searched for ‘`.opt`’ files.

`--no-libgrade-install-check`

Do not check that libraries have been installed before attempting to use them. (This option is only meaningful with ‘`mmc -make`’.)

`--use-subdirs`

Create intermediate files in a ‘`Mercury`’ subdirectory, rather than in the current directory.

`--use-grade-subdirs`

Generate intermediate files in a ‘`Mercury`’ subdirectory, laid out so that multiple grades can be built simultaneously. Executables and libraries will be symlinked or copied into the current directory. ‘`--use-grade-subdirs`’ does not work with Mmake (it does work with ‘`mmc --make`’).

`--order-make-by-timestamp`

Make ‘`mmc --make`’ compile more recently modified source files first.

`--show-make-times`

Report run times for commands executed by ‘`mmc --make`’.

`--extra-library-header file`

`--extra-lib-header file`

Install the specified C header file along with a Mercury library. (This option is only supported by ‘`mmc --make`’)

--restricted-command-line

Enable this option if your shell doesn't support long command lines. This option uses temporary files to pass arguments to sub-commands. (This option is only supported by 'mmc --make')

9.12 Miscellaneous options

-?

-h

--help Print a usage message.

--filenames-from-stdin

Read then compile a newline terminated module name or file name from the standard input. Repeat this until EOF is reached. (This allows a program or user to interactively compile several modules without the overhead of process creation for each one.)

--typecheck-ambiguity-warn-limit *n*

Set the number of type assignments required to generate a warning about highly ambiguous overloading to *n*.

--typecheck-ambiguity-error-limit *n*

Set the number of type assignments required to generate an error excessively ambiguous overloading to *n*. If this limit is reached, the typechecker will not process the predicate or function any further.

--feedback-file

Use the specified profiling feedback file which may currently only be processed for implicit parallelism.

9.13 Target code compilation options

If you are using Mmake, you need to pass these options to the target code compiler (e.g. 'mgnuc') rather than to 'mmc'.

--target-debug

Enable debugging of the generated target code. If the target language is C, this has the same effect as '--c-debug' (see below). If the target language is IL, this causes the compiler to pass '/debug' to the IL assembler.

--cc *compiler-name*

Specify which C compiler to use.

--c-include-directory *dir*

--c-include-dir *dir*

Append *dir* to the list of directories to be searched for C header files. Note that if you want to override this list, rather than append to it, then you can set the 'MERCURY_MC_ALL_C_INCL_DIRS' environment variable to a sequence of '--c-include-directory' options.

--c-debug

Pass the '-g' flag to the C compiler, to enable debugging of the generated C code, and also disable stripping of C debugging information from the executable. Since the generated C code is very low-level, this option is not likely to be useful to anyone except the Mercury implementors, except perhaps for debugging code that uses Mercury's C interface extensively.

--no-c-optimize

Don't enable the C compiler's optimizations.

--no-ansi-c

Don't specify to the C compiler that the ANSI dialect of C should be used. Use the full contents of system headers, rather than the ANSI subset.

--inline-alloc

Inline calls to 'GC_malloc()'. This can improve performance a fair bit, but may significantly increase code size. This option has no effect if '--gc boehm' is not set or if the C compiler is not GNU C.

--cflags *options***--cflag** *option*

Specify options to be passed to the C compiler. '--cflag' should be used for single words which need to be quoted when passed to the shell.

--javac *compiler-name***--java-compiler** *compiler-name*

Specify which Java compiler to use. The default is 'javac'.

--java-interpreter *interpreter-name*

Specify which Java interpreter to use. The default is 'java'.

--java-flags *options***--java-flag** *option*

Specify options to be passed to the Java compiler. '--java-flag' should be used for single words which need to be quoted when passed to the shell.

--java-classpath *dir*

Set the classpath for the Java compiler.

- `--java-object-file-extension extension`
Specify an extension for Java object (bytecode) files. By default this is `‘.class’`.
- `--erlang-compiler compiler-name`
Specify which Erlang compiler to use. The default is `‘erlc’`.
- `--erlang-interpreter interpreter-name`
Specify which Erlang interpreter to use. The default is `‘erl’`.
- `--erlang-flags options`
- `--erlang-flag option`
Specify options to be passed to the Erlang compiler. `‘--erlang-flag’` should be used for single words which need to be quoted when passed to the shell.
- `--erlang-include-directory dir`
- `--erlang-include-dir dir`
Append *dir* to the list of directories to be searched for Erlang header files (`.hrl`).
- `--erlang-native-code`
Add `‘+native’` to the start of flags passed to the Erlang compiler. Cancelled out by `‘--no-erlang-native-code’` so it’s useful when you wish to enable native code generation for all modules except for a select few.
- `--no-erlang-inhibit-trivial-warnings`
Do not add `‘+nowarn_unused_vars +nowarn_unused_function’` to the list of flags passed to the Erlang compiler.

9.14 Link options

- `-o filename`
- `--output-file filename`
Specify the name of the final executable. (The default executable name is the same as the name of the first module on the command line, but without the `‘.m’` extension.) This option is ignored by `‘mmc --make’`.
- `--ld-flags options`
- `--ld-flags option`
Specify options to be passed to the command invoked by `‘mmc’` to link an executable. Use `mmc --output-link-command` to find out which command is used. `‘--ld-flag’` should be used for single words which need to be quoted when passed to the shell.
- `--ld-libflags options`

--ld-libflag *option*

Specify options to be passed to the command invoked by ‘mmc’ to link a shared library. Use `mmc --output-shared-lib-link-command` to find out which command is used. ‘--ld-libflag’ should be used for single words which need to be quoted when passed to the shell.

-L *directory***--library-directory** *directory*

Append *directory* to the list of directories in which to search for libraries.

-R *directory***--runtime-library-directory** *directory*

Append *directory* to the list of directories in which to search for shared libraries at runtime.

--shlib-linker-install-name-path *directory*

Specify the path where a shared library will be installed. This option is useful on systems where the runtime search path is obtained from the shared library and not via the -R option (such as Mac OS X).

-l *library***--library** *library*

Link with the specified library.

--link-object *object*

Link with the specified object file.

--search-lib-files-dir *directory***--search-library-files-directory** *directory*

Search *directory* for Mercury library files have not yet been installed. Similar to adding *directory* using all of the ‘--search-directory’, ‘--intermod-directory’, ‘--library-directory’, ‘--init-file-directory’, ‘--c-include-directory’ and ‘--erlang-include-directory’ options, but does the right thing when ‘--use-subdirs’ or ‘--use-grade-subdirs’ options are used.

--mld *directory***--mercury-library-directory** *directory*

Append *directory* to the list of directories to be searched for Mercury libraries. This will add ‘--search-directory’, ‘--library-directory’, ‘--init-file-directory’, ‘--c-include-directory’ and ‘--erlang-include-directory’ options as needed. See [\[Using installed libraries with mmc -make\]](#), [page](#) [\[Using installed libraries with mmc -make\]](#).

- `--ml library`
`--mercury-library library`
Link with the specified Mercury library. See [\[Using installed libraries with `mmc -make`\]](#), page [\[undefined\]](#).
- `--mercury-standard-library-directory directory`
`--mercury-stdlib-dir directory`
Search *directory* for the Mercury standard library. Implies '`--mercury-library-directory directory`' and '`--mercury-configuration-directory directory`'.
- `--no-mercury-standard-library-directory`
`--no-mercury-stdlib-dir`
Don't use the Mercury standard library. Implies '`--no-mercury-configuration-directory`'.
- `--init-file-directory directory`
Append *directory* to the list of directories to be searched for '.init' files by 'c2init'.
- `--init-file file`
Append *file* to the list of '.init' files to be passed to 'c2init'.
- `--trace-init-file file`
Append *file* to the list of '.init' files to be passed to 'c2init' when tracing is enabled.
- `--linkage {shared,static}`
Specify whether to use shared or static linking for executables. Shared libraries are always linked with '`--linkage shared`'.
- `--mercury-linkage {shared,static}`
Specify whether to use shared or static linking when linking an executable with Mercury libraries. Shared libraries are always linked with '`--mercury-linkage shared`'.
- `--no-strip`
Don't strip executables.
- `--no-demangle`
Don't pipe link errors through the Mercury demangler.
- `--no-main`
Don't generate a C `main()` function. The user's code must provide a `main()` function.

- `--no-allow-undefined`
Do not allow undefined symbols in shared libraries.
- `--no-use-readline`
Disable use of the readline library in the debugger.
- `--runtime-flags flags`
Specify flags to pass to the Mercury runtime.
- `--extra-initialization-functions`
- `--extra-inits`
Search ‘.c’ files for extra initialization functions. (This may be necessary if the C files contain hand-coded C code with ‘INIT’ comments, rather than containing only C code that was automatically generated by the Mercury compiler.)
- `--link-executable-command command`
Specify the command used to invoke the linker when linking an executable.
- `--link-shared-lib-command command`
Specify the command used to invoke the linker when linking a shared library.
- `--java-archive-command command`
Specify the command used to produce Java archive (JAR) files.

10 Environment variables

The shell scripts in the Mercury compilation environment will use the following environment variables if they are set. There should be little need to use these, because the default values will generally work fine.

MERCURY_DEFAULT_GRADE

The default grade to use if no ‘`--grade`’ option is specified.

MERCURY_STDLIB_DIR

The directory containing the installed Mercury standard library. ‘`--mercury-stdlib-dir`’ options passed to the ‘`mmc`’, ‘`ml`’, ‘`mgnuc`’ and ‘`c2init`’ scripts override the setting of the MERCURY_STDLIB_DIR environment variable.

MERCURY_NONSHARED_LIB_DIR

For IRIX 5, this environment variable can be used to specify a directory containing a version of libgcc.a which has been compiled with ‘`-mno-abicalls`’. See the file ‘`README.IRIX-5`’ in the Mercury source distribution.

MERCURY_OPTIONS

A list of options for the Mercury runtime system, which gets linked into every Mercury program. The options given in this environment variable apply to every program; the options given in an environment variable whose name is of the form ‘MERCURY_OPTIONS_*progrname*’ apply only to programs named *progrname*. Options may also be set for a particular executable at compile time by passing ‘--runtime-flags’ options to the invocations of ‘ml’ and ‘c2init’ which create that executable. These options are processed first, followed by those in ‘MERCURY_OPTIONS’, with the options in ‘MERCURY_OPTIONS_*progrname*’ being processed last.

The Mercury runtime system accepts the following options.

- C *size* Tells the runtime system to optimize the locations of the starts of the various data areas for a primary data cache of *size* kilobytes. The optimization consists of arranging the starts of the areas to differ as much as possible modulo this size.

- D *debugger* Enables execution tracing of the program, via the internal debugger if *debugger* is ‘i’ and via the external debugger if *debugger* is ‘e’. (The mdb script works by including ‘-Di’ in MERCURY_OPTIONS.) The external debugger is not yet available.

- p Disables profiling. This only has an effect if the executable was built in a profiling grade.

- P *num* Tells the runtime system to use *num* threads if the program was built in a parallel low-level C grade. The Mercury runtime attempts to automatically determine this value if support is available from the operating system. If it cannot or support is unavailable it defaults to ‘1’.

- max-contexts-per-thread *num* Tells the runtime system to create at most *num* contexts per POSIX thread. Each context created requires a set of stacks, setting this value too high can consume excess memory. This only has an effect if the executable was built in a low-level C parallel grade.

- thread-pinning Request that the runtime system attempts to pin Mercury engines (POSIX threads) to CPU cores/hardware threads. This only has an effect if the executable was built in a parallel low-level C grade. This is disabled by default but may be enabled by default in the future.

-T *time-method*

If the executable was compiled in a grade that includes time profiling, this option specifies what time is counted in the profile. *time-method* must have one of the following values:

- 'r'** Profile real (elapsed) time (using ITIMER_REAL).
- 'p'** Profile user time plus system time (using ITIMER_PROF). This is the default.
- 'v'** Profile user time (using ITIMER_VIRTUAL).

Currently, the **'-Tp'** and **'-Tv'** options don't work on Windows, so on Windows you must explicitly specify **'-Tr'**.

--heap-size *size*

Sets the size of the heap to *size* kilobytes.

--heap-size-kwords *size*

Sets the size of the heap to *size* kilobytes multiplied by the word size in bytes.

--detstack-size *size*

Sets the size of the det stack to *size* kilobytes.

--detstack-size-kwords *size*

Sets the size of the det stack to *size* kilobytes multiplied by the word size in bytes.

--nondetstack-size *size*

Sets the size of the nondet stack to *size* kilobytes.

--nondetstack-size-kwords *size*

Sets the size of the nondet stack to *size* kilobytes multiplied by the word size in bytes.

--small-detstack-size *size*

Sets the size of the det stack used for executing parallel conjunctions to *size* kilobytes. The regular det stack size must be equal or greater.

--small-detstack-size-kwords *size*

Sets the size of the det stack used for executing parallel conjunctions to *size* kilobytes multiplied by the word size in bytes. The regular det stack size must be equal or greater.

- small-nondetstack-size** *size*
Sets the size of the nondet stack for executing parallel computations to *size* kilobytes. The regular nondet stack size must be equal or greater.
- small-nondetstack-size-kwords** *size*
Sets the size of the nondet stack for executing parallel computations to *size* kilobytes multiplied by the word size in bytes. The regular nondet stack size must be equal or greater.
- solutions-heap-size** *size*
Sets the size of the solutions heap to *size* kilobytes.
- solutions-heap-size-kwords** *size*
Sets the size of the solutions heap to *size* kilobytes multiplied by the word size in bytes.
- trail-size** *size*
Sets the size of the trail to *size* kilobytes. This option is ignored in grades that use trail segments.
- trail-size-kwords** *size*
Sets the size of the trail to *size* kilobytes multiplied by the word size in bytes. This option is ignored in grades that use trail segments.
- trail-segment-size** *size*
Sets the size of each trail segment to be *size* kilobytes. This option is ignored in grades that do not use trail segments.
- trail-segment-size-kwords** *size*
Set the size of each trail segment to be *size* kilobytes multiplied by the words size in bytes. This option is ignored in grades that do not use trail segments.
- genstack-size** *size*
Sets the size of the generator stack to *size* kilobytes.
- genstack-size-kwords** *size*
Sets the size of the generator stack to *size* kilobytes multiplied by the word size in bytes.
- cutstack-size** *size*
Sets the size of the cut stack to *size* kilobytes.

`--cutstack-size-kwords size`
Sets the size of the cut stack to *size* kilobytes multiplied by the word size in bytes.

`--pneystack-size size`
Sets the size of the png stack to *size* kilobytes.

`--pneystack-size-kwords size`
Sets the size of the png stack to *size* kilobytes multiplied by the word size in bytes.

`-i filename`
`--mdb-in filename`
Read debugger input from the file or device specified by *filename*, rather than from standard input.

`-o filename`
`--mdb-out filename`
Print debugger output to the file or device specified by *filename*, rather than to standard output.

`-e filename`
`--mdb-err filename`
Print debugger error messages to the file or device specified by *filename*, rather than to standard error.

`-m filename`
`--mdb-tty filename`
Redirect all three debugger I/O streams — input, output, and error messages — to the file or device specified by *filename*.

`--debug-threads`
Output information to the standard error stream about the locking and unlocking occurring in each module which has been compiled with the C macro symbol ‘MR_DEBUG_THREADS’ defined.

`--tabling-statistics`
Prints statistics about tabling when the program terminates.

`--mem-usage-report prefix`
Print a report about the memory usage of the program when the program terminates. The report is printed to a new file named ‘.mem_usage_report*N*’ for the lowest value of *N* (up to 99) which doesn’t overwrite an existing file. Note that this capability is not supported on all operating systems.

--trace-count

When the program terminates, generate a trace counts file listing all the debugger events the program executed, if the program actually executed any debugger events. If MERCURY_OPTIONS includes the `-trace-count-output-file=filename` option, then the trace counts are put into the file *filename*. If MERCURY_OPTIONS includes the `-trace-count-summary-file=basename` option, then the trace counts are put either in the file *basename* (if it does not exist), or in *basename.N* for the lowest value of the integer *N* which doesn't overwrite an existing file. (There is a limit on the value of *N*; see the option `-trace-count-summary-max` below.) If neither option is specified, then the output will be written to a file with the prefix `'mercury_trace_counts'` and a unique suffix. Specifying both options is an error.

--coverage-test

Act as the `-trace-count` option, except include *all* debugger events in the output, even the ones that were not executed.

--trace-count-if-exec=*prog*

Act as the `-trace-count` option, but only if the executable is named *prog*. This is to allow the collection of trace count information from only one Mercury program even if several Mercury programs are executed with the same setting of MERCURY_OPTIONS.

--coverage-test-if-exec=*prog*

Act as the `-coverage-test` option, but only if the executable is named *prog*. This is to allow the collection of coverage test information from only one Mercury program even if several Mercury programs are executed with the same setting of MERCURY_OPTIONS.

--trace-count-output-file=*filename*

Documented alongside the `-trace-count` option.

--trace-count-summary-file=*basename*

Documented alongside the `-trace-count` option.

--trace-count-summary-max=*N*

If MERCURY_OPTIONS includes both the `-trace-count` option (or one of the other options that imply `-trace-count`) and the `-trace-count-summary-file=basename` option, then the generated program will put the generated trace counts either in *basename* (if it does not exist), or in *basename.N* for the lowest value of the integer *N* which doesn't overwrite an existing file. The `-trace-count-summary-max` option specifies the maximum value of this *N*. When this maximum is reached, the program will invoke the `'mtc_union'` program to

summarize *basename*, *basename.1*, ... *basename.N* into a single file *basename*, and delete the rest. By imposing a limit on the total number (and hence indirectly on the total size) of these trace count files, this mechanism allows the gathering of trace count information from a large number of program runs. The default maximum value of *N* is 20.

`--boehm-gc-free-space-divisor=N`

This option sets the value of the free space divisor in the Boehm garbage collector to *N*. It is meaningful only when using the Boehm garbage collector. The default value is 3. Increasing its value will reduce heap space but increase collection time. See the Boehm GC documentation for details.

`--boehm-gc-calc-time`

This option enables code in the Boehm garbage collector to calculate the time spent doing garbage collection so far. Its only effect is to enable the ‘`report_stats`’ predicate in the ‘`benchmarking`’ module of the standard library to report this information.

`--fp-rounding-mode mode`

Set the rounding mode for floating point operations to *mode*. Recognised modes are ‘`downward`’, ‘`upward`’, ‘`to_nearest`’ and ‘`toward_zero`’. Exactly what modes are available and even if it is possible to set the rounding mode is system dependent.

`MERCURY_COMPILER`

Filename of the Mercury Compiler.

`MERCURY_MKINIT`

Filename of the program to create the ‘`*_init.c`’ file.

`MERCURY_DEBUGGER_INIT`

Name of a file that contains startup commands for the Mercury debugger. This file should contain documentation for the debugger command set, and possibly a set of default aliases.

11 Using a different C compiler

The Mercury compiler takes special advantage of certain extensions provided by GNU C to generate much more efficient code. We therefore recommend that you use GNU C for compiling Mercury programs. However, if for some reason you wish to use another compiler, it is possible to do so. Here’s what you need to do.

- Create a new configuration for the Mercury system using the `mercury_config` script, specifying the different C compiler, e.g. `mercury_config --output-prefix=/usr/local/mercury-cc --with-cc=cc`.
- Add the `bin` directory of the new configuration to the beginning of your `PATH`.
- You must use a grade beginning with `none`, `hlc` or `hl` (e.g. `hlc.gc`). You can specify the grade in one of three ways: by setting the `MERCURY_DEFAULT_GRADE` environment variable, by adding a line `GRADE=...` to your `Mmake` file, or by using the `--grade` option to `mmc`. (You will also need to install those grades of the Mercury library, if you have not already done so.)
- If your compiler is particularly strict in enforcing ANSI compliance, you may also need to compile the Mercury code with `--no-static-ground-terms`.

12 Foreign language interface

The Mercury foreign language interfaces allows `pragma foreign_proc` to specify multiple implementations (in different foreign programming languages) for a procedure.

If the compiler generates code for a procedure using a back-end for which there are multiple applicable foreign languages, it will choose the foreign language to use for each procedure according to a builtin ordering.

If the language specified in a `foreign_proc` is not available for a particular backend, it will be ignored.

If there are no suitable `foreign_proc` clauses for a particular procedure but there are Mercury clauses, they will be used instead.

<code>'C'</code>	This is the default foreign language on all backends which compile to C or assembler. Only available on backends that compile to C or assembler.
<code>'C#'</code>	Only available on backends that compile to IL. This is the second preferred foreign language for IL code generation.
<code>'IL'</code>	IL is the intermediate language of the .NET Common Language Runtime (sometimes also known as CIL or MSIL). Only available on backends that compile to IL. This is the preferred foreign language for IL code generation.
<code>'Java'</code>	This is the only foreign language for backends which compile to Java.

13 Stand-alone Interfaces

Programs written in a language other than Mercury should not make calls to foreign exported Mercury procedures unless the Mercury runtime has been initialised. (In the case where the Mercury runtime has not been initialised the behaviour of these calls is undefined.) Such programs must also ensure that any module specific initialisation is performed before calling foreign exported procedures in Mercury modules. Likewise, module specific finalisation may need to be performed after all calls to Mercury procedures have been made.

A stand-alone interface provides a mechanism by which non-Mercury programs may initialise (and shut down) the Mercury runtime plus a specified set of Mercury libraries.

A stand-alone interface is created by invoking the compiler with the ‘`--generate-standalone-interface`’ option. The set of Mercury libraries to be included in the stand-alone interface is given via one of the usual mechanisms for specifying what libraries to link against, e.g. the ‘`--ml`’ and ‘`--mld`’ options. (see [\[Libraries\]](#), page [\(undefined\)](#)). The Mercury standard library is always included in this set.

In C grades the ‘`--generate-standalone-interface`’ option causes the compiler to generate an object file that should be linked into the executable. This object file contains two functions: `mercury_init()` and `mercury_terminate()`. The compiler also generates a C header file that contains the prototypes of these functions. (This header file may be included in C++ programs.) The roles of the two functions are described below.

- `mercury_init()`

Prototype:

```
void mercury_init(int argc, char **argv, void *stackbottom);
```

Initialise the Mercury runtime, standard library and any other Mercury libraries that were specified when the stand-alone interface was generated. *argc* and *argv* are the argument count and argument vector, as would be passed to the function `main()` in a C program. *stackbottom* is the address of the base of the stack. In grades that use conservative garbage collection this is used to tell the collector where to begin tracing. This function must be called before any Mercury procedures and must only be called once. It is recommended that the value of *stackbottom* be set by passing the address of a local variable in the `main()` function of a program, for example:

```
int main(int argc, char **argv) {
    void *dummy;
    mercury_init(argc, argv, &dummy);
    ...
}
```

Note that the address of the stack base should be word aligned as some garbage collectors rely upon this. (This is why the type of the dummy variable in the above example is `void *`.)

- `mercury_terminate()`

Prototype:

```
int mercury_terminate(void);
```

Shut down the Mercury runtime. The value returned by this function is Mercury’s exit status (as set by the predicate `io.set_exit_status/3`). This function will also invoke any finalisers contained in the set of libraries for which the stand-alone interface was generated.

The basename of the object and header file are provided as the argument of ‘`--generate-standalone-interface`’ option.

Stand-alone interfaces are not currently supported for target languages other than C. For an example of using a stand-alone interface see the ‘`samples/c_interface/standalone_c`’ directory in the Mercury distribution.

Index

(Index is nonexistent)